

Template Matrix/Vector Library for C++

User Manual

Version 0.64

Mike Jarvis

June 1, 2010

Contents

1	Overview	5
1.1	Features of TMV	5
1.2	Basic usage	8
1.3	Data types	8
1.4	Notations used in this document	9
2	Vectors	10
2.1	Constructors	10
2.2	Initialization	11
2.3	Access	12
2.4	Views	13
2.5	Functions of a vector	14
2.5.1	Non-modifying functions	14
2.5.2	Modifying functions	16
2.6	Arithmetic	18
2.6.1	Operators	18
2.6.2	Subroutines	18
2.7	Input/Output	19
3	Dense Rectangular Matrices	20
3.1	Constructors	20
3.2	Initialization	22
3.3	Access	22
3.4	Views	23
3.5	Functions of a matrix	26
3.5.1	Non-modifying functions	26
3.5.2	Modifying functions	28
3.6	Arithmetic	29
3.6.1	Basic Operators	29
3.6.2	Outer Products	30
3.6.3	Subroutines	30
3.7	Matrix division	31
3.7.1	Operators	31
3.7.2	Least-square solutions	32

3.7.3	Decompositions	32
3.7.4	Pseudo-inverse	33
3.7.5	Efficiency issues	33
3.7.6	Determinants	34
3.7.7	Accessing the decompositions	35
3.7.8	Singular matrices	36
3.8	Input/Output	38
3.9	Small matrices	38
3.9.1	Constructors	39
3.9.2	Access	40
3.9.3	Functions	40
3.9.4	Limitations	40
4	Diagonal matrices	42
4.1	Constructors	42
4.2	Access	43
4.3	Functions	43
4.4	Arithmetic	44
4.5	Division	44
4.6	Input/Output	45
5	Upper/lower triangle matrices	46
5.1	Constructors	46
5.2	Access	47
5.3	Functions	48
5.4	Arithmetic	48
5.5	Division	49
5.6	Input/Output	50
6	Band-diagonal matrices	51
6.1	Constructors	52
6.2	Access	54
6.3	Functions	55
6.4	Arithmetic	56
6.5	Division	56
6.6	Input/Output	58
7	Symmetric and hermitian matrices	59
7.1	Constructors	59
7.2	Access	60
7.3	Functions	61
7.4	Arithmetic	62
7.5	Division	64
7.6	Input/Output	66
8	Symmetric and hermitian band matrices	67
8.1	Constructors	67
8.2	Access	69
8.3	Functions	70
8.4	Arithmetic	71
8.5	Division	71

8.6	Input/Output	73
9	Permutations	74
9.1	Constructors	74
9.2	Access	74
9.3	Functions	74
9.4	Arithmetic	75
9.5	Input/Output	75
10	Errors and Exceptions	77
10.1	FailedAssert	77
10.2	Singular	78
10.3	ReadError	78
10.4	NonPosDef	78
10.5	Warnings	79
11	Advanced Usage	80
11.1	Eigenvalues and eigenvectors	80
11.2	Matrix decompositions	81
11.3	Update a QR decomposition	86
11.4	Downdate a QR decomposition	88
11.4.1	The update algorithm	89
11.4.2	The downdate algorithm	90
11.5	Other SymMatrix operations	90
11.6	Element-by-element product	91
11.7	BaseMatrix views	92
11.8	Iterators	92
11.9	Direct memory access	93
11.10	“Linear” views	94
11.11	Getting the Version of TMV	95
12	Obtaining and Compiling the Library	96
12.1	Obtaining the source code	96
12.2	Installing using SCons	96
12.3	Installing using make	101
12.4	Compiler flags	103
12.5	Installing using CMake	105
12.6	Installation using Microsoft Visual C++	106
12.7	Known compiler issues	107
13	Example Code	112
13.1	Vector	113
13.2	Matrix	118
13.3	Division	126
13.4	BandMatrix	132
13.5	SymMatrix	140
14	Known Bugs and Deficiencies (aka To Do List)	149
15	History	152

1 Overview

First, this library is provided without any warranty of any kind. There is no guarantee that the code will produce accurate results for all inputs. If someone dies because my code gave you a wrong result, do not blame me.

OK, that was mostly for the crazies out there. Really, I think this is a pretty good product, and I've been using it for my own research extensively. So at least for the routines that I use, they are probably reasonably well debugged. I also have a test suite, which tries to be comprehensive, although occasionally I find bugs that I hadn't thought to test for in the test suite, so there may still be a few lurking in there. That means the code should definitely be considered a beta-type release. Hence the "0." version number. I'm also still adding functionality, so there may be interface changes from one version to the next if I decide I want to do something a different way. Just be warned. Any such changes will be mentioned in §15 of subsequent releases.

1.1 Features of TMV

The Template Matrix/Vector (TMV) Library is a C++ class library designed to make writing code with vectors and matrices both transparent and fast. Transparency means that when you look at your code months later, it is obvious what the code does, making it easier to debug. Fast means the execution time of the code should be as fast as possible - this is mostly algorithm dependent, so we want the underlying library code to use the fastest algorithms possible.

If there were another free C++ Matrix library available that satisfied these requirements and provided all (or even most) of the functionality I wanted, I probably would not have written this. But, at least when I started writing this, the available matrix libraries were not very good. Either they didn't have good operator overloading, or they didn't do complex matrices well, or they didn't include singular value decompositions, or something. Anyway, since I did decide to write my own library, hopefully other people can benefit from my efforts and will find this to be a useful product.

Given the above basic guidelines, the specific design features that I have incorporated into the code include:

1. Operator overloading

Matrix equations look like real math equations in the code. For example, one can write

```
v2 = m * v1;  
v2 += 5 * m.transpose() * v1;  
m *= 3.;  
v2 += m1*v1 + 3*v2 + 8*m2.transpose()*v3;
```

to perform the corresponding mathematical operations.

I also define division:

```
x = b/A;
```

to mean solve the matrix equation $Ax = b$. If A has more rows than columns, then the solution will be a least-square solution.

2. Delayed evaluation

Equations like the above

```
v2 = m * v1;  
v2 += 5 * m.transpose() * v1;
```

do not create a temporary vector before assigning or adding the result to `v2`. This makes the TMV code just as fast as code which uses a functional form, such as:

```

dgemv('N',m,n,1.,mp,ldm,vlp,s1,0.,v2p,s2);
dgemv('T',m,n,5.,mp,ldm,vlp,s1,1.,v2p,s2);

```

In fact, when installed with a BLAS library, the first version with TMV just calls those exact BLAS functions given by the second version. So speed is not sacrificed for the sake of code legibility.

However, a more complicated equation like

```

v2 += m1*v1 + 3*v2 + 8*m2.transpose()*v3;

```

does not have a specialized routine, so it will require a couple temporary `Vectors`. Generally, if a statement performs just one operation, no temporary will be needed. (This includes all operations with corresponding BLAS functions along with some others that are not included in BLAS.) More complicated equations like this last example will give the right answer, but may not be quite as efficient as if you expand the code to perform one operation per line.

3. Template classes

Of course, all of the matrix and vector classes are templates, so you can have

```

Matrix<float>
Matrix<double>
Matrix<complex<double> >
Matrix<long double>
Matrix<MyQuadPrecisionType>

```

or whatever.

4. Mix complex/real

One can multiply a real matrix by a complex vector without having to copy the matrix to a complex one for the sake of the calculation and deal with the concomitantly slower calculation. Likewise for the other arithmetic operations.

However, it does not allow mixing of underlying data types (float with double, for example), with the exception of simple assignments.

5. Views

Operations like `m.transpose()` or `v.subVector(3,8)` return “views” of the underlying data rather than copying to new storage. This model helps delay calculations, which increases efficiency. And the syntax is fairly obvious. For example:

```

v.subVector(3,8) *= 3.;
m.row(3) += 4. * m.row(0);
m *= m.transpose();

```

modifies the underlying `v` and `m` in the obvious ways.

Note that in the last equation, `m.transpose()` uses the same storage as `m`, which is getting overwritten in the process. The code recognizes this conflict and uses temporary storage to obtain the correct result. See “Alias checking” below for more about this.

6. C- or Fortran-style indexing

Both C- and Fortran-style (i.e. zero-based or one-based) indexing are possible for element access of the matrices and vectors.

With C-style indexing, all matrix and vector indexing starts with 0. So the upper-left element of a matrix is `m(0,0)`, not `m(1,1)`. Likewise, the lower right element of an $M \times N$ matrix is `m(M-1,N-1)`. For

element ranges, such as `v.subVector(0,10)`, the first number is the index of the first element, and the second number is “one-past-the-end” of the range. So, this would return a 10 element vector from `v(0)` to `v(9)` inclusive, not an 11 element one including `v(10)`.

With Fortran-style indexing, all matrix and vector indexing starts with 1. So the upper-left element of a matrix is `m(1,1)`. Likewise, the lower right element of an $M \times N$ matrix is `m(M,N)`. For element ranges, such as `v.subVector(1,10)`, the first number is the index of the first element, and the second number is the last element. So, this would return a 10 element vector from `v(1)` to `v(10)` inclusive, which represents the same actual elements as the C-style example above.

7. Special matrices

Many applications use matrices with known sparsity or a specific structure. The code is able to exploit a number of these structures for increased efficiency in both speed and storage. So far the following special matrices are available: diagonal, upper/lower triangle, symmetric, hermitian, banded, and symmetric or hermitian banded. Special types of banded matrices, such as upper and lower banded, tridiagonal, or Hessenberg, may all be declared as a regular `BandMatrix`. The code checks the number of sub- and super-diagonals and uses the right algorithm when such specialization is advantageous for a particular calculation.

8. Flexible storage

Both row-major and column-major storage are possible as an optional extra template parameter. For banded matrices, there is also diagonal-major storage. This can aid I/O, which may require a particular format to mesh with another program. Also, some algorithms are faster for one storage than the other, so it can be worth switching storage and checking the speed difference.

9. Alias checking

Expressions such as `m *= m` pose a problem for many matrix libraries, since no matter what order you do the calculation, you will necessarily overwrite elements that you need for later stages of the calculation. The TMV code automatically recognizes the conflict (generally known as an alias) and creates the needed temporary storage.

The code only checks the addresses of the first elements of the different objects. So expressions such as `m = m.lowerTri() * m.upperTri()` will work correctly, even though there are three types of matrices involved, since the address of the upper-left corner of each matrix is the same. (This particular expression does not even need a temporary. The code orders the steps of this calculation so it can be done in place.)

However, `v.subVector(5,15) += v.subVector(0,10)` will be calculated incorrectly, since the subvectors start at different locations, so the code doesn’t notice the aliasing. Here, elements 5-9 will be overwritten before they are added to the left-side vector.

Therefore, some care is still needed. But this limited check is sufficient for most applications.

10. BLAS

For the combinations of types for which there are existing BLAS routines, the code can call the optimized BLAS routines instead of its own code. For other combinations (or for user defined types like `MyQuadPrecisionType` or `somesuch`), the code does its best to order the steps of the calculation, including using blocking techniques, to be reasonably efficient, but it is definitely not as fast as BLAS for most machines.

This feature can be turned off at compile time if desired with the compilation flag `-DNOBLAS`, although this is not generally recommended if BLAS is available on your machine, and speed is important for your application. For some operations, the BLAS routines can be a factor of 10 or more faster than the native TMV code. (I’m working on reducing this gap, but I don’t really expect to ever do better than a factor of 2 slower than an optimized BLAS library.)

11. LAPACK

When possible, the code can call LAPACK routines, which may be faster than the native TMV code. For types that don't have LAPACK routines, the code uses blocked and/or recursive algorithms, which are similarly fast. Again, this feature can be turned off at compile time, this time with `-DNOLAP` (`-DNOBLAS` necessarily turns off the LAPACK calls as well.)

For almost all algorithms, the TMV code is approximately as fast as LAPACK routines - sometimes faster, since most LAPACK distributions do not use recursive algorithms yet, which are generally slightly faster on modern machines with good compilers. So if you don't want to deal with getting LAPACK up and running, it won't generally be too bad, speedwise, to turn off the LAPACK calls.

The only real exception to this statement is the eigenvector calculation for a hermitian matrix. I have not yet implemented the new RRR (Relatively Robust Representation) algorithm by Dhillon. So if the code is spending a significant time calculating eigenvectors, it may be worth having it call the LAPACK routines.

1.2 Basic usage

All of the basic TMV classes and functions, including the `Vector` and `Matrix` classes, can be accessed with

```
#include "TMV.h"
```

This file includes all the other files for the basic TMV routines. The special matrices described below (other than diagonal and triangular matrices) are not included in this header file. See their sections for the names of the files to include to access those classes.

All of the TMV classes and functions reside in the namespace `tmv`. And of course, they are all templates. So if you want to declare a 10×10 `Matrix`, one would write:

```
tmv::Matrix<double> m(10,10);
```

If writing `tmv::` all the time is cumbersome, one can use `using` statements near the top of the code:

```
using tmv::Matrix;
using tmv::Vector;
```

Or, while generally considered bad coding style, one can import the whole namespace:

```
using namespace tmv;
```

Or, you could use `typedef` to avoid having to write the template type as well:

```
typedef tmv::Matrix<double> DMatrix;
typedef tmv::Vector<double> DVector;
```

In this documentation, I will usually write `tmv::` with the class names to help remind the reader that it is necessary, especially near the beginnings of the sections. But for the sake of brevity and readability, I sometimes omit it.

1.3 Data types

Throughout most of the documentation, I will write `T` for the underlying type. Wherever you see `T`, you should put `double` or `std::complex<float>` or whatever.

For a user defined type, like `MyQuadPrecisionType` for example, the main requirements are that in addition to the usual arithmetic operators, the functions:

```
std::numeric_limits<T>::epsilon()
sqrt(T x)
exp(T x)
log(T x)
```


need to be defined appropriately, where `T` is your type name. See §12 for details about compiling the library for types other than `double` and `float`.

Some functions in this documentation will return a real value or require a real argument, even if `T` is complex. In these cases, I will write `RT` to indicate “the real type associated with `T`”. Similarly, there are a couple places where `CT` indicates “the complex type associated with `T`”.

It may be worth noting that `Matrix<int>` is possible as well if you compile the library with the `-DINST_INT` flag. However, only simple routines like multiplication and addition will give correct answers. If you try to divide by a `Matrix<int>`, for example, the required calculations are impossible for `int`’s, so the result will not be correct. But since the possibility of multiplication of integer matrices seemed desirable, we do allow them to be used. *Caveat programor*. If debugging is turned on (or more accurately, not turned off via the compile flag `-DNDEBUG`), then trying to do anything that requires `sqrt` or `epsilon` for `ints` will result in a runtime error.

1.4 Notations used in this document

There are three fonts used in this document. First, times is used for the main text. Second, as you have no doubt already noticed, typewriter font is used to indicate bits of code. And finally, when I discuss the math about matrices, I use *italics* – for example, $v_2 = m * v_1$.

Also, my code syntax in this documentation is not very rigorous, aiming to maximize readability of the code, rather than including all of the type specifications for everything.

I tend to freely mix the syntax of how a function is used with how it is declared in order to try to provide all of the information you will need on one line. For example, the constructor listed as:

```
tmv::Vector<T,index> v(size_t n, const T* vv)
```

is actually declared in the code as:

```
tmv::Vector<T,index>::Vector(size_t n, const T* vv);
```

and when it is used in your source code, you would write something like:

```
size_t n = 5;
const double vv[n] = {1.2, 3.1, 9.1, 9.2, -3.5};
tmv::Vector<double,tmv::CStyle> v(n,vv);
```

So, the notation that I use in the documentation for this constructor is kind of a hybrid between the declaration syntax and the use syntax. The intent is to improve readability, but if you are ever confused about how to use a particular method, you should look at the `.h` header files themselves, since they obviously have the exactly accurate declarations.

2 Vectors

The `Vector` class is our mathematical vector. Not to be confused with the standard template library's `vector` class. Our `Vector` class name is capitalized, while the STL `vector` is not. If this is not enough of a difference for you, and you are using both extensively in your code, we recommend keeping the full `tmv::Vector` designation for ours and `std::vector` for theirs to distinguish them. Or you might want to `typedef tmv::Vector` to some other name.

`Vector` inherits from its base class `GenVector` (i.e. “generic vector”). Most operations that do not modify the data are actually defined in `GenVector` rather than `Vector`, although some routines are overridden in `Vector` for speed reasons.

The other classes that inherit from `GenVector` are `VectorView`, `ConstVectorView` (both described in more detail below - see §2.4), and `VectorComposite`, which is the base class for all arithmetic operations that return a (logical) `Vector`. This means that any of these other objects can be used any time a `Vector` can be used in a non-assignable capacity. For example, `Norm(v1+m*v2)` is completely valid, and will automatically create the necessary temporary `Vector` to store the result of the mathematical operation in the parentheses.

There is another template argument for a `Vector` in addition to `T` (which represents the data type of the elements). The second template argument may be either `tmv::CStyle` or `tmv::FortranStyle`. Or it may be omitted, in which case `CStyle` is assumed. This argument governs how the element access is performed.

With C-style indexing, the first element of a `Vector` of length `N` is `v(0)` and the last is `v(N-1)`. Also, methods that take range arguments use the common C convention of “one-past-the-end” for the last element; so `v.subVector(0, 3)` returns a 3-element vector, not 4.

With Fortran-style indexing, the first element of the vector is `v(1)` and the last is `v(N)`. And ranges indicate the first and last elements, so the same subvector as above would now be accessed using `v.subVector(1, 3)` to return the first three elements.

All views of a `Vector` keep the same indexing style as the original unless you explicitly change it with a cast. You can cast a `VectorView<T, CStyle>` as a `VectorView<T, FortranStyle>` and vice versa. Likewise for `ConstVectorView`.

The only thing to watch out for about the indexing is that `GenVector` and `VectorComposite` do not have the extra `index` template argument and are always indexed using the C-style convention. Therefore, if you want to index a `GenVector` using the Fortran-style convention, you would need to recast it as an object of type `ConstVectorView<T, FortranStyle>`. A `VectorComposite` would not generally be indexed, but if you did want to do so using the Fortran-style conventions, you would need to explicitly instantiate it as a `Vector<T, FortranStyle>`.

2.1 Constructors

Here, `T` is used to represent the data type of the elements of the `Vector` (e.g. `double`, `complex<double>`, `int`, etc.) and `index` is either `tmv::CStyle` or `tmv::FortranStyle`. In all of the constructors the `index` template argument may be omitted, in which case `CStyle` is assumed.

- `tmv::Vector<T, index> v(size_t n)`

Makes a `Vector` of size `n` with uninitialized values. If debugging is turned on (this is actually the default - turn off debugging by compiling with `-DNDEBUG`), then the values are in fact initialized to 888. This should help you notice when you have neglected to initialize the `Vector` correctly.

- `tmv::Vector<T, index> v(size_t n, T x)`

Makes a `Vector` of size `n` with all values equal to `x`

- `tmv::Vector<T, index> v(size_t n, const T* vv)`
`tmv::Vector<T, index> v(const std::vector<T>& vv)`

Make a `Vector` which copies the elements of `vv`. For the first one, `n` specifies the length. The second gets the length from `vv`.

- `tmv::Vector<T,index> v = tmv::BasisVector<T,index>(size_t n, int i)`

Makes a `Vector` whose elements are all 0, except `v(i) = 1`. Note the `BasisVector` also has the `index` template argument to indicate which element is meant by `v(i)`. Again, if it is omitted, `CStyle` is assumed.

- `tmv::Vector<T,index> v1(const tmv::GenVector<T2>& v2)`
`v1 = v2`

Copy the `Vector` `v2`, which may be of any type `T2` so long as values of type `T2` are convertible into type `T`. The assignment operator has the same flexibility.

- `tmv::VectorView<T,index> v =`
 `tmv::VectorViewOf(T* vv, size_t n, int step = 1)`
`tmv::ConstVectorView<T,index> v =`
 `tmv::VectorViewOf(const T* vv, size_t n, int step = 1)`

Makes a `VectorView` (see §2.4 below) which refers to the exact elements of `vv`, not copying them to new storage. The parameter `n` is the number of values to include in the view. The optional `step` parameter allows a non-unit spacing between successive vector elements in memory.

2.2 Initialization

There are a number of ways to initialize the elements of a `Vector`:

The most straightforward way is to simply set each element individually:

```
for(int i=0; i<N; ++i) {  
    v[i] = /* whatever */  
}
```

But this is often inconvenient, especially if you want to assign a particular list of values that does not lend itself to being put in a `for` loop. Of course you can assign each element one line at a time, but that can be a bit unwieldy. So another way is to use a C array initializer, and then create the `Vector` from these values. For example:

```
double var[5] = { 1.2, 3.2, -9.2, -1, 3.3 };  
tmv::Vector<double> v(5,var);
```

This works, but it seems a bit inefficient to use a temporary variable. The ideal notation in my opinion would be to put the list initializer as a parameter in the constructor:

```
tmv::Vector<double> v(5,  
    { 1.2, 3.2, -9.2, -1, 3.3 }); // Error - invalid
```

Unfortunately, this is not valid C++ syntax¹. One way to get around this is to use:

```
tmv::Vector<double> v(5,  
    (double[]) { 1.2, 3.2, -9.2, -1, 3.3 }); // Usually OK
```

although not all compilers accept this code as valid. I'm not even sure whether it is ANSI-compliant C++ code or not².

¹ Apparently they are working on this shortcoming as part of the new so-called C++0x standards update.

² The common editor, vim, doesn't seem to think so. It flags the braces in red, which is its indicator for a syntax error. But most compilers seem to compile it correctly.

So to address this issue, we decided to follow the syntax of `std::ostream` and use the `<<` operator to initialize a `Vector`. The idea is that we are sending information into the vector.

After the first `<<`, the elements may be separated by either commas or more `<<` operators:

```
tmv::Vector<double> v(5);
v << 1.2, 3.2, -9.2, -1, 3.3;
v << 1.2 << 3.2 << -9.2 << -1 << 3.3;
```

There must be precisely as many values as there are elements in the `Vector`, or a `tmv::ReadError` will be thrown.

One advantage of this method for initialization is that the values do not need to be simple numbers. Normally the elements in a list initializer need to be numerical literals, or at least to be computable at compile time. With the above syntax, the elements can be variables, or even values returned from functions. Anything that is convertible into the element type of the `Vector` will work.

Also, `v` can be re-initialized with this method at any time, so it is not limited to initialization at construction. And finally, `v` can be a `VectorView`, not just a `Vector`.

2.3 Access

- `size_t v.size() const`

Returns the size (length) of `v`.

- `T v[int i] const`
`T v(int i) const`
`T v.cref(int i) const`
`typename Vector<T>::reference v[int i]`
`typename Vector<T>::reference v(int i)`
`typename Vector<T>::reference v.ref(int i)`

The `[]` and `()` forms are equivalent. Each returns the `i`-th element of `v`. With `index = CStyle`, the first element is `v(0)`, and the last element is `v(n-1)`. With `index = FortranStyle`, they are `v(1)` and `v(n)`.

If `v` is a `const` `Vector`, a `ConstVectorView`, or a `GenVector`, then the return type is just the value, not a reference.

If `v` is a non-`const` `Vector`, then the return type is a normal reference, `T&`.

If `v` is a `VectorView`, then the return type is an object, which is an lvalue (i.e. it is assignable), but which may not be `T&`. Specifically, it has the type `typename VectorView<T>::reference`. For a real-typed `VectorView`, it is just `T&`. But for a complex-typed `VectorView`, the return type is an object that keeps track of the possibility of a conjugation.

The main difference between the operator forms and `cref` or `ref` is that the latter versions do not check for the validity of the parameter `i`, even when compiling with debugging turned on. Also, `cref` and `ref` always use `CStyle` indexing.

- `typename tmv::Vector<T>::iterator v.begin()`
`typename tmv::Vector<T>::iterator v.end()`
`typename tmv::Vector<T>::const_iterator v.begin() const`
`typename tmv::Vector<T>::const_iterator v.end() const`
`typename tmv::Vector<T>::reverse_iterator v.rbegin()`
`typename tmv::Vector<T>::reverse_iterator v.rend()`
`typename tmv::Vector<T>::const_reverse_iterator v.rbegin() const`
`typename tmv::Vector<T>::const_reverse_iterator v.rend() const`

These provide iterator-style access into a `Vector`, which works just like the standard template library's iterators. If `v` is a `VectorView`, the iterator types are slightly different from the `Vector` iterators, so you should declare them as `typename tmv::VectorView<T>::iterator`, etc. instead.

2.4 Views

A `VectorView<T>` object refers to the elements of some other object, such as a regular `Vector<T>` or `Matrix<T>`, so that altering the elements in the view alters the corresponding elements in the original object. A `VectorView` can have non-unit steps between elements (for example, a view of a column of a row-major matrix). It can also be a conjugation of the original elements, so that

```
tmv::VectorView<double> cv = v.conjugate();
cv(3) = z;
```

would actually set the original element, `v(3)` to `conj(z)`.

Also, we have to keep track of whether we are allowed to alter the original values or just look at them. Since we want to be able to pass these views around, it turns out that the usual `const`-ing doesn't work the way you would want. Thus, there are two objects that are views of a `Vector`: `ConstVectorView` and `VectorView`. The first is only allowed to view, not modify, the original elements. The second is allowed to modify them. This distinction is akin to the `const_iterator` and `iterator` types in the standard template library.

One slightly non-intuitive thing about `VectorViews` is that a `const VectorView` is still mutable. The `const` in this case means that one cannot change the components to which the view refers. But a `VectorView` is inherently an object that can be used to modify the underlying data, regardless of any `const` in front of it.

The following methods return views to portions of a `Vector`. If `v` is either a (non-`const`) `Vector` or a `VectorView`, then a `VectorView` is returned. If `v` is a `const Vector`, a `ConstVectorView`, or any other `GenVector`, then a `ConstVectorView` is returned.

- `v.subVector(int i1, int i2, int istep=1)`

This returns a view to a subset of the original vector. `i1` is the first element in the subvector. `i2` is either "one past the end" (C-style) or the last element (Fortran-style) of the subvector. `istep` is an optional step size. Thus, if you have a `Vector~v` of length 10, and you want to multiply the first 3 elements by 2, with C-style indexing, you could write:

```
v.subVector(0,3) *= 2.;
```

To set all the even elements to 0, you could write:

```
v.subVector(0,10,2).setZero();
```

And then to output the last 4 elements of `v`, you could write:

```
std::cout << v.subVector(6,10);
```

For Fortran-style indexing, the same steps would be accomplished by:

```
v.subVector(1,3) *= 2.;
```

```
v.subVector(1,9,2).setZero();
```

```
std::cout << v.subVector(7,10);
```

- `v.reverse()`

This returns a view whose elements are the same as `v`, but in the reverse order

- `v.conjugate()`

This returns the conjugate of a `Vector` as a view, so it still points to the same physical elements, but modifying this will set the actual elements in memory to the conjugate of what you set. Likewise, accessing an element will return the conjugate of the value in memory.

- `v.view()`

Returns a view of a `Vector`. This seems at first like a silly function to have, but if you write a function that takes a mutable `Vector` argument, and you want to be able to pass it views in addition to regular `Vectors`, it is easier to write the function once with a `VectorView` parameter. Then you only need a second function with a `Vector` parameter, which calls the first function using `v.view()` as the argument:

```
double foo(const tmv::VectorView<double>& v)
{ ... [modifies v] ... }
double foo(tmv::Vector<double>& v)
{ return foo(v.view()); }
```

If you are not going to be modifying `v` in the function, you only need to write one function, and you should use the base class `GenVector` for the argument type:

```
double foo(const tmv::GenVector<double>& v)
{ ... [doesn't modify v] ... }
```

The arguments could then be a `const Vector`, a `ConstVectorView`, or even a `VectorComposite`.

- `v.cView()`
`v.fView()`

Like `view()` but forces the result to have C- or Fortran-style indexing respectively.

- `v.realPart()`
`v.imagPart()`

These return views to the real and imaginary parts of a complex `Vector`. Note the return type is a real view in each case:

```
tmv::Vector<std::complex<double>> v(10, std::complex<double>(1,4));
tmv::VectorView<double> vr = v.realPart();
tmv::VectorView<double> vi = v.imagPart();
```

2.5 Functions of a vector

Functions that do not modify the `Vector` are defined in `GenVector`, and so can be used for any type derived from `GenVector`: `Vector`, `ConstVectorView`, `VectorView`, or `VectorComposite`.

Functions that modify the `Vector` are only defined for `Vector` and `VectorView`.

2.5.1 Non-modifying functions

Each of the following functions can be written in two ways, either as a method or a function. For example, the expressions:

```
double normv = v.norm();
double normv = Norm(v);
```

are equivalent. In either case, `v` can be any `GenVector`. Also, several of the functions below have multiple equivalent names. For example, `norm1` and `sumAbsElements` are equivalent, so you can use whichever one is clearer to you in your situation. And just to remind you, `RT` refers to the real type associated with `T`. So if `T` is either `double` or `complex<double>`, `RT` would be `double`.

- `RT v.norm1() const`
`RT Norm1(v)`
`RT v.sumAbsElements() const`
`RT SumAbsElements(v)`

The 1-norm of `v`: $\|v\|_1 = \sum_i |v(i)|$.

- `RT v.norm2() const`
`RT Norm2(v)`
`RT v.norm() const`
`RT Norm(v)`

The 2-norm of `v`: $\|v\|_2 = (\sum_i |v(i)|^2)^{1/2}$. This is the most common meaning for the norm of a vector, so we define the `norm` function to be the same as `norm2`.

- `RT v.normSq(const RT scale=1) const`
`RT NormSq(v)`

The square of the 2-norm of `v`: $(\|v\|_2)^2 = \sum_i |v(i)|^2$. In the method version of this function, you may provide an optional scale factor, in which case the return value is equal to `NormSq(scale*v)` instead, which can help avoid underflow or overflow problems.

- `RT v.normInf() const`
`RT NormInf(v)`
`RT v.maxAbsElement() const`
`RT MaxAbsElement(v)`

The infinity-norm of `v`: $\|v\|_\infty = \max_i |v(i)|$.

- `T v.maxElement(int* i=0) const`
`T MaxElement(v)`
`T v.minElement(int* i=0) const`
`T MinElement(v)`

The maximum/minimum element. For complex values, there is no way to define a max or min element, so just the real component of each element is used. The `i` argument is available in the method versions of these function, and it is optional. If it is present (and not 0), then `*i` is set to the index of the max/min element returned.

- `RT v.maxAbsElement(int* i=0) const`
`RT MaxAbsElement(v)`
`RT v.minAbsElement(int* i=0) const`
`RT MinAbsElement(v)`
`RT v.maxAbs2Element(int* i=0) const`
`RT MaxAbs2Element(v)`
`RT v.minAbs2Element(int* i=0) const`
`RT MinAbs2Element(v)`

The maximum/minimum element by absolute value. The `i` argument is available in the method versions of these function, and it is optional. If it is present (and not 0), then `*i` is set to the index of the max/min element returned.

For real vectors, the last four cases with `Abs2` instead of `Abs` are identical, but for complex vectors, they use $|real(v(i))| + |imag(v(i))|$ instead. This is faster than doing the normal absolute value, and for many purposes (such as finding a suitably large value with which to scale a vector), it is just as useful.

- `T v.sumElements() const`
`T SumElements(v)`

The sum of the elements of $v = \sum_i v(i)$.

2.5.2 Modifying functions

The following functions are methods of both `Vector` and `VectorView`, and they work the same way in the two cases, although there may be speed differences between them. All of these are usually written on a line by themselves. However, they do return the (modified) `Vector`, so you can string them together if you want. For example:

```
v.clip(1.e-10).conjugateSelf().reverseSelf();
```

would first clip the elements at $1.e-10$, then conjugate each element, then finally reverse the order of the elements. (This would probably not be considered very good programming style, however.) Likewise, the expression:

```
foo(v.clip(1.e-10));
```

which would first clip the elements at $1.e-10$, then pass the resulting `Vector` to the function `foo`.

- `v.setZero();`

Clear the `Vector` `v`. i.e. Set each element to 0.

- `v.setAllTo(T x);`

Set each element to the value `x`.

- `v.clip(RT thresh)`

Set each element whose absolute value is less than `thresh` equal to 0. Note that `thresh` should be a real value even for complex valued `Vectors`.

- `v.addToAll(T x)`

Add the value `x` to each element.

- `v.conjugateSelf()`

Change each element into its complex conjugate. Note the difference between this and `v.conjugate()`, which returns a view to a conjugated version of `v` without actually changing the underlying data. This function, `v.conjugateSelf()`, does change the underlying data.

- `v.reverseSelf()`

Reverse the order of the elements. Note the difference between this and `v.reverse()` which returns a view to the elements in reversed order.

- `v.makeBasis(int i)`

Set all elements to 0, except for $v(i) = 1$.

- `v.swap(int i1, int i2)`

Swap elements $v(i1)$ and $v(i2)$.

- `v.permute(const int* p)`
`v.permute(const int* p, int i1, int i2)`
`v.reversePermute(const int* p)`
`v.reversePermute(const int* p, int i1, int i2)`

The first one performs a series of swaps: $(v(0), v(p[0])), (v(1), v(p[1])), \dots$ The second starts at $i1$ and ends at $i2-1$ rather than doing the whole range from 0 to $n-1$. The last two work the same way, but do the swaps in the opposite order.

Note: The indices listed in a permutation array (p) always use the C-style convention, even if v uses Fortran-style indexing.

These functions should not generally be needed in normal usage, since we now have a `Permutation` class to encapsulate the most common of these kinds of operations. See §9

- `v.sort(Permutation& p, tmv::ADType ad=tmv::ASCEND,
tmv::COMPTYPE comp=tmv::REAL_COMP)`
`v.sort(tmv::ADType ad=tmv::ASCEND,
tmv::COMPTYPE comp=tmv::REAL_COMP)`

Sorts the vector v , optionally returning the corresponding permutation in p . If you do not care about the permutation, the second form is slightly more efficient.

The next parameter, ad , determines whether the sorted `Vector` will have its elements in ascending or descending order. The possible values are `ASCEND` and `DESCEND`. The default if omitted is to sort in ascending order.

The final parameter, $comp$, determines what component of the elements to use for the sorting. This is especially relevant if T is complex, since complex values are not intrinsically sortable. The possible values are `REAL_COMP`, `ABS_COMP`, `IMAG_COMP`, and `ARG_COMP`. Only the first two make sense for non-complex vectors. The default if omitted is to sort the real values of the elements.

- `Swap(v1, v2)`

Swap the corresponding elements of $v1$ and $v2$. Note that if $v1$ and/or $v2$ are views, this does physically swap the data elements, not just some pointers to the data. This is the intuitive meaning of a statement like

```
Swap(m.row(4), m.row(5));
```

Clearly what is desired by that is to swap the actual values, and this is what we actually do.

However, if $v1$ and $v2$ are both `tmv::Vector` objects, rather than views, then the swap efficiently swaps the pointers to the data, and so takes $O(1)$ time, rather than $O(N)$.

2.6 Arithmetic

2.6.1 Operators

All the usual operators work the way you would expect for `Vectors`. For shorthand in the following list, I use `x` for a scalar of type `T` or `RT`, and `v` for a `Vector`. When there are two `Vectors` listed, they may either be both of the same type `T`, or one may be of type `T` and the other of `complex<T>`. Whenever `v` is an lvalue, it may be either a `Vector` or a `VectorView`. Otherwise, it may be any `GenVector`.

Also, I use the notation `[+-]` to mean either `+` or `-`, since the syntax is generally the same for these operators. Likewise, I use `[*/]` when their syntax is equivalent.

```
v2 = -v1;
v2 = x * v1;
v2 = v1 [*/] x;
v3 = v1 [+-] v2;
v [*/] = x;
v2 [+-] = v1;
x = v1 * v2;
```

The last one, `v1 * v2`, returns the inner product of two vectors, which is a scalar. That is, the product is a row vector times a column vector.

This is the only case (so far) where the specific row or column orientation of a vector matters. For the others listed here, the left side and the right side are implied to be of the same orientation, but that orientation is otherwise arbitrary. Later, when we get to a matrix times a vector, the orientation of the vector will be inferred from context.

2.6.2 Subroutines

Each of the above equations use deferred calculation so that the sum or product is not calculated until the storage is known. The equations can even be a bit more complicated without requiring a temporary. Here are some equations that do not require a temporary `Vector` for the calculation:

```
v2 = -(x1*v1 + x2*v2);
v2 += x1*(x2*(-v1));
v2 -= x1*(v1/=x2);
```

The limit to how complicated the right hand side can be without using a temporary is set by the functions that the code eventually calls to perform the calculation. While you shouldn't ever need to use these directly, it may help you understand when the code will require temporary `Vectors`. If you do use these, note that the `v` parameters are `VectorViews`, rather than `Vectors`. So you would need to call them with `v.view()` if `v` is a `Vector`.

- `MultXV(T x, const VectorView<T>& v)`

Performs the calculation `v *= x`.

- `MultXV(T x, const GenVector<T1>& v1, const VectorView<T>& v2)`

Performs the calculation `v2 = x*v1`.

- `AddVV(T x, const GenVector<T1>& v1, const VectorView<T>& v2)`

Performs the calculation `v2 += x*v1`.

- `AddVV(T x1, const GenVector<T1>& v1, T x2, const GenVector<T2>& v2, const VectorView<T>& v3)`

Performs the calculation $v3 = x1*v1 + x2*v2$.

- `T MultVV(const GenVector<T>& v1, const GenVector<T2>& v2)`

Performs the calculation $v1*v2$.

More complicated arithmetic equations such as

```
v1 += x*(v1+v2+v3) + (x*v3-v1)
```

will require one or more temporary vectors, and so may be less efficient than you might like, but the code should return the correct result, no matter how complicated the equation is.

2.7 Input/Output

The simplest output is the usual:

```
os << v
```

where `os` is any `std::ostream`. The output format is:

```
n ( v(0)  v(1)  v(2)  ...  v(3) )
```

where `n` is the length of the `Vector`.

The same format can be read back in one of two ways:

```
tmv::Vector<T> v(n);  
is >> v;  
std::auto_ptr<tmv::Vector<T> > vptr;  
is >> vptr;
```

For the first version, the `Vector` must already be declared, which requires knowing how big it needs to be. If the input `Vector` does not match in size, an exception of type `tmv::ReadError` is thrown. The second version allows you to automatically get the size from the input. The `Vector` pointed to by `vptr` will be created according to whatever size the input `Vector` is.

Often, it is convenient to output only those values that aren't very small. This can be done using

```
v.write(std::ostream& os, RT thresh)
```

which is equivalent to

```
os << tmv::Vector<T>(v).clip(thresh);
```

but without requiring the temporary `Vector`.

3 Dense Rectangular Matrices

The `Matrix` class is our dense matrix class. It inherits from `GenMatrix`, which has the definitions of all the methods that do not modify the `Matrix`.

The other classes that inherit from `GenMatrix` are `ConstMatrixView`, `MatrixView` (see §3.4 below), and `MatrixComposite`, which is the base class for the various arithmetic operations that return a (logical) `Matrix`.

`GenMatrix` in turn inherits from `BaseMatrix`. All of the various special `Matrix` classes also inherit from `BaseMatrix`. `BaseMatrix` has virtual declarations for the functions that can be performed on any kind of `Matrix` regardless of its structure or sparsity.

In addition to the data type template parameter (indicated here by `T` as usual), there is also a storage template parameter, which may be either `RowMajor` or `ColMajor`.

Finally, there is also a template argument indicating which indexing convention you want the matrix to use, which may be either `CStyle` or `FortranStyle`.

With C-style indexing, the upper-left element of an $M \times N$ `Matrix` is `m(0,0)`, the lower-left is `m(M-1,0)`, the upper-right is `m(0,N-1)`, and the lower-right is `m(M-1,N-1)`. Also, methods that take a pair of indices to define a range use the common C convention of “one-past-the-end” for the meaning of the second index. So `m.subMatrix(0,3,3,6)` returns a 3×3 submatrix.

With Fortran-style indexing, the upper-left element of an $M \times N$ `Matrix` is `m(1,1)`, the lower-left is `m(M,1)`, the upper-right is `m(1,N)`, and the lower-right is `m(M,N)`. Also, methods that take range arguments take the pair of indices to be the actual first and last elements in the range. So `m.subMatrix(1,3,4,6)` returns the same 3×3 submatrix as given above.

All views of a `Matrix` keep the same indexing style as the original unless you explicitly change it with a cast. You can cast a `MatrixView<T,CStyle>` as a `MatrixView<T,FortranStyle>` and vice versa. (Likewise for `ConstMatrixView`.) However, as for `GenVector`, you should be aware that `GenMatrix` and `MatrixComposite` do not have the extra template argument and are always indexed using the C-style convention. So if you want to index one of these using the Fortran-style convention, you need to (respectively) cast the `GenMatrix` as a `ConstMatrixView<T,FortranStyle>` or instantiate the `MatrixComposite` as a `Matrix<T,FortranStyle>`.

You may omit the indexing template argument, in which case `CStyle` is assumed. And if so, you may also then omit the storage argument, in which case `ColMajor` is assumed. If you want to specify `FortranStyle` indexing, you need to include the storage argument.

3.1 Constructors

We use `stor` to indicate the storage template argument. This argument must be either `tmv::RowMajor` or `tmv::ColMajor`. And `index` indicates either `tmv::CStyle` or `tmv::FortranStyle`. The default values for these are `tmv::ColMajor` and `tmv::CStyle` if the arguments are omitted.

- `tmv::Matrix<T,stor,index> m(size_t nrows, size_t ncols)`

Makes a `Matrix` with `nrows` rows and `ncols` columns with uninitialized values. If debugging is turned on (this is actually the default - turn off debugging by compiling with `-DNDEBUG`), then the values are in fact initialized to 888. This should help you notice when you have neglected to initialize the `Matrix` correctly.

- `tmv::Matrix<T,stor,index> m(size_t nrows, size_t ncols, T x)`

Makes a `Matrix` with `nrows` rows and `ncols` columns with all values equal to `x`

- `tmv::Matrix<T,stor,index> m(size_t nrows, size_t ncols, const T* vv)`
`tmv::Matrix<T,stor,index> m(size_t nrows, size_t ncols,`
`const std::vector<T>& vv)`

Makes a `Matrix` with `nrows` rows and `ncols` columns, which copies the elements of `vv`.

If `stor` is `RowMajor`, then the elements of `vv` are taken to be in row-major order: first the `ncols` elements of the first row, then the `ncols` elements of the second row, and so on for the `nrows` rows. Likewise if `stor` is `ColMajor`, then the elements of `vv` are taken to be in column-major order - the elements of the first column, then the second, and so on.

- `tmv::Matrix<T,stor,index> m(const std::vector<std::vector<T> >& vv)`

Makes a `Matrix` with elements `m(i, j) = vv[i][j]`. The size of the `Matrix` is taken from the sizes of the vectors. (If `index` is `FortranStyle`, then `m(i, j) = vv[i-1][j-1]`.)

- `tmv::Matrix<T,index> m1(const GenMatrix<T2>& m2)`
`m1 = m2`

Copy the `Matrix` `m2`, which may be of any type `T2` so long as values of type `T2` are convertible into type `T`. The assignment operator has the same flexibility.

- `tmv::MatrixView<T,index> m =`
`tmv::MatrixViewOf(T* mm, size_t nrows, size_t ncols,`
`StorageType stor)`
`tmv::ConstMatrixView<T,index> m =`
`tmv::MatrixViewOf(const T* mm, size_t nrows, size_t ncols,`
`StorageType stor)`
`tmv::MatrixView<T,index> m =`
`tmv::MatrixViewOf(T* mm, size_t nrows, size_t ncols,`
`int stepi, int stepj)`
`tmv::ConstMatrixView<T,index> m =`
`tmv::MatrixViewOf(const T* mm, size_t nrows, size_t ncols,`
`int stepi, int stepj)`

Makes a `MatrixView` (see §3.4 below) that refers to the exact elements of `mm`, not copying them to new storage. The first two versions indicate the ordering of the elements using the normal `StorageType` designation, which must be either `RowMajor` or `ColMajor`.

The next two versions allow you to provide an arbitrary step through the data in the *i* and *j* directions. This means that (for C-style indexing):

`m(i, j) == *(mm + i*stepi + j*stepj)`

- `tmv::MatrixView<T,index> m =`
`RowVectorViewOf(VectorView<T,index> v)`
`tmv::ConstMatrixView<T,index> m =`
`RowVectorViewOf(ConstVectorView<T,index> v)`

Makes a view of the `Vector`, which treats it as a $1 \times n$ `Matrix` (i.e. a single row).

- `tmv::MatrixView<T,index> m =`
`ColVectorViewOf(VectorView<T,index> v)`
`tmv::ConstMatrixView<T,index> m =`
`ColVectorViewOf(ConstVectorView<T,index> v)`

Makes a view of the `Vector`, which treats it as an $n \times 1$ `Matrix` (i.e. a single column).

3.2 Initialization

A `Matrix` can be initialized with a comma-delimited list using the `<<` operator, just like with a `Vector`. For example:

```
tmv::Matrix<double,tmv::RowMajor> m(3,5);
m << 1.2,  3.2, -9.2, -1.0,  3.3,
      2.0, -2.7,  2.2, -6.5, -7.6,
      -8.2, 3.2, -8.7,  5.1,  1.6;
```

There must be precisely as many values as there are elements in the `Matrix`, or a `tmv::ReadError` will be thrown.

This initialization is most readable when the `Matrix` `m` is `RowMajor`, since then the elements in the list can be arranged as in the above example, which shows the correct matrix structure of the elements. If `m` is `ColMajor`, then the list elements are assigned in order of the columns instead.

Also, the list initialization works for a `MatrixView` as well. However, it only works if the `MatrixView`'s elements are all contiguous in memory. You can check if this is the case with the `canLinearize()` method described in §11.10.

3.3 Access

- `size_t` `m.nrows()` `const`
 `size_t` `m.ncols()` `const`
 `size_t` `m.colsize()` `const`
 `size_t` `m.rowsize()` `const`

Returns the size of each dimension of `m`. `nrows()` and `colsize()` are equivalent. Likewise `ncols` and `rowsize()` are equivalent.

- `T` `m(int i, int j)` `const`
 `T` `m[int i][int j]` `const`
 `T` `m.cref(int i, int j)` `const`
 `typename` `tmv::Matrix<T>::reference` `m(int i, int j)`
 `typename` `tmv::Matrix<T>::reference` `m[int i][int j]`
 `typename` `tmv::Matrix<T>::reference` `m.ref(int i,int j)`

Returns the `i, j` element of `m`. i.e. the `i`th element in the `j`th column. Or equivalently, the `j`th element in the `i`th row.

With C-style indexing, the upper-left element of a `Matrix` is `m(0,0)`, the lower-left is `m(nrows-1,0)`, The upper-right is `m(0,ncols-1)`, and the lower-right is `m(nrows-1,ncols-1)`.

With Fortran-style indexing, these four elements would instead be `m(1,1)`, `m(nrows,1)`, `m(1,ncols)`, and `m(nrows,ncols)`, respectively.

If `m` is a `const` `Matrix`, a `ConstMatrixView`, or a `GenMatrix`, then the return type is just the value, not a reference.

If `m` is a non-`const` `Matrix`, then the return type is a normal reference `T&`.

If `m` is a `MatrixView`, then the return type is an object, which is an `Ivalue` (i.e. it is assignable), but which may not be `T&`. It has the type `typename MatrixView<T>::reference`, (which is the same as `typename VectorView<T>::reference`). It is equal to `T&` for real `MatrixViews`, but is more complicated for complex `MatrixViews` since it needs to keep track of the possibility of conjugation.

The main difference between the operator forms and `cref` or `ref` is that the latter versions do not check for the validity of the parameters `i` and `j`, even when compiling with debugging turned on. Also, `cref` and `ref` always use CStyle indexing.

- `ConstVectorView<T> m.row(int i) const`
`ConstVectorView<T> m.col(int j) const`
`VectorView<T> m.row(int i)`
`VectorView<T> m.col(int j)`

Return a view of the `i`th row or `j`th column respectively. If `m` is mutable (either a non-`const` `Matrix` or a `MatrixView`), then a `VectorView` is returned. Otherwise, a `ConstVectorView` is returned.

- `ConstVectorView<T> m.row(int i, int j1, int j2) const`
`ConstVectorView<T> m.col(int j, int i1, int i2) const`
`VectorView<T> m.row(int i, int j1, int j2)`
`VectorView<T> m.col(int j, int i1, int i2)`

Variations on the above, where only a portion of the row or column is returned.

For example, with C-style indexing, `m.col(3, 2, 6)` returns a 4-element vector view containing the elements `[m(2, 3), m(3, 3), m(4, 3), m(5, 3)]`.

With Fortran-style indexing, the same elements are returned by `m.col(4, 3, 6)`. (And these elements would be called: `[m(3, 4), m(4, 4), m(5, 4), m(6, 4)]`.)

- `ConstVectorView<T> m.diag() const`
`ConstVectorView<T> m.diag(int i) const`
`ConstVectorView<T> m.diag(int i, int k1, int k2) const`
`VectorView<T> m.diag()`
`VectorView<T> m.diag(int i)`
`VectorView<T> m.diag(int i, int k1, int k2)`

Return the diagonal or one of the sub- or super-diagonals. This first one returns the main diagonal. For the second and third, `i=0` refers to the main diagonal; `i>0` are the super-diagonals; and `i<0` are the sub-diagonals. The last version is equivalent to the expression `m.diag(i).subVector(k1, k2)`.

- `ConstVectorView<T> m.subVector(int i, int j, int istep, int jstep,`
`int size) const`
`VectorView<T> m.subVector(int i, int j, int istep, int jstep, int size)`

If the above methods aren't sufficient to obtain the `VectorView` you need, this function is available, which returns a view through the `Matrix` starting at `m(i, j)`, stepping by `(istep, jstep)` between elements, for a total of `size` elements. For example, the diagonal from the lower-left to the upper-right of an $n \times n$ `Matrix` would be obtained by: `m.subVector(n-1, 0, -1, 1, n)` for C-style or `m.subVector(n, 1, -1, 1, n)` for Fortran-style.

3.4 Views

A `MatrixView` object refers to some or all of the elements of a regular `Matrix`, so that altering the elements in the view alters the corresponding elements in the original object. A `MatrixView` can be either row-major, column-major, or neither. That is, the view can span a `Matrix` with non-unit steps in both directions. It can also be a conjugation of the original elements.

There are two view classes for a `Matrix`: `ConstMatrixView` and `MatrixView`. The first is only allowed to view, not modify, the original elements. The second is allowed to modify them.

It is worth pointing out again that a `const` `MatrixView` is still mutable, just like a `const` `VectorView`. The `const` just means that you cannot change which elements the view references.

The following methods return views to portions of a `Matrix`. If `m` is either a (non-`const`) `Matrix` or a `MatrixView`, then a `MatrixView` is returned. If `m` is a `const` `Matrix`, `ConstMatrixView`, or any other `GenMatrix`, then a `ConstMatrixView` will be returned.

- `m.subMatrix(int i1, int i2, int j1, int j2)`
`m.subMatrix(int i1, int i2, int j1, int j2, int istep, int jstep)`

This returns a view to a submatrix contained within the original matrix.

If `m` uses C-style indexing, the upper-left corner of the returned view is `m(i1, j1)`, the lower-left corner is `m(i2-1, j1)`, the upper-right corner is `m(i1, j2-1)`, and the lower-right corner is `m(i2-1, j2-1)`.

If `m` uses Fortran-style indexing, the upper-left corner of the view is `m(i1, j1)`, the lower-left corner is `m(i2, j1)`, the upper-right corner is `m(i1, j2)`, and the lower-right corner is `m(i2, j2)`.

The second version allows for non-unit steps in the two directions. To set a `Matrix` to be a checkerboard of 1's, you could write (for C-style indexing):

```
tmv::Matrix<int> board(8,8,0)
board.subMatrix(0,8,0,8,2,2).setAllTo(1);
board.subMatrix(1,9,1,9,2,2).setAllTo(1);
```

For Fortran-style indexing, the same thing would be accomplished by:

```
tmv::Matrix<int, tmv::ColMajor, tmv::FortranStyle> board(8,8,0)
board.subMatrix(1,7,1,7,2,2).setAllTo(1);
board.subMatrix(2,8,2,8,2,2).setAllTo(1);
```

- `m.rowRange(int i1, int i2)`
`m.colRange(int j1, int j2)`

Since pulling out a bunch of contiguous rows or columns is a common submatrix use, we provide these functions. They are shorthand for

```
m.subMatrix(i1,i2,0,ncols)
m.subMatrix(0,nrows,j1,j2)
```

respectively. (For Fortran-style indexing, replace the 0 with a 1.)

- `m.rowPair(i1,i2)`
`m.colPair(i1,i2)`

Another common submatrix is to select a pair of rows or columns, not necessarily adjacent to each other. These are short hand for:

```
m.subMatrix(i1,i2+(i2-i1),0,ncols,i2-i1,1)
m.subMatrix(0,nrows,j1,j2+(j2-j1),1,j2-j1)
```

respectively. The equivalent in Fortran-style indexing would be:

```
m.subMatrix(i1,i2,1,ncols,i2-i1,1)
m.subMatrix(1,nrows,j1,j2,1,j2-j1).
```


- `m.transpose()`
`m.conjugate()`
`m.adjoint()`

These return the transpose, conjugate, and adjoint (aka conjugate-transpose) of a `Matrix`. They point to the same physical elements as the original matrix, so modifying these will correspondingly modify the original matrix.

Note that some people define the adjoint of a matrix as the determinant times the inverse. This combination is also called the adjugate or the cofactor matrix. It is not the same as our `m.adjoint()`. What we call the adjoint is usually written as m^\dagger , or variously as m^H or m^* , and is sometimes referred to as the hermitian conjugate or (rarely) tranjugate. This definition of the adjoint seems to be the more modern usage. Older texts tend to use the other definition. However, if this is confusing for you, it may be clearer to explicitly write out `m.conjugate().transpose()`, which will not produce any efficiency reduction in your code compared with using `m.adjoint()` (assuming your compiler inlines these methods properly).

- `m.view()`

Returns a view of a `Matrix`. As with the `view()` function for a `Vector`, it is mostly useful for passing a `Matrix` to a function that takes a `MatrixView` argument. This lets you convert the first into the second.

- `m.cView()`
`m.fView()`

Like `view()` but forces the result to have C- or Fortran-style indexing respectively.

- `m.realPart()`
`m.imagPart()`

These return views to the real and imaginary parts of a complex `Matrix`. Note the return type is a real view in each case:

```
tmv::Matrix<std::complex<double>> m(10, std::complex<double>(1, 4));
tmv::MatrixView<double> mr = m.realPart();
tmv::MatrixView<double> mi = m.imagPart();
```

- `m.upperTri(DiagType dt = NonUnitDiag)`
`m.lowerTri(DiagType dt = NonUnitDiag)`
`m.unitUpperTri()`
`m.unitLowerTri()`

These return an `UpperTriMatrixView` or a `LowerTriMatrixView` which views either the upper triangle or the lower triangle of a square `Matrix`. If `m` has more rows than columns, then only `UpperTri` is valid, since the portion below the diagonal is not triangular. Likewise, if `m` has more columns than rows, then only `LowerTri` is valid.

In first two cases, you may provide an optional parameter `dt`, which declares whether the diagonal elements are treated as all 1s (`dt = UnitDiag`) or as their actual values (`dt = NonUnitDiag`). See §5 for more details about this parameter and triangular matrices in general. For now, the latter two versions are equivalent to `m.upperTri(UnitDiag)` and `m.lowerTri(UnitDiag)` respectively. However, future versions of TMV (specifically version 0.70, which is in development) will be able to exploit the compile-time knowledge of `dt` to make this more efficient.

3.5 Functions of a matrix

Functions that do not modify the `Matrix` are defined in `GenMatrix`, and so can be used for any type derived from `GenMatrix`: `Matrix`, `ConstMatrixView`, `MatrixView`, or `MatrixComposite`. Functions that modify the `Matrix` are only defined for `Matrix` and `MatrixView`.

3.5.1 Non-modifying functions

Each of the following functions can be written in two ways, either as a method or a function. For example, the expressions `m.norm()` and `Norm(m)` are equivalent. In each case, `m` can be any `GenMatrix<T>`. As a reminder, `RT` refers to the real type associated with `T`. In other words, `T` is either the same as `RT` or it is `std::complex<RT>`.

- `RT m.norm1()` `const`
`RT Norm1(m)`

The 1-norm of `m`: $\|m\|_1 = \max_j(\sum_i |m(i, j)|)$.

- `RT m.norm2()` `const`
`RT Norm2(m)`
`RT m.doNorm2()` `const`

The 2-norm of `m`: $\|m\|_2$ = the largest singular value of `m`, which is also the square root of the largest eigenvalue of $(m^\dagger m)$.

This function can be fairly expensive if you have not already performed an SV decomposition of `m`, so the first two versions output a warning to `stdout` if there is no SV decomposition already set.

If you understand that you are asking TMV to perform an SV decomposition to calculate $\|m\|_2$, and you are ok with it, you can either change the warning behavior – see §10.5 for how to do this – or you can just call `doNorm2()` instead, which will do the necessary SV decomposition without any warnings.

In fact, if you are not going to need the decomposition for anything else, then `m.doNorm2()` will be faster than `m.svd()`, `m.norm2()`, because the former will not bother to accumulate the singular vectors.

- `RT m.normInf()` `const`
`RT NormInf(m)`

The infinity-norm of `m`: $\|m\|_\infty = \max_i(\sum_j |m(i, j)|)$.

- `RT m.normF()` `const`
`RT NormF(m)`
`RT m.norm()` `const`
`RT Norm(m)`

The Frobenius norm of `m`: $\|m\|_F = (\sum_{i,j} |m(i, j)|^2)^{1/2}$.

This is the most common meaning for the norm of a matrix, so we define the `norm` function to be the same as `normF`.

- `RT m.normSq(RT scale=1)` `const`
`RT NormSq(m)`

The square of the Frobenius norm of `m`: $(\|m\|_F)^2 = \sum_{i,j} |m(i, j)|^2$.

In the method version of this function, you may provide an optional scale factor, in which case the return value is equal to `NormSq(scale*v)` instead, which can help avoid underflow or overflow problems.

- `RT m.maxAbsElement() const`
`RT MaxAbsElement(m)`

The element of m with the maximum absolute value: $\|m\|_{\Delta} = \max_{i,j} |m(i,j)|$.

- `RT m.maxAbs2Element() const`
`RT MaxAbs2Element(m)`

The same as `maxAbsElement`, but using $|real(m(i,j))| + |imag(m(i,j))|$ instead of the normal absolute value for complex matrices.

- `T m.trace() const`
`T Trace(m)`

The trace of m : $\text{Tr}(m) = \sum_i m(i,i)$.

- `T m.sumElements() const`
`T SumElements(m)`

The sum of all elements of m : $\sum_{i,j} m(i,j)$.

- `RT m.sumAbsElements() const`
`RT SumAbsElements(m)`

The sum of the absolute values of all elements of m : $\sum_{i,j} |m(i,j)|$.

- `T m.det() const`
`T Det(m)`

The determinant of m , $\det(m)$. For speed issues regarding this function, see §3.7 below on division.

- `RT m.logDet(T* sign=0) const`
`RT LogDet(m)`

The log of the absolute value of the determinant of m . If the optional argument `sign` is provided (possible in the method version only), then on output `*sign` records the sign of the determinant. See 3.7.6 for more details.

- `bool m.isSingular() const`

Return whether m is singular, i.e. $\det(m) = 0$. (Singular matrices are discussed in more detail in §3.7.8.)

- `RT m.condition() const`
`RT m.doCondition() const`

The condition (technically the 2-condition) is the ratio of the largest singular value of m to the smallest.

Like `norm2`, this function requires a singular value decomposition to be performed, so it can be fairly expensive if you have not already performed an SV decomposition of m . So the first version outputs a warning to `stdout` if there is no SV decomposition already set for the matrix.

And as with `norm2`, you can bypass the warning by calling `doCondition()` instead, which will do the necessary SV decomposition without any warnings. Also, see §10.5 for how to change the warning behavior of TMV.

- `tmv::Matrix<T> minv = m.inverse()`
`tmv::Matrix<T> minv = Inverse(m)`
`void m.makeInverse(Matrix<T>& minv)`

Set `minv` to the inverse of `m`.

If `m` is not square, then `minv` is set to the pseudo-inverse, or an approximate pseudo-inverse. If `m` is singular, then an error may result, or the pseudo-inverse may be returned, depending on the division method specified for the matrix. See §3.7.4 and §3.7.8 on pseudo-inverses and singular matrices for more details.

Note that the first two forms do not actually require a temporary (despite appearances), so they are just as efficient as the third version. This is because `m.inverse()` actually returns an object whose type is derived from `MatrixComposite`. The calculation of the inverse is then delayed until there is a place to store the result.

- `void m.makeInverseATA(Matrix<T>& cov) const`

Set `cov` to be $(m^\dagger m)^{-1}$.

If `m` has more rows than columns, then using it to solve a system of equations really amounts to finding the least-square solution, since there is (typically) no exact solution. When you do this, `m` is known as the “design matrix” of the system, and is commonly called A . Solving $Ax = b$ gives x as the least-square solution. And the covariance matrix for the solution vector x is $\Sigma = (A^\dagger A)^{-1}$. It turns out that computing this matrix is generally easy to do once you have performed the decomposition needed to solve for x (either a QR or SV decomposition - see §3.7.3). Thus this function is provided, which sets the argument `cov` to the equivalent of `Inverse(m.adjoint()*m)`, but generally does so much more efficiently than doing this explicitly, and also probably more accurately.

3.5.2 Modifying functions

The following functions are methods of both `Matrix` and `MatrixView`, and they work the same way for each. As with the `Vector` modifying functions, these all return a reference to the newly modified `Matrix`, so you can string them together if you want.

- `m.setZero()`

Set all elements to 0.

- `m.setAllTo(T x)`

Set all elements to the value `x`.

- `m.clip(RT thresh)`

Set each element whose absolute value is less than `thresh` equal to 0. Note that `thresh` should be a real value even for complex valued `Matrixes`.

- `m.setToIdentity(T x = 1)`

Set `m` to `x` times the identity matrix. If the argument `x` is omitted, it is taken to be 1, so `m` is set to the identity matrix. This is equivalent to `m.setZero().diag().setAllTo(x)`.

- `m.conjugateSelf()`

Conjugate each element. Note the difference between this and `m.conjugate()`, which returns a view to the conjugate of `m` without actually changing the underlying data. Contrariwise, `m.conjugateSelf()` does change the underlying data.

- `m.transposeSelf()`

Transpose the `Matrix`. Note the difference between this and `m.transpose()`, which returns a view to the transpose without actually changing the underlying data.

- `m.swapRows(int i1, int i2)`
`m.swapCols(int j1, int j2)`

Swap the corresponding elements of two rows or two columns.

- `m.permuteRows(const int* p)`
`m.permuteRows(const int* p, int i1, int i2)`
`m.reversePermuteRows(const int* p)`
`m.reversePermuteRows(const int* p, int i1, int i2)`

These are equivalent to the corresponding routines for `Vectors` (`permute` and `reversePermute`), performing a series of `swapRows` commands.

- `m.permuteCols(const int* p)`
`m.permuteCols(const int* p, int j1, int j2)`
`m.reversePermuteCols(const int* p)`
`m.reversePermuteCols(const int* p, int j1, int j2)`

Same as above, but performing the a series of `swapCols` commands.

- `Swap(m1, m2)`

Swap the corresponding elements of `m1` and `m2`. Note that this does physically swap the data elements, not just some pointers to the data, so it takes $O(N^2)$ time.

3.6 Arithmetic

3.6.1 Basic Operators

We'll start with the simple operators that require little explanation aside from the notation: `x` is a scalar, `v` is a `Vector`, and `m` is a `Matrix`. As a reminder, the notation `[+-]` is used to indicate either `+` or `-`. Likewise for `[*/]`.

```
m2 = -m1
m2 = x * m1
m2 = m1 [*/] x
m3 = m1 [+ -] m2
m [*/] = x
m2 [+ -] = m1
v2 = m * v1
v2 = v1 * m
v *= m
m3 = m1 * m2
m2 *= m1
```

Note that the orientation of a vector is inferred from context. $m*v$ involves a column vector, and $v*m$ involves a row vector.

It is sometimes convenient to be able to treat scalars as the scalar times an identity matrix. So the following operations are allowed and use that convention:

```
m2 = m1 [+-] x
m2 = x [+-] m1
m [+-] = x
m = x
```

For example, you could check if a matrix is numerically close to the identity matrix with:

```
if (Norm(m-1.) < 1.e-8) { [...] }
```

3.6.2 Outer Products

For the product of two vectors, there are two orientation choices that make sense mathematically. It could mean a row vector times a column vector, which is called the inner product. Or it could mean a column vector times a row vector, which is called the outer product. We chose to let $v1*v2$ indicate the inner product, since this is far more common. For the outer product, we use a different symbol:

```
m = v1 ^ v2
```

One problem with this choice is that the order of operations in C++ for $^$ is not the same as for $*$. So, one needs to be careful when combining it with other calculations. For example

```
m2 = m1 + v1 ^ v2 // ERROR!
```

will give a compile time error indicating that you can't add a `Matrix` and a `Vector`, because the operator $+$ has higher precedence in C++ than $^$. So you need to write:

```
m2 = m1 + (v1 ^ v2)
```

3.6.3 Subroutines

As with `Vectors`, we try to defer calculations until we have a place to store them. So $m*v$ returns an object that can be assigned to a `Vector`, but hasn't performed the calculation yet.

The limit to how complicated an expression can be without resorting to a temporary object is set by the functions that the code eventually calls to perform the calculation. While you shouldn't ever need to use these directly, it may help you understand when the code will create a temporary `Matrix`³.

- `MultXM(T x, const MatrixView<T>& m)`

Performs the calculation $m *= x$.

- `MultMV<bool add>(T x, const GenMatrix<Tm>& m, const GenVector<Tv1>& v1, const VectorView<T>& v2)`

Performs the calculation $v2 \text{ (+)= } x*m*v1$ where “(+)=” means “+=” if `add` is true and “=” if `add` is false.

- `AddMM(T x, const GenMatrix<T1>& m1, const MatrixView<T>& m2)`

Performs the calculation $m2 += x*m1$.

³If you do use these, note that the last parameters are `MatrixViews`, rather than `Matrixes`. So you would need to call them with `m.view()` if `m` is a `Matrix`. (For `MultMV`, it would be `v.view()`.)

- `AddMM(T x1, const GenMatrix<T1>& m1, T x2, const GenMatrix<T2>& m2, const MatrixView<T>& m3)`

Performs the calculation $m3 = x1 * m1 + x2 * m2$.

- `MultMM<bool add>(T x, const GenMatrix<T1>& m1, const GenMatrix<T2>& m2, const MatrixView<T>& m3)`

Performs the calculation $m3 \text{ (+)= } x * m1 * m2$ where “(+)=” means “+=” if `add` is true and “=” if `add` is false.

- `Rank1Update<bool add>(T x, const GenVector<T1>& v1, const GenVector<T2>& v2, const MatrixView<T>& m)`

Performs the calculation $m \text{ (+)= } x * (v1 \wedge v2)$ where “(+)=” means “+=” if `add` is true and “=” if `add` is false.

3.7 Matrix division

One of the main things people often want to do with a matrix is use it to solve a set of linear equations. The set of equations can be written as a single matrix equation:

$$Ax = b$$

where A is a matrix and x and b are vectors. A and b are known, and one wants to solve for x . Sometimes there are multiple systems to be solved using the same coefficients, in which case x and b become matrices as well.

3.7.1 Operators

Using the TMV classes, one would solve this equations by writing simply:

```
x = b / A
```

Note that this really means $x = A^{-1}b$, which is different from $x = bA^{-1}$. Writing the matrix equation as we did ($Ax = b$) is much more common than writing $xA = b$, so “left-division” is correspondingly much more common than “right-division”. Therefore, it makes sense to use left-division for our definition of the `/` operator.

However, we do allow for the possibility of wanting to right-multiply a vector by A^{-1} (in which case the vector is inferred to be a row-vector). We designate this operation by:

```
x = b % A
```

which means $x = bA^{-1}$.

Given this explanation, the rest of the division operations should be self-explanatory, where we use the notation `[/ %]` to indicate that either `/` or `%` may be used with the above difference in meaning:

```
v2 = v1 [ / % ] m
m3 = m1 [ / % ] m2
v [ / % ] = m
m2 [ / % ] = m1
m2 = x [ / % ] m1
```

If you feel uncomfortable using the `/` and `%` symbols, you can also explicitly write things like

```
v2 = m.inverse() * v1
v3 = v1 * m.inverse()
```

which delay the calculation in exactly the same way that the above forms do. These forms do not ever explicitly calculate the matrix inverse, since this is not (numerically) a good way to perform these calculations. Instead, the appropriate decomposition (see §3.7.3) is used to calculate `v2` and `v3`.

3.7.2 Least-square solutions

If A is not square, then the equation $Ax = b$ does not have a unique solution. If A has more rows than columns, then there is in general no solution. And if A has more columns than rows, then there are an infinite number of solutions.

The former case is more common and represents an overdetermined system of equations. In this case, one is not looking for an exact solution for x , but rather the value of x that minimizes $\|b - Ax\|_2$. This is the meaning of the least-square solution, and is the value returned by $x = b/A$ for the TMV classes.

The matrix A is called the “design” matrix. For example, assume you are doing a simple quadratic fit to some data (x_i, y_i) , and the model you are fitting for can be written as $y = c + dx + ex^2$. Furthermore, assume that each y_i value has a measurement error of s_i . Then the rows of A should be set to: $(1/s_i \ x_i/s_i \ x_i^2/s_i)$. The corresponding element of the vector b should be (y_i/s_i) . It is easily verified that $\|b - Ax\|_2^2$ is the normal χ^2 expression. The solution returned by $x = b/A$ would then be the least-square fit solution: $(c \ d \ e)$, which would be the solution which minimizes χ^2 .

The underdetermined case is not so common, but can still be defined reasonably. As mentioned above, there are infinitely many solutions to such an equation, so the value returned by $x = b/A$ in this case is the value of x that satisfies the equation and has minimum 2-norm, $\|x\|_2$.

When you have calculated a least-square solution for x , it is common to want to know the covariance matrix of the returned values. It turns out that this matrix is $(A^\dagger A)^{-1}$. It is not very efficient to calculate this matrix explicitly and then invert it. But once you have calculated the decomposition needed for the division, it is quite easy. So we provide the routine

```
A.makeInverseATA(Matrix<T>& cov)
```

to perform the calculation efficiently. (Make sure you save the decomposition with `A.saveDiv()` - see §3.7.5 for more about this.)

3.7.3 Decompositions

There are quite a few ways to go about solving the equations written above. The more efficient ways involve decomposing A into a product of matrices with special structures or properties. You can select which decomposition to use with the method:

```
m.divideUsing(tmv::DivType dt)
```

where `dt` can be any of `{tmv::LU, tmv::QR, tmv::QRP, tmv::SV}`. If you do not specify which decomposition to use, LU is the default for square matrices, and QR is the default for non-square matrices.

1. **LU Decomposition:** (`dt = tmv::LU`) $A = PLU$, where L is a lower-triangle matrix with all 1's along the diagonal, U is an upper-triangle matrix, and P is a permutation matrix. This decomposition is only available for square matrices.
2. **QR Decomposition:** (`dt = tmv::QR`) $A = QR$ where Q is a unitary matrix and R is an upper-triangle matrix. (Note: real unitary matrices are also known as orthogonal matrices.) A unitary matrix is such that $Q^\dagger Q = I$.
If A has dimensions $M \times N$ with $M > N$, then R has dimensions $N \times N$, and Q has dimensions $M \times N$. In this case, Q will only be column-unitary. That is $QQ^\dagger \neq I$.
If $M < N$, then A^T is actually decomposed into QR .
3. **QRP Decomposition:** (`dt = tmv::QRP`) $A = QRP$ where Q is unitary, R is upper-triangle, and P is a permutation. This decomposition is somewhat slower than a simple QR decomposition, but it is numerically more stable if A is singular, or nearly so. (Singular matrices are discussed in more detail in §3.7.8.)

There are two slightly different algorithms for doing a QRP decomposition, controlled by a global `bool` variable: `tmv::QRPDiv<T>::StrictQRP`. If this is set to true, then the decomposition will make the diagonal elements of R be strictly decreasing (in absolute value) from upper-left to lower-right.

If it is false, however (the default), then there will be no diagonal element of R below and to the right of one which is much smaller in absolute value, where “much” means the ratio will be at most $\epsilon^{1/4}$, where ϵ is the machine precision for the type in question. This restriction is almost always sufficient to make the decomposition useful for singular or nearly singular matrices, and it is much faster than the strict algorithm.

4. **Singular Value Decomposition:** (dt = `tmv::SV`) $A = USV$ where U is unitary (or column-unitary if $M > N$), S is diagonal and real, and V is unitary (or row-unitary if $M < N$). The values of S will be such that all the values will be non-negative and will decrease along the diagonal. The singular value decomposition is most useful for matrices that are singular or nearly so. We will discuss this decomposition in more detail in §3.7.8 on singular matrices below.

3.7.4 Pseudo-inverse

If m is not square, then `m.inverse()` should return what is called the pseudo-inverse. If m has more rows than columns, then `m.inverse() * m` is the identity matrix, but `m * m.inverse()` is not an identity. If m has more columns than rows, then the opposite holds.

Here are some features of the pseudo-inverse (we use X to represent the pseudo-inverse of M):

$$MXM = M$$

$$XMX = X$$

$$(MX)^T = MX$$

$$(XM)^T = XM$$

For singular square matrices, one can also define a pseudo-inverse with the same properties.

In the first sentence of this section, I used the word “should”. This is because the different decompositions calculate the pseudo-inverse differently and result in slightly different answers. For QR or QRP, the matrix returned by `m.inverse()` for non-square matrices isn’t quite correct. When M is not square, but is also not singular, then X will satisfy the first three of the above equations, but not the last one. With the SV decomposition, however, X is the true pseudo-inverse and all four equations are satisfied.

For singular matrices, QR will fail to give a good pseudo-inverse (and may throw the `tmv::Singular` exception - c.f. §10), QRP will be close to correct (again failing only the last equation) and will not throw an exception, and SV will be correct.

3.7.5 Efficiency issues

Let’s say you compute a matrix decomposition for a particular division calculation, and then later want to use it again for a different right hand side:

```
x = b / m;
[ . . . ]
y = c / m;
```

Ideally, the code would just use the same decomposition that had already been calculated for the `x` assignment when it gets to the later assignment of `y`, so this second division would be very fast. However, what if somewhere in the `[. . .]`, the matrix `m` is modified? Then using the same decomposition would be incorrect - a new decomposition would be needed for the new matrix values.

One solution might be to try to keep track of when a `Matrix` gets changed. We could set an internal flag whenever it is changed to indicate that any existing decomposition is invalid. While not impossible, this type of

check is made quite difficult by the way we have allowed different view objects to point to the same data. It would be difficult, and probably very slow, to make sure that any change in one view invalidates the decompositions of all other views to the same data.

Our solution is instead to err on the side of correctness over efficiency and to always recalculate the decomposition by default. Of course, this can be quite inefficient, so we allow the programmer to override this behavior for a specific `Matrix` object with the method:

```
m.saveDiv()
```

After this call, whenever a decomposition is set, it is saved for any future uses. You are telling the program to assume that the values of `m` will not change after that point (technically after the next decomposition is calculated).

If you do modify `m` after a call to `saveDiv()`, you can manually reset the decomposition with

```
m.resetDiv()
```

which deletes any current saved decomposition, and recalculates it. Similarly,

```
m.unsetDiv()
```

will delete any current saved decomposition, but not calculate a new one. This can be used to free up the memory that the decomposition had been using.

Sometimes you may want to set a decomposition before you actually need it. For example, the division may be in a speed critical part of the code, but you have access to the `Matrix` well before then. You can tell the object to calculate the decomposition with

```
m.setDiv()
```

This may also be useful if you just want to perform and access the decomposition separate from any actual division statement (e.g. SVD for principal component analysis). You can also determine whether the decomposition has been set yet with:

```
bool m.divIsSet()
```

Also, if you change what kind of decomposition the `Matrix` should use by calling `divideUsing(...)`, then this will also delete any existing decomposition that might be saved (unless you “change” it to the same thing).

Finally, there is another efficiency issue, which can sometimes be important. The default behavior is to use extra memory for calculating the decomposition, so the original matrix is left unchanged. However, it is often the case that once you have calculated the decomposition, you don’t need the original matrix anymore. In that case, it is ok to overwrite the original matrix. For very large matrices, the savings in memory may be significant. (The $O(N^2)$ steps in copying the `Matrix` is generally negligible compared to the $O(N^3)$ steps in performing the decomposition. So memory issues are probably the only reason to do this.)

Therefore, we provide another routine that lets the decomposition be calculated in place, overwriting the original `Matrix`⁴:

```
m.divideInPlace()
```

3.7.6 Determinants

Aside from very small matrices, the calculation of the determinant typically requires calculations similar to those performed in the above decompositions. Since a determinant only really makes sense for square matrices, one would typically perform an LU decomposition to calculate the determinant. Then the determinant of A is just the determinant of U (possibly times -1 depending on the details of P), which in turn is simply the product of the values along the diagonal.

⁴ Note: for a regular `Matrix`, this is always possible. However, for some of the special matrix varieties, there are decompositions which cannot be done in place. Whenever that is the case, this directive will be ignored.

Therefore, calling `m.det()` involves calculating the LU decomposition, and then finding the determinant of U . If you are also performing a division calculation, you should probably use `m.saveDiv()` to avoid calculating the decomposition twice.

If you have set `m` to use some other decomposition using `m.divideUsing(...)`, then the determinant will be determined from that decomposition instead (which is always similarly easy).

For large matrices, the value of the determinant can be extremely large, which can easily lead to overflow problems, even for only moderately large matrices. Therefore, we also provide the method `m.logDet()` which calculates the natural logarithm of the absolute value of the determinant. This method can be given an argument, `sign`, which returns the sign of the determinant. The actual determinant can then be reconstructed from

```
T sign;
T logdet = m.logDet(&sign);
T det = sign * exp(logdet);
```

If `m` is a complex matrix, then the “sign” is really a complex number whose absolute value is 1. It is defined to be the value of `(det/abs(det))`.

This alternate calculation is especially useful for non-linear maximum likelihood calculations. Since `log` is a monotonic function, the maximum likelihood is coincident with the maximum of its logarithm. And, since likelihoods in matrix formulation often involve a determinant, the `logDet()` function is exactly what is needed to calculate the log likelihood.

3.7.7 Accessing the decompositions

Sometimes, you may want to access the components of the decomposition directly, rather than just use them for performing the division or calculating the determinant.

- For the LU decomposition, we have:

```
tmv::ConstLowerTriMatrixView<T> m.lud().getL()
tmv::ConstUpperTriMatrixView<T> m.lud().getU()
Permutation m.lud().getP()
bool m.lud().IsTrans()
```

`getL()` and `getU()` return L and U . `getP()` returns the permutation, P . Finally, `IsTrans()` returns whether the product PLU is equal to `m` or `m.transpose()`.

The following should result in a `Matrix m2`, which is numerically very close to the original `Matrix m`:

```
tmv::Matrix<T> m2 = m.lud().getP() * m.lud().getL() * m.lud().getU();
if (m.lud().IsTrans()) m2.transposeSelf();
```

- For the QR decomposition, we have:

```
tmv::PackedQ<T> m.qrd().getQ();
tmv::ConstUpperTriMatrix<T> m.qrd().getR();
bool m.qrd().IsTrans();
```

`getQ()` and `getR()` return Q and R . `IsTrans()` returns whether the decomposition is equal to `m` or `m.transpose()`.

Note: the `PackedQ` class is convertible into a regular `Matrix`, but if you are doing arithmetic with it, then these can generally be done without any conversion, so it is more efficient.

The following should result in a `Matrix m2`, which is numerically very close to the original `Matrix m`:

```
tmv::Matrix<T> m2(m.nrows(),m.ncols());
tmv::MatrixView<T> m2v =
    m.qrd().IsTrans() ? m2.transpose() : m2.view();
m2v = m.qrd().getQ() * m.qrd().getR();
```

- For the QRP decomposition, we have:

```
tmv::PackedQ<T> m.qrpd().getQ()
tmv::ConstUpperTriMatrixView<T> m.qrpd().getR()
Permutation m.qrpd().getP()
bool m.qrpd().IsTrans()
```

`getQ()` and `getR()` return Q and R . `getP()` returns the permutation, P . `IsTrans()` returns whether the decomposition is equal to m or $m.transpose()$.

The following should result in a `Matrix m2`, which is numerically very close to the original `Matrix m`:

```
tmv::Matrix<T> m2(m.nrows(),m.ncols());
tmv::MatrixView<T> m2v =
    m.qrpd().IsTrans() ? m2.transpose() : m2.view();
m2v = m.qrpd().getQ() * m.qrpd().getR() * m.qrpd().getP();
```

- For the SV decomposition, we have:

```
tmv::ConstMatrixView<T> m.svd().getU()
tmv::ConstDiagMatrixView<RT> m.svd().getS()
tmv::ConstMatrixView<T> m.svd().getV()
```

`getU()`, `getS()`, and `getV()` return U , S , and V .

The following should result in a `Matrix m2` that is numerically very close to the original `Matrix m`:

```
tmv::Matrix<T> m2 = m.svd().getU() * m.svd().getS() * m.svd().getV();
```

Each of the above access methods may also be used to perform the decomposition. You do not need to have used the matrix to perform a division before calling `m.lud()`, `m.qrd()`, `m.qrpd()`, or `m.svd()`. Also, calling `m.lud()` implicitly calls `m.divideUsing(tmv::LU)` (likewise for the other decompositions). Thus, if you use one of these decomposition methods, any subsequent division statement will use that decomposition, unless you explicitly call `divideUsing` or call a different decomposition method.

Note that a matrix can only store one decomposition at a time. So if you call `m.lud()` then `m.qrd()`, the LU decomposition is deleted and the QR decomposition is calculated. If you call `m.lud()` again, then it will have to be recalculated.

3.7.8 Singular matrices

If a matrix is singular (i.e. its determinant is 0), then LU and QR decompositions will fail when you attempt to divide by the matrix, since the calculation involves division by 0. When this happens, the TMV code throws an exception, `tmv::Singular`. Furthermore, with numerical rounding errors, a matrix that is close to singular may also end up with 0 in a location that gets used for division and thus throw `tmv::Singular`.

Or, more commonly, a singular or nearly singular matrix will just have numerically very small values rather than actual 0's. In this case, there won't be an error, but the results will be numerically very unstable, since the calculation will involve dividing by numbers that are comparable to the machine precision, ϵ .

You can check whether a `Matrix` is (exactly) singular with the method

```
bool m.isSingular()
```

which basically just returns whether `m.det() == 0`. But this will not tell you if a matrix is merely close to singular, so it does not guard against unreliable results.

Singular value decompositions provides a way to deal with singular and nearly singular matrices. There are a number of methods for the object returned by `m.svd()`, which can help diagnose and fix potential problems with a singular matrix.

First, the so-called “singular values”, which are the elements of the diagonal S matrix, tell you how close the matrix is to being singular. Specifically, if the ratio of the smallest and the largest singular values, $S(N-1)/S(0)$, is close to the machine precision, ϵ , for the underlying type (`double`, `float`, etc.), then the matrix is singular, or effectively so. Note: the value of ϵ is accessible with:

```
std::numeric_limits<T>::epsilon()
```

The inverse of this ratio, $S(0)/S(N-1)$, is known as the “condition” of the matrix (specifically the 2-condition, or κ_2), which can be obtained by:

```
m.svd().condition()
```

The larger the condition, the closer the matrix is to singular, and the less reliable any calculation would be.

So, how does SVD help in this situation? (So far we have diagnosed the possible problem, but not fixed it.)

Well, we need to figure out what solution we want from a singular matrix. If the matrix is singular, then there are not necessarily any solutions to $Ax = b$. Furthermore, if we are looking for a least squares solution (rather than an exact solution) then there are an infinite number of choices for x that give the same minimum value of $\|Ax - b\|_2$.

Another way of looking at it is that there will be particular values of y for which $Ay = 0$. Then given a solution x , the vector $x' = x + \alpha y$ for any α will produce the same solution: $Ax' = Ax = b$.

The usual desired solution is the x with minimum 2-norm, $\|x\|_2$. With SVD, we can get this solution by setting to 0 all of the values in S^{-1} that would otherwise be infinity (or at least large compared to $1/\epsilon$). It is somewhat ironic that the best way to deal with an infinite value is to set it to 0, but that is actually the solution we want.

There are two methods that can be used to control which singular values are set to 0⁵:

```
m.svd().thresh(RT thresh)
m.svd().top(int nsing)
```

`thresh` sets to 0 any singular values with $S(i)/S(0) < \text{thresh}$. `top` uses only the largest `nsing` singular values, and sets the rest to 0.

The default behavior is equivalent to:

```
m.svd().thresh(std::numeric_limits<T>::epsilon());
```

since at least these values are unreliable. For different applications, you may want to use a larger threshold value.

You can check how many singular values are currently considered non-zero with

```
int m.svd().getKMax()
```

A QRP decomposition can deal with singular matrices similarly, but it doesn’t have the flexibility in checking for not-quite-singular but somewhat ill-conditioned matrices like SVD does. QRP will put all of the small elements of R ’s diagonal in the lower right corner. Then it ignores any that are less than ϵ when doing the division. For actually singular matrices, this should produce essentially the same result as the SVD solution.

We should also mention again that the 2-norm of a matrix is the largest singular value, which is just $S(0)$ in our decomposition. So this norm requires an SVD calculation, which is relatively expensive compared to the other norms if you have not already calculated the singular value decomposition (and saved it with `m.saveDiv()`). On the other hand, if you have already computed the SVD for division, then `norm2` is trivial and is the fastest norm to compute.

⁵Technically, the actual values are preserved, but an internal value, `kmax`, keeps track of how many singular values to use.

If you just want to calculate the singular values, but don't need to do the actual division, then you don't need to accumulate the U and V matrices. This saves a lot of the calculation time. Or you might want U or V , but not both for some purpose. See section §11.2 for how to do this.

3.8 Input/Output

The simplest output is the usual:

```
os << m
```

where `os` is any `std::ostream`. The output format is:

```
nrows ncols
( m(0,0) m(0,1) m(0,2) ... m(0,ncols-1) )
( m(1,0) m(1,1) m(1,2) ... m(1,ncols-1) )
...
( m(nrows-1,0) ... ... m(nrows-1,ncols-1) )
```

The same format can be read back in one of two ways:

```
tmv::Matrix<T> m(nrows,ncols);
is >> m;
std::auto_ptr<tmv::Matrix<T> > mptr;
is >> mptr;
```

For the first version, the `Matrix` must already be declared, which requires knowing how big it needs to be. If the input `Matrix` does not match in size, a runtime error will occur. The second version allows you to get the size from the input. `mptr` will be created (with `new`) according to whatever size the input `Matrix` is.

Often, it is convenient to output only those values that aren't very small. This can be done using

```
m.write(std::ostream& os, RT thresh)
```

which is equivalent to

```
os << tmv::Matrix<T>(m).clip(thresh);
```

but without requiring the temporary `Matrix`.

3.9 Small matrices

The algorithms for regular `Matrix` operations are optimized to be fast for large matrices. Usually, this makes sense, since any code with both large and small matrices will probably have its performance dominated by the speed of the large matrix algorithms.

However, it may be the case that a particular program spends all of its time using 2×2 or 3×3 matrices. In this case, many of the features of the TMV code are undesirable. For example, the alias checking in the operation `v2 = m * v1` becomes a significant fraction of the operating time. Even the simple act of performing a function call, rather than doing the calculation inline may be a big performance hit.

So we include the alternate matrix class called `SmallMatrix`, along with the corresponding vector class called `SmallVector`, for which most of the operations are done inline. Furthermore, the sizes are template arguments, rather than normal parameters. This allows the compiler to easily optimize simple calculations that might only be 2 or 3 arithmetic operations, which may significantly speed up your code.

All the `SmallMatrix` and `SmallVector` routines are included by:

```
#include "TMV_Small.h"
```

These classes do not inherit from the regular versions, but they do have essentially all the same methods, functions, and arithmetic operators.

3.9.1 Constructors

The template arguments `M` and `N` below are both integers and represent the size of the matrix or vector as indicated. The template argument `stor` may be either `tmv::RowMajor` or `tmv::ColMajor`, and `index` may be either `tmv::CStyle` or `tmv::FortranStyle`. These both have the same meanings as they do for a regular `Matrix`. The defaults are `ColMajor` and `CStyle` if they are omitted.

- `tmv::SmallVector<T,N,index> v()`

Makes a `SmallVector` of size `N` with uninitialized values. If debugging is turned on (i.e. not turned off with `-DNDEBUG`), then the values are in fact initialized to 888.

- `tmv::SmallVector<T,N,index> v(T x)`

Makes a `SmallVector` with all values equal to `x`.

- `tmv::SmallVector<T,N,index> v(const T* vv)`
`tmv::SmallVector<T,N,index> v(const std::vector<T>& vv)`

Makes a `SmallVector` with values copied from `vv`.

- `tmv::SmallVector<T,N,stor,index> v1(const GenVector<T>& v2)`

Makes a `SmallVector` from a regular `Vector`.

- `tmv::SmallVector<T,N,stor,index> v1(
 const SmallVector<T2,N,stor2,index2>& v2)
v1 = v2`

Copy the `SmallVector` `v2`, which may be of any type `T2` so long as values of type `T2` are convertible into type `T`. The assignment operator has the same flexibility.

- `v << v0 , v1 , v2 , v3 ...`
`v << v0 << v1 << v2 << v3 ...`

Initialize the `SmallVector` `v` with a list of values.

- `tmv::SmallMatrix<T,M,N,stor,index> m()`

Makes an $M \times N$ `SmallMatrix` with uninitialized values. If debugging is turned on (i.e. not turned off with `-DNDEBUG`), then the values are in fact initialized to 888.

- `tmv::SmallMatrix<T,M,N,stor,index> m(T x)`

Makes an $M \times N$ `SmallMatrix` with all values equal to `x`.

- `tmv::SmallMatrix<T,M,N,stor,index> m(const T* vv)`
`tmv::SmallMatrix<T,M,N,stor,index> m(const std::vector<T>& vv)`
`tmv::SmallMatrix<T,M,N,stor,index> m(
 const std::vector<std::vector<T>> & vv)`

Makes an $M \times N$ `SmallMatrix` with values copied from `vv`.

- `tmv::SmallMatrix<T,M,N,stor,index> m1(const GenMatrix<T>& m2)`

Makes an $M \times N$ `SmallMatrix` from a regular `Matrix`.

- `tmv::SmallMatrix<T,M,N,stor,index> m1(
 const SmallMatrix<T2,M,N,stor2,index2>& m2)
m1 = m2`

Copy the `SmallMatrix` `m2`, which may be of any type `T2` so long as values of type `T2` are convertible into type `T`. The assignment operator has the same flexibility.

- `tmv::SmallMatrix<T,M,N,tmv::RowMajor,index> m;
m << m00 , m01 , m02 ...
 m10 , m11 , m12 ...
 ...`

Initialize the `SmallMatrix` `m` with a list of values. The elements are assigned in order according to the storage order of `m`. This is most intuitive if `stor = RowMajor` as shown.

3.9.2 Access

The basic access methods are the same as for a regular `Vector` or `Matrix`. (See 2.3, 3.3.) However, since the size is known to the compiler, the inline calculation is able to be a lot faster, often reducing to a trivial memory access.

The various view methods, like `row`, `col`, `transpose`, `conjugate`, etc. do not return a `SmallVector` or `SmallMatrix`, so operations with the returned views will not be done inline. However, you can copy the view back to a “Small” object, which will be done inline, so that should be fast.

Also, views may be combined with `Small` objects in arithmetic statements. In fact, simple views – those where all the memory elements are contiguous in memory – will usually have the arithmetic calculation done inline. In this case, a good optimizing compiler will probably produce code that eliminates the temporary non-`Small` view. For example:

```
SmallMatrix<T,M,N> m3 = m1.transpose() * m2;  
SmallVector<T,N> v2 = m1 * m3.col(2);
```

will both be done inline, even though `m1.transpose()` is not a `SmallMatrix` and `m3.col(2)` is not a `SmallVector`.

3.9.3 Functions

`SmallVectors` and `SmallMatrixes` all have exactly the same function methods as the regular varieties. (See 2.5, 3.5.) Likewise, the syntax of the arithmetic is identical. There are only a few methods that are not done inline.

First, reading a `SmallMatrix` or `SmallVector` from a file uses the regular `Matrix` I/O methods. Also, there is no `auto_ptr` version of these read operations, since you need to know the size of a `SmallMatrix` or `SmallVector` at compile time anyway, so there is no way to wait until the file is read to determine the size.

Also, the `sort` command for a `SmallVector` just uses the regular `Vector` version.

3.9.4 Limitations

There are a few limitations on `SmallMatrix` and `SmallVector` objects that we impose in order to give the compiler the ability to optimize them as much as possible.

1. No alias checking

Statements such as

```
v = m * v;
```


will not produce correct code if `v` and `m` are a `SmallVector` and `SmallMatrix`.

For a regular `Vector` and `Matrix`, the TMV library checks whether any of the objects in an arithmetic statement use the same memory. Then it uses the correct algorithm to deal with it correctly.

This alias checking is fairly expensive for small matrices, so we don't do it if `m` and `v` are a `SmallMatrix` and a `SmallVector`.

If you need to perform this operation you can write:

```
v = m * tmv::SmallVector<N>(v);
```

to create a temporary copy of the vector on the right hand side.

2. No virtual functions

The `vtable` required for `SmallMatrix` to use virtual functions is potentially a detriment to performance, so we don't have any virtual functions in either `SmallVector` or `SmallMatrix`.

A consequence of this decision is that the `SmallMatrix` class is not in the usual class hierarchy. It does not inherit from either `GenMatrix` or `BaseMatrix` or even `AssignableToMatrix`, since they have virtual functions. Likewise, `SmallVector` does not inherit from `GenVector` or `AssignableToVector`.

If you want to view a `SmallMatrix` as a `GenMatrix`, you can type

```
ConstMatrixView<T> mv = m.view();
```

and use `mv` wherever you need a `GenMatrix`. Of course, you will lose the inlining advantages of a `SmallMatrix`, but sometimes that is what you need.

Similarly, a non-`const` `SmallMatrix` can be viewed as a `MatrixView` with the same method.

3. Views are not “Small”

All the operations that return some kind of view, either a `VectorView` or a `MatrixView`, are not “Small”. As described above, this means that they don't have all the inlining advantages of `SmallVector` and `SmallMatrix`.

However, you can copy them back to a `Small` object or combine with `Small` objects in an arithmetic statement, which will be done inline in some cases.

4. Division method is fixed

A `SmallMatrix` does not have the various division control methods like `divideUsing`, `saveDiv`, etc. So a square `SmallMatrix` will always use LU decomposition, and a non-square one will always use QR decomposition. And if you are doing multiple division statements with the same matrix, the library will not save the decomposition between statements.

4 Diagonal matrices

The `DiagMatrix` class is our diagonal matrix class. A diagonal matrix is only non-zero along the main diagonal of the matrix.

The class `DiagMatrix` inherits from `GenDiagMatrix`, which in turn inherits from `BaseMatrix`. The various views and composite classes described below also inherit from `GenDiagMatrix`.

Most functions and methods for `DiagMatrix` work the same as they do for `Matrix`. In these cases, we will just list the functions that are allowed with the effect understood to be the same as for a regular `Matrix`. Of course, there are almost always algorithmic speed-ups, which the code will use to take advantage of the diagonal structure. Whenever there is a difference in how a function works, we will explain the difference.

4.1 Constructors

As usual, the optional `index` template argument specifies which indexing style to use.

- `tmv::DiagMatrix<T,index> d(size_t n)`

Makes an $n \times n$ `DiagMatrix` with uninitialized values. If debugging is turned on (i.e. not turned off with `-DNDEBUG`), then the values along the diagonal are in fact initialized to 888.

- `tmv::DiagMatrix<T,index> d(size_t n, T x)`

Makes an $n \times n$ `DiagMatrix` with all values along the diagonal equal to `x`.

- `tmv::DiagMatrix<T,index> m(size_t n, const T* vv)`
`tmv::DiagMatrix<T,index> m(const std::vector<T>& vv)`

Makes a `DiagMatrix` which copies the elements of `vv`.

- `tmv::DiagMatrix<T,index> d(const GenVector<T>& v)`

Makes a `DiagMatrix` with `v` as the diagonal.

- `tmv::DiagMatrix<T,index> d(const GenMatrix<T>& m)`

Makes a `DiagMatrix` with the diagonal of `m` as the diagonal.

- `tmv::DiagMatrix<T,index> d1(const GenDiagMatrix<T2>& d2)`
`d1 = d2`

Copy the `DiagMatrix` `d2`, which may be of any type `T2` so long as values of type `T2` are convertible into type `T`. The assignment operator has the same flexibility.

- `tmv::DiagMatrixView<T,index> d =`
 `DiagMatrixViewOf(VectorView<T,index> v)`
`tmv::ConstDiagMatrixView<T,index> d =`
 `DiagMatrixViewOf(ConstVectorView<T,index> v)`

Makes a `DiagMatrixView` whose diagonal is `v`.

- `tmv::DiagMatrixView<T,index> d =`
 `tmv::DiagMatrixViewOf(T* vv, size_t n)`
`tmv::ConstDiagMatrixView<T,index> d =`
 `tmv::DiagMatrixViewOf(const T* vv, size_t n)`

Make a `DiagMatrixView` whose diagonal consists of the actual memory elements `vv`.

4.2 Access

```
d.nrows() = d.ncols() = d.colsize() = d.rowsize() = d.size()
d(i, j)
d(i) = d(i, i)
```

For the mutable `d(i, j)` version, `i` must equal `j`. If `d` is not mutable, then `d(i, j)` with $i \neq j$ returns the value 0.

```
d.diag()
```

```
d.subDiagMatrix(int i1, int i2, int istep = 1)
```

This is equivalent to `DiagMatrixViewOf(d.diag().subVector(i1, i2, istep))`.

```
d.transpose() = d.view()
d.conjugate() = d.adjoint()
d.cview()
d.fview()
d.realPart()
d.imagPart()
```

4.3 Functions

```
RT d.norm1() = Norm1(d)
RT d.norm2() = Norm2(d) = d.doNorm2()
RT d.normInf() = NormInf(d)
RT d.maxAbsElement() = MaxAbsElement(d)
```

(Actually for a diagonal matrix, all of the above norms are equal.)

```
RT d.maxAbs2Element() = MaxAbs2Element(d)
RT d.normF() = NormF(d) = d.norm() = Norm(d)
RT d.normSq() = NormSq(d)
RT d.normSq(RT scale)
T d.trace() = Trace(d)
T d.sumElements() = SumElements(d)
RT d.sumAbsElements() = SumAbsElements(d)
T d.det() = Det(d)
RT d.logDet(T* sign=0) = LogDet(d)
bool d.isSingular()
RT d.condition()
RT d.doCondition()
dinv = d.inverse() = Inverse(d)
d.makeInverse(Matrix<T>& minv)
d.makeInverse(DiagMatrix<T>& dinv)
d.makeInverseATA(Matrix<T>& cov)
d.makeInverseATA(DiagMatrix<T>& cov)
```

Since the inverse of a `DiagMatrix` is a `DiagMatrix`, we also provide a version of the `makeInverse` syntax, which allows `dinv` to be a `DiagMatrix`. (Likewise for `makeInverseATA`.) The same option is available with the operator version: `dinv = d.inverse()`.

```
d.setZero()
d.setAllTo(T x)
```

```

d.clip(RT thresh)
d.setToIdentity(T x = 1)
d.conjugateSelf()
d.transposeSelf() // null operation
d.invertSelf()
Swap(d1, d2)

```

`invertSelf` is new for `DiagMatrix` and calculates d^{-1} in place. It is equivalent to `d = d.inverse()`.

4.4 Arithmetic

In addition to `x`, `v`, and `m` from before, we now add `d` for a `DiagMatrix`.

```

d2 = -d1
d2 = x * d1
d2 = d1 [*/] x
d3 = d1 [+ -] d2
m2 = m1 [+ -] d
m2 = d [+ -] m1
d [*/] = x
d2 [+ -] = d1
m [+ -] = d
v2 = d * v1
v2 = v1 * d
v *= d
d3 = d1 * d2
m2 = d * m1
m2 = m1 * d
d2 *= d1
m *= d
d2 = d1 [+ -] x
d2 = x [+ -] d1
d [+ -] = x
d = x

```

4.5 Division

The division operations are:

```

v2 = v1 [/ %] d
m2 = m1 [/ %] d
m2 = d [/ %] m1
d3 = d1 [/ %] d2
d2 = x [/ %] d1
v [/ %] = d
d2 [/ %] = d1
m [/ %] = d

```

There is only one allowed `DivType` for a `DiagMatrix`: `LU`. And, since it is also the default behavior, there is no reason to ever use this function. Furthermore, since a `DiagMatrix` is already a U matrix, the decomposition requires no work at all. Hence, it is always done in place; no extra storage is needed, and the methods `m.divideInPlace()`, `m.saveDiv()`, etc. are irrelevant.

If a `DiagMatrix` is singular, you can find out with `m.isSingular()`, but there is no direct way to use SVD for the division and skip any divisions by 0. If you want to do this, you should use `BandMatrixViewOf(d)` to treat the `DiagMatrix` as a `BandMatrix`, which can use SVD.

4.6 Input/Output

The simplest output is the usual:

```
os << d
```

where `os` is any `std::ostream`. The output format is the same as for a `Matrix`, including all the 0's. (See 3.8.)

There is also a compact format:

```
d.writeCompact(os)
```

which outputs in the format:

```
D n ( d(0,0)  d(1,1)  d(2,2)  ...  d(n-1,n-1) )
```

The same (compact, that is) format can be read back in the usual two ways:

```
tmv::DiagMatrix<T> d(n);  
is >> d;  
std::auto_ptr<tmv::DiagMatrix<T> > dptr;  
is >> dptr;
```

where the first gives an error if `d` is the wrong size and the second allocates a new `DiagMatrix` that is the correct size.

One can also write small values as 0 with

```
m.write(std::ostream& os, RT thresh)  
m.writeCompact(std::ostream& os, RT thresh)
```

5 Upper/lower triangle matrices

The `UpperTriMatrix` class is our upper triangle matrix class, which is non-zero only on the main diagonal and above. `LowerTriMatrix` is our class for lower triangle matrices, which are non-zero only on the main diagonal and below.

The class `UpperTriMatrix` inherits from `GenUpperTriMatrix`, and the class `LowerTriMatrix` inherits from `GenLowerTriMatrix`, both of which in turn inherit from `BaseMatrix`. The various views and composite classes described below also inherit from `GenUpperTriMatrix` and `GenLowerTriMatrix` as appropriate.

In addition to the `T` template parameter, there are three other template parameters: `dt`, which can be either `tmv::UnitDiag` or `tmv::NonUnitDiag`; `stor`, which can be `tmv::RowMajor` or `tmv::ColMajor`; and `index`, which can be `tmv::CStyle` or `tmv::FortranStyle`. The default values for these template parameters are `NonUnitDiag`, `ColMajor`, and `CStyle` respectively.

The storage of both an `UpperTriMatrix` and a `LowerTriMatrix` takes $N \times N$ elements of memory, even though approximately half of them are never used. Someday, I'll write the packed storage versions, which allow for efficient storage of the matrices.

All of the routines are analogous for `UpperTriMatrix` and `LowerTriMatrix`, so we only list each routine once (the `UpperTriMatrix` version for definiteness).

Most functions and methods for `UpperTriMatrix` and `LowerTriMatrix` work the same as they do for `Matrix`. In these cases, we will just list the functions that are allowed with the effect understood to be the same as for a regular `Matrix`. Of course, there are almost always algorithmic speed-ups, which the code will use to take advantage of the triangle structure. Whenever there is a difference in how a function works, we will explain the difference.

5.1 Constructors

- `tmv::UpperTriMatrix<T,dt,stor,index> U(size_t n)`

Makes an $n \times n$ `UpperTriMatrix` with uninitialized values. If debugging is turned on (i.e. not turned off with `-DNDEBUG`), then the values are in fact initialized to 888.

- `tmv::UpperTriMatrix<T,dt,stor,index> U(size_t n, T x)`

Makes an $n \times n$ `UpperTriMatrix` with all values equal to `x`.

- `tmv::UpperTriMatrix<T,dt,stor,index> U(size_t n, const T* vv)`
`tmv::UpperTriMatrix<T,dt,stor,index> U(size_t n,`
`const std::vector<T>& vv)`

Makes an `UpperTriMatrix` which copies the elements of `vv`. Note: the elements in `vv` need to include the zero's of the lower triangle portion of `U`.

- `tmv::UpperTriMatrix<T,dt,stor,index> U(const GenMatrix<T>& m)`
`tmv::UpperTriMatrix<T,dt,stor,index> U(const GenUpperTriMatrix<T>& U2)`

Make an `UpperTriMatrix` which copies the corresponding values of `U2`. Note that the second one is allowed to have `U2` be `NonUnitDiag` but `dt = UnitDiag`, in which case only the off-diagonal elements are copied. The converse would set the diagonal of the new `UpperTriMatrix` to all 1's.

- `tmv::UpperTriMatrix<T> U1(const GenUpperTriMatrix<T2>& U2)`
`U1 = U2`

Copy the `UpperTriMatrix U2`, which may be of any type `T2` so long as values of type `T2` are convertible into type `T`. The assignment operator has the same flexibility.

- `tmv::UpperTriMatrixView<T,index> U =`
`tmv::UpperTriMatrixViewOf(T* vv, size_t n,`
`StorageType stor, DiagType dt=NonUnitDiag)`
`tmv::ConstUpperTriMatrixView<T,index> U =`
`tmv::UpperTriMatrixViewOf(const T* vv, size_t n,`
`StorageType stor, DiagType dt=NonUnitDiag)`

Make a `UpperTriMatrixView` of the actual memory elements, `vv`. One wrinkle here is that if `dt` is `UnitDiag`, then `vv` is still the location of the upper left corner, even though that value is never used (since the value is just taken to be 1). Also, `vv` must be of length $n \times n$, so all of the lower triangle elements must be in memory, even though they are never used.

5.2 Access

```
U.nrows() = U.ncols() = U.colsize() = U.rowsize() = U.size()
U(i,j)
U.row(int i, int j1, int j2)
U.col(int i, int j1, int j2)
U.diag()
U.diag(int i)
U.diag(int i, int k1, int k2)
```

For the mutable `d(i,j)` version, Note that the versions of `row` and `col` with only one argument are missing, since the full row or column isn't accessible as a `VectorView`. You must specify a valid range within the row or column that you want, given the upper triangle shape of `U`. Likewise for the `LowerTriMatrix` versions of these. If `dt` is `UnitDiag`, then the range may not include the diagonal element. Similarly, `U.diag()` is valid only if `dt` is `NonUnitDiag`.

```
U.subVector(int i, int j, int istep, int jstep, int size)
U.subMatrix(int i1, int i2, int j1, int j2)
U.subMatrix(int i1, int i2, int j1, int j2, int istep, int jstep)
```

This works the same as for `Matrix` (See 3.3), except that all of the elements in the subvector or submatrix must be completely within the upper or lower triangle, as appropriate. If `dt` is `UnitDiag`, then no elements may be on the main diagonal.

```
U.subTriMatrix(int i1, int i2, int istep = 1)
```

This returns the upper or lower triangle matrix whose upper-left corner is `U(i1,i1)`, and whose lower-right corner is `U(i2-istep,i2-istep)` for C-style indexing or `U(i2,i2)` for Fortran-style indexing. If `istep` $\neq 1$, then it is the step in both the `i` and `j` directions.

```
U.offDiag()
```

This returns a view to the portion of the triangle matrix that does not include the diagonal elements. It will always be `NonUnitDiag`. Internally, it provides an easy way to deal with the `UnitDiag` triangle matrices for many routines. But it may be useful for some users as well.

```
U.viewAsUnitDiag()
```

This returns a view to a `NonUnitDiag` triangle matrix that treats it instead as a `UnitDiag` triangle matrix.

```

U.transpose()
U.conjugate()
U.adjoint()
U.view()
U.cView()
U.fView()
U.realPart()
U.imagPart()

```

Note that the transpose and adjoint of an `UpperTriMatrix` is an `LowerTriMatrixView` and vice versa.

5.3 Functions

```

RT U.norm1() = Norm1(U)
RT U.norm2() = Norm2(U) = U.doNorm2()
RT U.normInf() = NormInf(U)
RT U.normF() = NormF(U) = U.norm() = Norm(U)
RT U.normSq() = NormSq(U)
RT U.normSq(RT scale)
RT U.maxAbsElement() = MaxAbsElement(U)
RT U.maxAbs2Element() = MaxAbs2Element(U)
T U.trace() = Trace(U)
T U.sumElements() = SumElements(U)
RT U.sumAbsElements() = SumAbsElements(U)
T U.det() = Det(U)
RT U.logDet(T* sign=0) = LogDet(U)
bool U.isSingular()
RT U.condition()
RT U.doCondition()
Uinv = U.inverse() = Inverse(U)
U.makeInverse(Matrix<T>& minv)
U.makeInverse(UpperTriMatrix<T>& Uinv)
U.makeInverseATA(Matrix<T>& cov)

```

Since the inverse of an `UpperTriMatrix` is also upper triangular, the object returned by `U.inverse()` is assignable to an `UpperTriMatrix`. Of course you can also assign it to a regular `Matrix` if you prefer. Similarly, there are versions of `U.makeInverse(minv)` for both argument types. Of course, similar statements hold for `LowerTriMatrix` as well.

```

U.setZero()
U.setAllTo(T x)
U.clip(RT thresh)
U.setToIdentity(T x = 1)
U.conjugateSelf()
U.invertSelf()
Swap(U1, U2)

```

Like for `DiagMatrix`, `invertSelf` calculates U^{-1} in place. It is equivalent to `U = U.inverse()`.

5.4 Arithmetic

In addition to `x`, `v`, and `m` from before, we now add `U` and `L` for a `UpperTriMatrix` and `LowerTriMatrix` respectively. Where the syntax is identical for the two cases, only the `U` form is listed.


```

U2 = -U1
U2 = x * U1
U2 = U1 [*/] x
U3 = U1 [+ -] U2
m2 = m1 [+ -] U
m2 = U [+ -] m1
m = L [+ -] U
m = U [+ -] L
U [*/] = x
U2 [+ -] = U1
m [+ -] = U
v2 = U * v1
v2 = v1 * U
v *= U
U3 = U1 * U2
m2 = U * m1
m2 = m1 * U
m = U * L
m = L * U
U2 *= U1
m *= U
U2 = U1 [+ -] x
U2 = x [+ -] U1
U [+ -] = x

```

5.5 Division

The division operations are: (again omitting the L forms when redundant)

```

v2 = v1 [/ %] U
m2 = m1 [/ %] U
m2 = U [/ %] m1
U3 = U1 [/ %] U2
U2 = x [/ %] U1
m = U [/ %] L
m = L [/ %] U
v [/ %] = U
U2 [/ %] = U1
m [/ %] = U

```

There is only one allowed `DivType` for an `UpperTriMatrix` or a `LowerTriMatrix`: `LU`. And, since it is also the default behavior, there is no reason to ever specify it. Furthermore, as with a `DiagMatrix`, the decomposition requires no work at all. In fact, the ease of dividing by a upper or lower triangle matrix is precisely why the LU decomposition is useful. Hence, it is always done in place. i.e. no extra storage is needed, and all of the `m.divideInPlace()`, `m.saveDiv()`, etc. are irrelevant.

If an `UpperTriMatrix` or `LowerTriMatrix` is singular, you can check easily with `m.isSingular()`, but there is no direct way to use SVD for the division and avoid any divisions by 0. If you want to do this use `BandMatrixViewOf(m)` to treat the `TriMatrix` as a `BandMatrix`, which can use SVD.

5.6 Input/Output

The simplest output is the usual:

```
os << U << L
```

where `os` is any `std::ostream`. The output format is the same as for a `Matrix`, including all the 0's. (See 3.8.)

There is also a compact format. For an `UpperTriMatrix`,

```
U.writeCompact(os)
```

outputs in the format:

```
U n
( U(0,0)  U(0,1)  U(0,2)  ...  U(0,n-1) )
( U(1,1)  U(1,2)  ...  U(1,n-1) )
...
( U(n-1,n-1) )
```

For a `LowerTriMatrix`,

```
L.writeCompact(os)
```

outputs in the format:

```
L n
( L(0,0) )
( L(1,0)  L(1,1) )
...
( L(n-1,0)  L(n-1,1)  ...  L(n-1,n-1) )
```

In each case, the compact format can be read back in the usual two ways:

```
tmv::UpperTriMatrix<T> U(n);
tmv::LowerTriMatrix<T> L(n);
is >> U >> L;
std::auto_ptr<tmv::UpperTriMatrix<T> > Uptr;
std::auto_ptr<tmv::LowerTriMatrix<T> > Lptr;
is >> Uptr >> Lptr;
```

One can write small values as 0 with

```
m.write(std::ostream& os, RT thresh)
m.writeCompact(std::ostream& os, RT thresh)
```

6 Band-diagonal matrices

The `BandMatrix` class is our band-diagonal matrix, which is only non-zero on the main diagonal and a few sub- and super-diagonals. While band-diagonal matrices are usually square, we allow for non-square banded matrices as well. You may even have rows or columns that are completely outside of the band structure, and hence are all 0. For example a 10×5 band matrix with 2 sub-diagonals is valid even though the bottom 3 rows are all 0.

Throughout, we use `nlo` to refer to the number of sub-diagonals (below the main diagonal) stored in the `BandMatrix`, and `nhi` to refer to the number of super-diagonals (above the main diagonal).

`BandMatrix` inherits from `GenBandMatrix`, which in turn inherits from `BaseMatrix`. The various views and composite classes described below also inherit from `GenBandMatrix`.

All the `BandMatrix` routines are included by:

```
#include "TMV_Band.h"
```

In addition to the `T` template parameter, we also have `stor` to indicate which storage you want to use, and the usual `index` parameter. For this class, we add an additional storage possibility: along with `RowMajor` and `ColMajor` storage, a `BandMatrix` may also be `DiagMajor`, which has unit step along the diagonals. The `index` parameter has the usual options of `CStyle` or `FortranStyle`. The default values for `stor` and `index` are `ColMajor` and `CStyle`.

For each type of storage, we require that the step size in each direction be uniform within a given row, column or diagonal. This means that we require a few extra elements of memory that are not actually used. To demonstrate the different storage orders and why extra memory is required, here are three 6×6 band-diagonal matrices, each with $nlo = 2$ and $nhi = 3$ in each of the different storage types. The number in each place indicates the offset in memory from the top left element.

$$\begin{array}{lcl}
 \text{ColMajor:} & & \begin{pmatrix} 0 & 5 & 10 & 15 & & \\ 1 & 6 & 11 & 16 & 21 & \\ 2 & 7 & 12 & 17 & 22 & 27 \\ & 8 & 13 & 18 & 23 & 28 \\ & & 14 & 19 & 24 & 29 \\ & & & 20 & 25 & 30 \end{pmatrix} \\
 \text{RowMajor:} & & \begin{pmatrix} 0 & 1 & 2 & 3 & & \\ 5 & 6 & 7 & 8 & 9 & \\ 10 & 11 & 12 & 13 & 14 & 15 \\ & 16 & 17 & 18 & 19 & 20 \\ & & 22 & 23 & 24 & 25 \\ & & & 28 & 29 & 30 \end{pmatrix} \\
 \text{DiagMajor:} & & \begin{pmatrix} 0 & 6 & 12 & 18 & & \\ -5 & 1 & 7 & 13 & 19 & \\ -10 & -4 & 2 & 8 & 14 & 20 \\ & -9 & -3 & 3 & 9 & 15 \\ & & -8 & -2 & 4 & 10 \\ & & & -7 & -1 & 5 \end{pmatrix}
 \end{array}$$

First, notice that all three storage methods require 4 extra locations in memory, which do not hold any actual matrix data. (They require a total of 31 memory addresses for only 27 that are used.) This is because we want to have the same step size between consecutive row elements for every row. Likewise for the columns (which in turn implies that it is also true for the diagonals).

For $N \times N$ square matrices, the total memory needed is $(N - 1) * (nlo + nhi + 1) + 1$, which wastes $(nlo - 1) * nlo/2 + (nhi - 1) * nhi/2$ locations. For non-square matrices, the formula is more complicated,

and changes slightly between the three storages. If you want to know the memory used by a `BandMatrix`, we provide the routine:

```
size_t BandStorageLength(StorageType stor, size_t nrows, size_t ncols,
                        int nlo, int nhi)
```

For the mutable `d(i, j)` version,

For square matrices, all three methods always need the same amount of memory (and for non-square, they aren't very different), so the decision about which method to use should generally be based on performance considerations rather than memory usage. The speed of the various matrix operations are different for the different storage types. If the matrix calculation speed is important, it may be worth trying all three to see which is fastest for the operations you are using.

Second, notice that the `DiagMajor` storage doesn't start with the upper left element as usual. Rather, it starts at the start of the lowest sub-diagonal. So for the constructors that take the matrix information from an array (`T*` or `vector<T>`), the start of the array needs to be at the start of the lowest sub-diagonal.

Most functions and methods for `BandMatrix` work the same as they do for `Matrix`. In these cases, we will just list the functions that are allowed with the effect understood to be the same as for a regular `Matrix`. Of course, there are almost always algorithmic speed-ups, which the code will use to take advantage of the banded structure. Whenever there is a difference in how a function works, we will explain the difference.

6.1 Constructors

- `tmv::BandMatrix<T, stor, index> b(size_t nrows, size_t ncols, int nlo, int nhi)`

Makes a `BandMatrix` with `nrows` rows, `ncols` columns, `nlo` sub-diagonals, and `nhi` super-diagonals with uninitialized values. If debugging is turned on (i.e. not turned off with `-DNDEBUG`), then the values are initialized to 888.

- `tmv::BandMatrix<T, stor, index> b(size_t nrows, size_t ncols, int nlo, int nhi, T x)`

Makes a `BandMatrix` with all values equal to `x`.

- `tmv::BandMatrix<T, stor, index> b(size_t nrows, size_t ncols, int nlo, int nhi, const T* vv)`
`tmv::BandMatrix<T, stor, index> b(size_t nrows, size_t ncols, int nlo, int nhi, const std::vector<T>& vv)`

Makes a `BandMatrix` which copies the elements from `vv`. See the discussion above about the different storage types to see what order these elements should be. The function `BandStorageLength` will tell you how long `vv` must be. The elements that don't fall in the bounds of the actual matrix are not used and may be left undefined.

- `tmv::BandMatrix<T, stor, index> b(const GenMatrix<T>& m, int nlo, int nhi)`
`tmv::BandMatrix<T, stor, index> b(const GenBandMatrix<T>& m, int nlo, int nhi)`
`tmv::BandMatrix<T, stor, index> b(const GenUpperTriMatrix<T>& m, int nhi)`
`tmv::BandMatrix<T, stor, index> b(const GenLowerTriMatrix<T>& m, int nlo)`

Make a `BandMatrix` the same size as `m`, which copies the values of `m` that are within the band defined by `nlo` and `nhi`. For the second one, `nlo` and `nhi` must not be larger than those for `m`. For the last two, `nlo` and `nhi` (respectively) are taken to be 0.

- `tmv::BandMatrix<T,tmv::DiagMajor> m = UpperBiDiagMatrix(`
`const GenVector<T>& v1, const GenVector<T>& v2)`
`tmv::BandMatrix<T,tmv::DiagMajor> m = LowerBiDiagMatrix(`
`const GenVector<T>& v1, const GenVector<T>& v2)`
`tmv::BandMatrix<T,tmv::DiagMajor> m = TriDiagMatrix(`
`const GenVector<T>& v1, const GenVector<T>& v2,`
`const GenVector<T>& v3)`

Shorthand to create bi- or tri-diagonal `BandMatrixes` if you already have the `Vectors`. The `Vectors` are in order from bottom to top in each case.

- `tmv::BandMatrix<T,stor,index> b1(const GenBandMatrix<T2>& b2)`
`b1 = b2`

Copy the `BandMatrix` `b2`, which may be of any type `T2` so long as values of type `T2` are convertible into type `T`. The assignment operator has the same flexibility.

- `tmv::BandMatrixView<T,index> b =`
`BandMatrixViewOf(MatrixView<T,index> m, int nlo, int nhi)`
`tmv::BandMatrixView<T,index> b =`
`BandMatrixViewOf(BandMatrixView<T,index> m, int nlo, int nhi)`
`tmv::BandMatrixView<T,index> b =`
`BandMatrixViewOf(DiagMatrixView<T,index> m)`
`tmv::BandMatrixView<T,index> b =`
`BandMatrixViewOf(UpperTriMatrixView<T,index> m)`
`tmv::BandMatrixView<T,index> b =`
`BandMatrixViewOf(UpperTriMatrixView<T,index> m, int nhi)`
`tmv::BandMatrixView<T,index> b =`
`BandMatrixViewOf(LowerTriMatrixView<T,index> m)`
`tmv::BandMatrixView<T,index> b =`
`BandMatrixViewOf(LowerTriMatrixView<T,index> m, int nlo)`

Make an `BandMatrixView` of the corresponding portion of `m`. There are also `ConstBandMatrixView` versions of all of these.

- `tmv::BandMatrixView<T,index> b(MatrixView<T,index> m,`
`int nlo, int nhi)`
`tmv::BandMatrixView<T,index> b(BandMatrixView<T,index> m,`
`int nlo, int nhi)`

For square matrices `m`, these (and the corresponding `ConstBandMatrixView` versions) work the same as the above `BandMatrixViewOf` commands. However, this version preserves the values of `nrows` and `ncols` from `m` even if some of the rows or columns do not include any of the new band. This is only important if `m` is not square.

For example, if `m` is 10×8 , then

```
tmv::BandMatrixView<T> b1(m,0,2);
```

will create a 10×8 `BandMatrixView` of `m`'s diagonal plus two super-diagonals, but

```
tmv::BandMatrixView<T> b2 = BandMatrixViewOf(m, 0, 2);
```

will instead create an 8×8 BandMatrixView of the same portion of m.

Note that the same difference holds for the BandMatrix constructor:

```
tmv::BandMatrix<T> b1(m, 0, 2);
```

will create a 10×8 BandMatrix, but

```
tmv::BandMatrix<T> b2 = BandMatrixViewOf(m, 0, 2);
```

will create an 8×8 BandMatrix.

- `tmv::BandMatrixView<T> b =`
`tmv::BandMatrixViewOf(T* vv, size_t ncols, size_t nrows,`
`int nlo, int nhi, StorageType stor)`
`tmv::ConstBandMatrixView<T> b =`
`tmv::BandMatrixViewOf(const T* vv, size_t ncols, size_t nrows,`
`int nlo, int nhi, StorageType stor)`

Make a BandMatrixView of the actual memory elements, vv.

6.2 Access

```
b.nrows() = b.colsize()
b.ncols() = b.rowsize()
b(i, j)
b.row(int i, int j1, int j2)
b.col(int i, int j1, int j2)
b.diag()
b.diag(int i)
b.diag(int i, int k1, int k2)
```

The versions of `row` and `col` with only one argument are missing, since the full row or column isn't accessible as a `VectorView`. You must specify a valid range within the row or column that you want, given the banded storage of `b`.

```
b.subVector(int i, int j, int istep, int jstep, int size)
b.subMatrix(int i1, int i2, int j1, int j2)
b.subMatrix(int i1, int i2, int j1, int j2, int istep, int jstep)
```

These work the same as for a `Matrix` (See 3.3), except that the entire subvector or submatrix must be completely within the band.

```
b.subBandMatrix(int i1, int i2, int j1, int j2, int newnlo, int newnhi)
b.subBandMatrix(int i1, int i2, int j1, int j2, int newnlo, int newnhi,
int istep, int jstep)
```

This returns a `BandMatrixView` of a subset of a `BandMatrix`. The `newnlo` and `newnhi` parameters do not need to be different from the existing `nlo` and `nhi` if `i1 = j1` (i.e. the new main diagonal is part of the old main diagonal). However, if you are moving the upper left corner off of the diagonal, you need to adjust `nlo` and `nhi` appropriately. For example, if `b` is a 6×6 `BandMatrix` with 2 sub-diagonals and 3 super-diagonals (like our example above), the 3 super-diagonals may be viewed as a `BandMatrixView` with `b.subBandMatrix(0, 5, 1, 6, 0, 2)`.

```

b.rowRange(int i1, int i2)
b.colRange(int j1, int j2)
b.diagRange(int k1, int k2)

```

These return a `BandMatrixView` of the parts of these rows, columns or diagonals that appear within the original banded structure. For our example of viewing just the super-diagonals of a 6×6 `BandMatrix` with 2 sub- and 3 super-diagonals, we could instead use `m.diagRange(1, 4)`. The last 3 rows would be `m.rowRange(3, 6)`. Note that this would be a 3×5 matrix with 0 sub-diagonals and 4 super-diagonals. These routines calculate the appropriate changes in the size and shape to include all of the relevant parts of the rows or columns.

```

b.upperBand()
b.lowerBand()

```

These return a `BandMatrixView` including the main diagonal and either the super- or sub-diagonals. The size is automatically set appropriately to include the entire band. (This is only non-trivial for non-square band matrices.) They are equivalent to `b.diagRange(0, b.nhi()+1)` and `b.diagRange(-b.nlo(), 1)` respectively.

```

b.upperBandOff()
b.lowerBandOff()

```

These return a `BandMatrixView` of only the off-diagonals of either the upper or lower half of the matrix. They are equivalent to `b.diagRange(1, b.nhi()+1)` and `b.diagRange(-b.nlo(), 0)` respectively. They are inspired by analogy with the combination `m.upperTri().offDiag()`. Since `BandMatrix` does not have the method `offDiag`, these provide the same functionality.

```

b.transpose()
b.conjugate()
b.adjoint()
b.view()
b.cView()
b.fView()
b.realPart()
b.imagPart()

```

6.3 Functions

```

RT b.norm1() = Norm1(b)
RT b.norm2() = Norm2(b) = b.doNorm2()
RT b.normInf() = NormInf(b)
RT b.normF() = NormF(b) = b.norm() = Norm(b)
RT b.normSq() = NormSq(b)
RT b.normSq(RT scale)
RT b.maxAbsElement() = MaxAbsElement(b)
RT b.maxAbs2Element() = MaxAbs2Element(b)
T b.trace() = Trace(b)
T b.sumElements() = SumElements(b)
RT b.sumAbsElements() = SumAbsElements(b)
T b.det() = Det(b)
RT b.logDet(T* sign=0) = LogDet(b)
bool b.isSingular()
RT b.condition()
RT b.doCondition()
minv = b.inverse() = Inverse(b)

```

```
b.makeInverse(Matrix<T>& minv)
b.makeInverseATA(Matrix<T>& cov)
```

The inverse of a `BandMatrix` is not (in general) banded. So `minv` here must be a regular `Matrix`.

```
b.setZero()
b.setTo(T x)
b.clip(RT thresh)
b.setToIdentity(T x = 1)
b.conjugateSelf()
b.transposeSelf()
Swap(b1, b2)
```

6.4 Arithmetic

In addition to `x`, `v`, and `m` from before, we now add `b` for a `BandMatrix`.

```
b2 = -b1
b2 = x * b1
b2 = b1 [*/] x
b3 = b1 [+ -] b2
m2 = m1 [+ -] b
m2 = b [+ -] m1
b [*/] = x
b2 [+ -] = b1
m [+ -] = b
v2 = b * v1
v2 = v1 * b
v *= b
b3 = b1 * b2
m2 = b * m1
m2 = m1 * b
m *= b
b2 = b1 [+ -] x
b2 = x [+ -] b1
b [+ -] = x
b = x
```

6.5 Division

The division operations are:

```
v2 = v1 [/ %] b
m2 = m1 [/ %] b
m2 = b [/ %] m1
m = b1 [/ %] b2
m = x [/ %] b
v [/ %] = b
m [/ %] = b
```

`BandMatrix` has three possible choices for the division decomposition:

1. `b.divideUsing(tm::LU)` does a normal LU decomposition, taking into account the band structure of the matrix, which greatly speeds up the calculation into the lower and upper (banded) triangles. This is the default decomposition to use for a square `BandMatrix` if you don't specify anything.

This decomposition can only really be done in place if either `nlo` or `nhi` is 0, in which case it is automatically done in place, since the `BandMatrix` is already lower or upper triangle. Thus, there is usually no reason to use the `divideInPlace()` method.

If this is not the case, and you really want to do the decomposition in place, you can declare a matrix with a wider band and view the portion that represents the matrix you actually want. This view then can be divided in place. More specifically, you need to declare the wider `BandMatrix` with `ColMajor` storage, with the smaller of $\{nlo, nhi\}$ as the number of sub-diagonals, and with $(nlo + nhi)$ as the number of super-diagonals. Then you can use `BandMatrixViewOf` to view the portion you want, transposing it if necessary. On the other hand, you are probably not going to get much of a speed gain from all of this finagling, so unless you are really memory starved, it's probably not worth it.

To access this decomposition, use:

```
bool b.lud().IsTrans()
tmv::LowerTriMatrix<T, UnitDiag> b.lud().getL()
tmv::ConstBandMatrixView<T> b.lud().getU()
Permutation b.lud().getP()
```

The following should result in a matrix numerically very close to `b`.

```
tmv::Matrix<T> m2 = b.lud().getP() * b.lud().getL() * b.lud().getU();
if (b.lud().IsTrans()) m2.transposeSelf();
```

2. `b.divideUsing(tm::QR)` will perform a QR decomposition. This is the default method for a non-square `BandMatrix`.

The same kind of convolutions need to be done to perform this in place as for the LU decomposition.

To access this decomposition, use:⁶

```
bool b.qrd().IsTrans()
tmv::Matrix<T> b.qrd().getQ()
tmv::ConstBandMatrixView<T> b.qrd().getR()
```

The following should result in a matrix numerically very close to `b`.

```
tmv::Matrix<T> m2(b.nrows, b.ncols);
tmv::MatrixView<T> m2v =
    b.qrd().IsTrans() ? b2.transpose() : b2.view();
m2v = b.qrd().getQ() * b.qrd().getR();
```

3. `b.divideUsing(tm::SV)` will perform a singular value decomposition.

This cannot be done in place.

To access this decomposition, use:

```
tmv::ConstMatrixView<T> b.svd().getU()
tmv::ConstDiagMatrixView<RT> b.svd().getS()
tmv::ConstMatrixView<T> b.svd().getV()
```

⁶ I have not yet made a version of the `PackedQ` class for `BandMatrix`. So unfortunately, here `getQ()` creates the matrix directly and is thus rather inefficient.

The product of these three should result in a matrix numerically very close to `b`.

There are the same control and access routines as for a regular SVD (See 3.7.7),

```
b.svd().thresh(RT thresh)
b.svd().top(int nsing)
RT b.svd().condition()
int b.svd().getKMax()
```

The routines

```
b.saveDiv()
b.setDiv()
b.resetDiv()
b.unsetDiv()
bool b.divIsSet()
b.divideInPlace()
```

work the same as for regular `Matrixes`. (See 3.7.5.)

6.6 Input/Output

The simplest output is the usual:

```
os << b
```

where `os` is any `std::ostream`. The output format is the same as for a `Matrix`. (See 3.8.)

There is also a compact format for a `BandMatrix`:

```
b.writeCompact(os)
```

outputs in the format:

```
B nrows ncols nlo nhi
( b(0,0)  b(0,1)  b(0,2)  ... b(0,nhi) )
( b(1,0)  b(1,1)  b(1,2)  ... b(1,nhi+1) )
...
( b(nlo,0) b(nlo,1) ... b(nlo,nlo+nhi) )
...
( b(nrows-nhi-1,nrows-nlo-nhi-1) ... b(nrows-nhi-1,ncols-1) )
...
( b(nrows-1,nrows-nlo-1) ... b(nrows-1,ncols-1) )
```

If `nrows` is not equal to `ncols`, then the above isn't exactly accurate. But the essence is the same: all the values in the band from each row are output one row at a time.

The same compact format can be read back in the usual two ways:

```
tmv::BandMatrix<T> b(nrows,ncols,nlo,nhi);
is >> b;
std::auto_ptr<tmv::BandMatrix<T> > bptr;
is >> bptr;
```

One can write small values as 0 with

```
b.write(std::ostream& os, RT thresh)
b.writeCompact(std::ostream& os, RT thresh)
```

7 Symmetric and hermitian matrices

The `SymMatrix` class is our symmetric matrix class. A symmetric matrix is one for which $m = m^T$. We also have a class called `HermMatrix`, which is our hermitian matrix class. A hermitian matrix is one for which $m = m^\dagger$. The two are exactly the same if T is real, but for complex T , they are different.

Both classes inherit from `GenSymMatrix`, which has an internal parameter to keep track of whether it is symmetric or hermitian. `GenSymMatrix` in turn inherits from `BaseMatrix`. The various views and composite classes described below also inherit from `GenSymMatrix`.

One general caveat about complex `HermMatrix` calculations is that the diagonal elements should all be real. Some calculations assume this, and only use the real part of the diagonal elements. Other calculations use the full complex value as it is in memory. Therefore, if you set a diagonal element to a non-real value at some point, the results will likely be wrong in unpredictable ways. Plus of course, your matrix won't actually be hermitian any more, so the right answer is undefined in any case.

All the `SymMatrix` and `HermMatrix` routines are included by:

```
#include "TMV_Sym.h"
```

In addition to the T template parameter, there are three other template parameters: `uplo`, which can be either `tmv::Upper` or `tmv::Lower`; `stor`, which can be either `tmv::RowMajor` or `tmv::ColMajor`; and `index`, which can be either `tmv::CStyle` or `tmv::FortranStyle`. The parameter `uplo` refers to which triangle the data are actually stored in, since the other half of the values are identical, so we do not need to reference them. The default values for these are `Lower`, `ColMajor`, and `CStyle` respectively.

The storage of a `SymMatrix` takes $N \times N$ elements of memory, even though approximately half of them are never used. Someday, I'll write the packed storage versions, which allow for efficient storage of the matrices.

Usually, the symmetric/hermitian distinction does not affect the use of the classes. (It does affect the actual calculations performed of course.) So we will use `s` for both, and just point out whenever a `HermMatrix` acts differently from a `SymMatrix`.

Most functions and methods for `SymMatrix` and `HermMatrix` work the same as they do for `Matrix`. In these cases, we will just list the functions that are allowed with the effect understood to be the same as for a regular `Matrix`. Of course, there are almost always algorithmic speed-ups, which the code will use to take advantage of the symmetric or hermitian structure. Whenever there is a difference in how a function works, we will explain the difference.

7.1 Constructors

- `tmv::SymMatrix<T,uplo,stor,index> s(size_t n)`
`tmv::HermMatrix<T,uplo,stor,index> s(size_t n)`

Makes an $n \times n$ `SymMatrix` or `HermMatrix` with uninitialized values. If debugging is turned on (i.e. not turned off with `-DNDEBUG`), then the values are in fact initialized to 888.

- `tmv::SymMatrix<T,uplo,stor,index> s(size_t n, T x)`
`tmv::HermMatrix<T,uplo,stor,index> s(size_t n, RT x)`

Makes an $n \times n$ `SymMatrix` or `HermMatrix` with all values equal to `x`. For the `HermMatrix` version of this, `x` must be real.

- `tmv::SymMatrix<T,uplo,stor,index> s(size_t n, const T* vv)`
`tmv::SymMatrix<T,uplo,stor,index> s(size_t n,`
 `const std::vector<T>& vv)`
`tmv::HermMatrix<T,uplo,stor,index> s(size_t n, const T* vv)`
`tmv::HermMatrix<T,uplo,stor,index> s(size_t n,`
 `const std::vector<T>& vv)`

Makes a SymMatrix or HermMatrix which copies the elements from vv.

- `tmv::SymMatrix<T,uplo,stor,index> s(const GenMatrix<T>& m)`
`tmv::HermMatrix<T,uplo,stor,index> s(const GenMatrix<T>& m)`
`tmv::SymMatrix<T,uplo,stor,index> s(const GenDiagMatrix<T>& d)`
`tmv::HermMatrix<T,uplo,stor,index> s(const GenDiagMatrix<T>& d)`

Makes a SymMatrix or HermMatrix which copies the corresponding values of m.

- `tmv::SymMatrix<T,uplo,stor,index> s1(const GenSymMatrix<T2>& s2)`
`tmv::HermMatrix<T,uplo,stor,index> s1(const GenSymMatrix<T2>& s2)`
`s1 = s2`

Copy the GenSymMatrix s2, which may be of any type T2 so long as values of type T2 are convertible into type T. The assignment operator has the same flexibility. If T and T2 are complex, then m1 and m2 need to be either both symmetric or both hermitian.

- `tmv::SymMatrixView<T,index> s =`
`tmv::SymMatrixViewOf(MatrixView<T,index> m, UpLoType uplo)`
`tmv::ConstSymMatrixView<T,index> s =`
`tmv::SymMatrixViewOf(ConstMatrixView<T,index> m, UpLoType uplo)`
`tmv::SymMatrixView<T,index> s =`
`tmv::HermMatrixViewOf(MatrixView<T,index>& m, UpLoType uplo)`
`tmv::ConstSymMatrixView<T,index> s =`
`tmv::HermMatrixViewOf(ConstMatrixView<T,index>& m, UpLoType uplo)`

Make a SymMatrixView of the corresponding portion of m.

- `tmv::SymMatrixView<T,index> s =`
`tmv::SymMatrixViewOf(T* vv, size_t n, UpLoType uplo,`
`StorageType stor)`
`tmv::ConstSymMatrixView<T,index> s =`
`tmv::SymMatrixViewOf(const T* vv, size_t n, UpLoType uplo,`
`StorageType stor)`
`tmv::SymMatrixView<T,index> s =`
`tmv::HermMatrixViewOf(T* vv, size_t n, UpLoType uplo,`
`StorageType stor)`
`tmv::ConstSymMatrixView<T,index> s =`
`tmv::HermMatrixViewOf(const T* vv, size_t n, UpLoType uplo,`
`StorageType stor)`

Make a SymMatrixView of the actual memory elements, vv, in either the upper or lower triangle. vv must be of length $n \times n$, even though only about half of the values are actually used,

7.2 Access

```
s.nrows() = s.ncols() = s.colsize() = s.rowsize() = s.size()
s(i,j)
s.row(int i, int j1, int j2)
s.col(int i, int j1, int j2)
s.diag()
s.diag(int i)
s.diag(int i, int k1, int k2)
```

For the mutable `d(i, j)` version, As for triangle matrices, the versions of `row` and `col` with only one argument are missing, since the full row or column isn't accessible as a `VectorView`. You must specify a valid range within the row or column that you want, given the upper or lower storage of `s`. The diagonal element may be in a `VectorView` with either elements in the lower triangle or the upper triangle, but not both. To access a full row, you would therefore need to use two steps:

```
s.row(i, 0, i) = ...
s.row(i, i, ncols) = ...

s.subVector(int i, int j, int istep, int jstep, int size)
s.subMatrix(int i1, int i2, int j1, int j2)
s.subMatrix(int i1, int i2, int j1, int j2, int istep, int jstep)
```

These work the same as for a `Matrix` (See 3.3.) except that the entire subvector or submatrix must be completely within the upper or lower triangle.

```
s.subSymMatrix(int i1, int i2, int istep = 1)
```

This returns a `SymMatrixView` of `s` whose upper-left corner is `s(i1, i1)`, and whose lower-right corner is `s(i2-istep, i2-istep)`. If `istep` \neq 1, then it is the step in both the `i` and `j` directions.

```
s.upperTri(DiagType dt=NonUnitDiag)
s.lowerTri(DiagType dt=NonUnitDiag)
s.unitUpperTri()
s.unitLowerTri()
```

Both of these are valid, regardless of which triangle stores the actual data for `s`.

```
s.transpose()
s.conjugate()
s.adjoint()
s.view()
s.cView()
s.fView()
s.realPart()
s.imagPart()
```

Note that the imaginary part of a complex hermitian matrix is skew-symmetric, so `s.imagPart()` is illegal for a `HermMatrix`. If you need to deal with the imaginary part of a `HermMatrix`, you would have to view them as a triangle matrix with `s.upperTri().imagPart()`. Or, since the diagonal elements are all real. you could also use `s.upperTri().offDiag().imagPart()`.

7.3 Functions

```
RT s.norm1() = Norm1(s)
RT s.norm2() = Norm2(s) = s.doNorm2()
RT s.normInf() = NormInf(s)
RT s.normF() = NormF(s) = s.norm() = Norm(s)
RT s.normSq() = NormSq(s)
RT s.normSq(RT scale)
RT s.maxAbsElement() = MaxAbsElement(s)
RT s.maxAbs2Element() = MaxAbs2Element(s)
T s.trace() = Trace(s)
T s.sumElements() = SumElements(s)
RT s.sumAbsElements() = SumAbsElements(s)
```

```

T s.det() = Det(s)
RT s.logDet(T* sign=0) = LogDet(s)
bool s.isSingular()
RT s.condition()
RT s.doCondition()
sinv = s.inverse() = Inverse(s)
s.makeInverse(Matrix<T>& sinv)
s.makeInverse(SymMatrix<T>& sinv)
s.makeInverseATA(Matrix<T>& cov)

```

Since the inverse of an `SymMatrix` is also symmetric, the object returned by `s.inverse()` is assignable to a `SymMatrix`. Of course you can also assign it to a regular `Matrix` if you prefer. Similarly, there are versions of `s.makeInverse(minv)` for both argument types.

```

s.setZero()
s.setAllTo(T x)
s.clip(RT thresh)
s.setToIdentity(T x = 1)
s.conjugateSelf()
s.transposeSelf()
Swap(s1,s2)

```

`transposeSelf` does nothing to a `SymMatrix` and is equivalent to `conjugateSelf` for a `HermMatrix`.

```

s.swapRowsCols(int i1, int i2)
s.permuteRowsCols(const int* p)
s.reversePermuteRowsCols(const int* p)
s.permuteRowsCols(const int* p, int i1, int i2)
s.reversePermuteRowsCols(const int* p, int i1, int i2)

```

The new method, `swapRowsCols`, would be equivalent to

```
s.swapRows(i1,i2).swapCols(i1,i2);
```

except that neither of these functions are allowed for a `SymMatrix`, since they result in non-symmetric matrices. Only the combination of both maintains the symmetry of the matrix. So this combination is included as a method. The `permute` methods performs a series of these combined swaps.

7.4 Arithmetic

In addition to `x`, `v`, and `m` from before, we now add `s` for a `SymMatrix`.

```

s2 = -s1
s2 = x * s1
s2 = s1 [*/] x
s3 = s1 [+ -] s2
m2 = m1 [+ -] s
m2 = s [+ -] m1
s [*/] = x
s2 [+ -] = s1
m [+ -] = s
v2 = s * v1
v2 = v1 * s
v *= s
m = s1 * s2

```

```

m2 = s * m1
m2 = m1 * s
m *= s
s2 = s1 [+-] x
s2 = x [+-] s1
s [+-] = x
s = x
s = v ^ v
s [+-] = v ^ v
s = m * m.transpose()
s [+-] = m * m.transpose()
s = U * U.transpose()
s [+-] = U * U.transpose()
s = L * L.transpose()
s [+-] = L * L.transpose()

```

For outer products, both v 's need to be the same actual data. If s is complex hermitian, then it should actually be $s = v \wedge v.conjugate()$. Likewise for the next three (called “rank-k updates”), the m 's, L 's and U 's need to be the same data, and for a complex hermitian matrix, `transpose()` should be replaced with `adjoint()`.

There is a minor issue with `SymMatrix` arithmetic that is worth knowing about. Because TMV uses the same base class for both hermitian and symmetric matrices, the compiler cannot tell the difference between them in arithmetic operations. This can become an issue when doing complicated arithmetic operations with complex hermitian or symmetric matrices.

Some operations with a `GenSymMatrix`, as far as the compiler can tell, may or may not result in another `SymMatrix`. For example, a complex scalar times a complex `HermMatrix` is not hermitian (or symmetric), but a complex scalar times a `SymMatrix` is still symmetric. Likewise the sum of two `GenSymMatrix`s may or may not be hermitian or symmetric.

So the complex `SymMatrixComposite` class⁷, which is the return type of these and similar operations, is derived from `GenMatrix` rather than from `GenSymMatrix`. This means that if you let it self-instantiate, rather than assign it to a `SymMatrix`, it will instantiate as a regular `Matrix`. However, it is assignable to a `SymMatrix`, despite deriving from `GenMatrix`, so if your operation should be allowed, then the assignment will work fine.

Most of the time, this won't matter, since you will generally assign the result to either a `SymMatrix` or a `Matrix` as appropriate. However, if you let your expression get more complicated than a single matrix addition, multiplication, etc., then some things that should be allowed give compiler errors. For example:

```
s3 += x*s1+s2;
```

is not legal for complex symmetric $s1$, $s2$, $s3$, even though this is valid mathematically. This is because there is no 3-matrix `Add` function. (The `AddMM` function just adds a multiple of one matrix to another.) So the right hand side needs to be instantiated before being added to the left side, and it will instantiate as a regular `Matrix`, which cannot be added to a `SymMatrix`. If the “+=” had been just “=”, then we wouldn't have any problem, since the composite object that stores $x*s1+s2$ is assignable to a `SymMatrix`.

One work-around is to explicitly tell the compiler to instantiate the right hand side as a `SymMatrix`:

```
s3 += SymMatrix<T>(x*s1+s2);
```

Another work-around, which I suspect will usually be preferred, is to break the equation into multiple statements, each of which are simple enough to not require any instantiation:

```
s3 += x*s1;
```

⁷ Note: all of these comments only apply to complex `SymMatrix` and `HermMatrix` arithmetic. The real version of the `SymMatrixComposite` does derive from `GenSymMatrix`, since there is no ambiguity in that case.

```
s3 += s2;
```

7.5 Division

The division operations are:

```
v2 = v1 [%] s
m2 = m1 [%] s
m2 = s [%] m1
m = s1 [%] s2
s1 = x [%] s2
v [%] = s
m [%] = s
```

`SymMatrix` has three possible choices for the decomposition to use for division:

1. `m.divideUsing(tmv::LU)` will perform something similar to the LU decomposition for regular matrices. But in fact, it does what is called an LDL or Bunch-Kaufman decomposition.

A permutation of `m` is decomposed into a lower triangle matrix (L) times a symmetric block diagonal matrix (D) times the transpose of L . D has either 1x1 and 2x2 blocks down the diagonal. For hermitian matrices, the third term is the adjoint of L rather than the transpose.

This is the default decomposition to use if you don't specify anything.

To access this decomposition, use:

```
ConstLowerTriMatrixView<T> s.lud().getL()
Matrix<T> s.lud().getD()
Permutation s.lud().getP()
```

The following should result in a matrix numerically very close to `s`.

```
Matrix<T> m2 = s.lud().getP() * s.lud().getL() * s.lud().getD() *
              s.lud().getL().transpose() * s.lud().getP().transpose();
```

For a complex hermitian `s`, you would need to replace `transpose` with `adjoint`.

2. `s.divideUsing(tmv::CH)` will perform a Cholesky decomposition. The matrix `s` must be hermitian (or real symmetric) to use `CH`, since that is the only kind of matrix that has a Cholesky decomposition.

It is also similar to an LU decomposition, where U is the adjoint of L , and there is no permutation. It can be a bit dangerous, since not all hermitian matrices have such a decomposition, so the decomposition could fail. Only so-called “positive-definite” hermitian matrices have a Cholesky decomposition. A positive-definite matrix has all positive real eigenvalues. In general, hermitian matrices have real, but not necessarily positive eigenvalues.

One example of a positive-definite matrix is $s = A^\dagger A$ where A is any matrix. Then s is guaranteed to be positive-semi-definite (which means some of the eigenvalues may be 0, but not negative). In this case, the routine will usually work, but still might fail from numerical round-off errors if s is nearly singular.

When the decomposition fails, it throws an object of type `NonPosDef`.

See §10.4 for some more discussion about positive-definite matrices.

The only advantage of Cholesky over Bunch-Kaufman is speed. (And only about 20 to 30% at that.) If you know your matrix is positive-definite, the Cholesky decomposition is the fastest way to do division.

To access this decomposition, use:


```
ConstLowerTriMatrixView<T> s.chd().getL()
```

The following should result in a matrix numerically very close to s .

```
Matrix<T> m2 = s.chd().getL() * s.chd().getL().adjoint()
```

3. `s.divideUsing(tmv::SV)` will perform either an eigenvalue decomposition (for hermitian and real symmetric matrices) or a regular singular value decomposition (for complex symmetric matrices).

For hermitian matrices (including real symmetric matrices), the eigenvalue decomposition is $H = USU^\dagger$, where U is unitary and S is diagonal. So this would be identical to a singular value decomposition where $V = U^\dagger$, except that the elements of S , the eigenvalues of H , may be negative.

However, this decomposition is just as useful for division, dealing with singular matrices just as elegantly. It just means that internally, we allow the values of S to be negative, taking the absolute value when necessary (e.g. for `norm2`). The below access commands finish the calculation of S, V so that the $S(i)$ values are positive.

To access this decomposition, use:

```
ConstMatrixView<T> s.svd().getU()
DiagMatrix<RT> s.svd().getS()
Matrix<T> s.svd().getV()
```

The following should result in a matrix numerically very close to s .

```
Matrix<T> m2 = s.svd().getU() * s.svd().getS() * s.svd().getV()
```

For a complex symmetric s , the situation is not as convenient. In principle, one could find a decomposition $s = USU^T$ where S is a complex diagonal matrix, but such a decomposition is not easy to find. So for complex symmetric matrices, we just do the normal SVD: $s = USV$, although the algorithm does use the symmetry of the matrix to speed up portions of the algorithm relative to that for a general matrix.

The access is also necessarily different, since the object returned by `s.svd()` implicitly assumes that $V = U.adjoint()$ (modulo some sign changes), so we need a new accessor: `s.symsvd()`. Its `getS` and `getV` methods return Views rather than instantiated matrices.

Both versions also have the same control and access routines as a regular SVD (See 3.7.7):

```
s.svd().thresh(RT thresh)
s.svd().top(int nsing)
RT s.svd().condition()
int s.svd().getKMax()
```

(Likewise for `s.symsvd()`.)

The routines

```
s.saveDiv()
s.setDiv()
s.resetDiv()
s.unsetDiv()
bool s.divIsSet()
s.divideInPlace()
```

work the same as for regular `Matrixes`. (See 3.7.5.)

7.6 Input/Output

The simplest output is the usual:

```
os << s
```

where `os` is any `std::ostream`. The output format is the same as for a `Matrix`.

There is also a compact format for a `SymMatrix`:

```
s.writeCompact(os)
```

outputs in the format:

```
H/S n
( s(0,0) )
( s(1,0)  s(1,1) )
...
( s(n-1,0)  s(n-1,1) ... s(n-1,n-1) )
```

where H/S means either the character H or S, which indicates whether the matrix is hermitian or symmetric. In each case, the same compact format can be read back in the usual two ways:

```
tmv::SymMatrix<T> s(n);
tmv::HermMatrix<T> h(n);
is >> s >> h;
std::auto_ptr<tmv::SymMatrix<T> > sptr;
std::auto_ptr<tmv::HermMatrix<T> > hptr;
is >> sptr >> hptr;
```

One can write small values as 0 with

```
s.write(std::ostream& os, RT thresh)
s.writeCompact(std::ostream& os, RT thresh)
```

8 Symmetric and hermitian band matrices

The `SymBandMatrix` class is our symmetric band matrix, which combines the properties of `SymMatrix` and `BandMatrix`; it has a banded structure and $m = m^T$. Likewise `HermBandMatrix` is our hermitian band matrix for which $m = m^\dagger$.

Both classes inherit from `GenSymBandMatrix`, which in turn inherits from `BaseMatrix`. The various views and composite classes described below also inherit from `GenSymBandMatrix`.

As with the documentation for `SymMatrix`/`HermMatrix`, the descriptions below will only be written for `SymBandMatrix` with the implication that a `HermBandMatrix` has the same functionality, but with the calculations appropriate for a hermitian matrix, rather than symmetric.

One general caveat about complex `HermBandMatrix` calculations is that the diagonal elements should all be real. Some calculations assume this, and only use the real part of the diagonal elements. Other calculations use the full complex value as it is in memory. Therefore, if you set a diagonal element to a non-real value at some point, the results will likely be wrong in unpredictable ways. Plus of course, your matrix will not actually be hermitian any more, so the right answer is undefined in any case.

All the `SymBandMatrix` and `HermBandMatrix` routines are included by:

```
#include "TMV_SymBand.h"
```

In addition to the `T` template parameter, there are three other template parameters: `uplo`, which can be either `tmv::Upper` or `tmv::Lower`; `stor`, which can be one of `tmv::RowMajor`, `tmv::ColMajor`, or `tmv::DiagMajor`; and `index`, which can be either `tmv::CStyle` or `tmv::FortranStyle`. The default values for these are `Lower`, `ColMajor`, `CStyle`.

The storage size required is the same as for the `BandMatrix` of the upper or lower band portion. (See 6.) As with square band matrices, all three storage methods always need the same amount of memory, so the decision about which method to use should generally be based on performance considerations rather than memory usage. The speed of the various matrix operations are different for the different storage types. If the matrix calculation speed is important, it may be worth trying all three to see which is fastest for the operations you are using.

Also, as with `BandMatrix`, the storage for `Lower`, `DiagMajor` does not start with the upper left element as usual. Rather, it starts at the start of the lowest sub-diagonal. So for the constructors that take the matrix information from an array (`T*` or `vector<T>`), the start of the array needs to be at the start of the lowest sub-diagonal.

Most functions and methods for `SymBandMatrix` and `HermBandMatrix` work the same as they do for `Matrix`. In these cases, we will just list the functions that are allowed with the effect understood to be the same as for a regular `Matrix`. Of course, there are almost always algorithmic speed-ups, which the code will use to take advantage of the symmetric (or hermitian) banded structure. Whenever there is a difference in how a function works, we will explain the difference.

8.1 Constructors

- `tmv::SymBandMatrix<T, uplo, stor, index> sb(size_t n, int nlo)`

Makes an $n \times n$ `SymBandMatrix` with `nlo` off-diagonals and with uninitialized values. If debugging is turned on (i.e. not turned off with `-DNDEBUG`), then the values are initialized to 888.

- `tmv::SymBandMatrix<T, uplo, stor, index> sb(size_t n, int nlo, T x)`

Makes an $n \times n$ `SymBandMatrix` with `nlo` off-diagonals and with all values equal to `x`.

- `tmv::SymBandMatrix<T, uplo, stor, index> sb(size_t n, int nlo, const T* vv)`
`tmv::SymBandMatrix<T, uplo, stor, index> sb(size_t n, int nlo, const std::vector<T>& vv)`

Makes a `SymBandMatrix` which copies the elements from `vv`.

- `tmv::SymBandMatrix<T,uplo,stor,index> sb(const GenMatrix<T>& m, int nlo)`
`tmv::SymBandMatrix<T,uplo,stor,index> sb(const GenSymMatrix<T>& m, int nlo)`
`tmv::SymBandMatrix<T,uplo,stor,index> sb(const GenBandMatrix<T>& m, int nlo)`
`tmv::SymBandMatrix<T,uplo,stor,index> sb(const GenSymBandMatrix<T>& m, int nlo)`

Makes a `SymBandMatrix` which copies the corresponding values of `m`. For the last two, `nlo` must not be larger than the number of upper or lower bands in `m`.

- `tmv::SymBandMatrix<T,uplo,DiagMajor> sb = SymTriDiagMatrix(const GenVector<T>& v1, const GenVector<T>& v2)`
`tmv::SymBandMatrix<T,uplo,DiagMajor> sb = SymTriDiagMatrix(const GenVector<T>& v1, const GenVector<T>& v2)`
`tmv::HermBandMatrix<T,uplo,DiagMajor> sb = HermTriDiagMatrix(const GenVector<T>& v1, const GenVector<T>& v2, UpLoType uplo)`
`tmv::HermBandMatrix<T,uplo,DiagMajor> sb = HermTriDiagMatrix(const GenVector<T>& v1, const GenVector<T>& v2, UpLoType uplo)`

Shorthand to create a symmetric tri-diagonal band matrix if you already have the Vectors. The main diagonal is `v1` and the off-diagonal is `v2`.

With a `HermTriDiagMatrix`, `v1` should be real, although it may be either a real-valued Vector or a complex-valued Vector whose imaginary components are all zero. Also, `HermTriDiagMatrix` takes an extra parameter, `uplo`, indicating whether `v2` should be used as the upper or lower off-diagonal.

- `tmv::SymBandMatrix<T> sb1(const GenSymBandMatrix<T2>& sb2)`
`sb1 = sb2`

Copy the `SymBandMatrix` `m2`, which may be of any type `T2` so long as values of type `T2` are convertible into type `T`. The assignment operator has the same flexibility.

- `tmv::SymBandMatrixView<T> sb = SymBandMatrixViewOf(MatrixView<T> m, UpLoType uplo, int nlo)`
`tmv::SymBandMatrixView<T> sb = SymBandMatrixViewOf(SymMatrixView<T> m, int nlo)`
`tmv::SymBandMatrixView<T> sb = SymBandMatrixViewOf(BandMatrixView<T> m, UpLoType uplo, int nlo)`
`tmv::SymBandMatrixView<T> sb = HermBandMatrixViewOf(MatrixView<T> m, UpLoType uplo, int nlo)`
`tmv::SymBandMatrixView<T> sb = HermBandMatrixViewOf(SymMatrixView<T> m, int nlo)`
`tmv::SymBandMatrixView<T> sb = HermBandMatrixViewOf(BandMatrixView<T> m, UpLoType uplo, int nlo)`

Make an `SymBandMatrixView` of the corresponding portion of `m`. To view these as a hermitian band matrix, use the command, `HermBandMatrixViewOf` instead. For the view of a `BandMatrix`, the

parameter `nlo` may be omitted, in which case either `m.nhi()` or `m.nlo()` is used according to whether `uplo` is Upper or Lower respectively. There are also `ConstSymBandMatrixView` versions of these.

- `tmv::SymBandMatrixView<T> sb =`
`tmv::SymBandMatrixViewOf(T* vv, size_t n, int nlo,`
`UpLoType uplo, StorageType stor)`
`tmv::ConstSymBandMatrixView<T> sb =`
`tmv::SymBandMatrixViewOf(const T* vv, size_t n, int nlo,`
`UpLoType uplo, StorageType stor)`
`tmv::SymBandMatrixView<T> sb =`
`tmv::HermBandMatrixViewOf(T* vv, size_t n, int nlo,`
`UpLoType uplo, StorageType stor)`
`tmv::ConstSymBandMatrixView<T> sb =`
`tmv::HermBandMatrixViewOf(const T* vv, size_t n, int nlo,`
`UpLoType uplo, StorageType stor)`

Make a `SymBandMatrixView` of the actual memory elements, `vv`.

8.2 Access

```
sb.nrows() = sb.ncols() = sb.colsize() = sb.rowsize() = sb.size()
sb.nlo() = sb.nhi()
sb(i, j)
sb.row(int i, int j1, int j2)
sb.col(int i, int j1, int j2)
sb.diag()
sb.diag(int i)
sb.diag(int i, int k1, int k2)
```

For the mutable `d(i, j)` version, Again, the versions of `row` and `col` with only one argument are missing, since the full row or column isn't accessible as a `VectorView`. You must specify a valid range within the row or column that you want, given the banded storage of `sb`. And, like for `SymMatrix`, a full row must be accessed in its two parts, one on each side of the diagonal.

```
sb.subVector(int i, int j, int istep, int jstep, int size)
sb.subMatrix(int i1, int i2, int j1, int j2)
sb.subMatrix(int i1, int i2, int j1, int j2, int istep, int jstep)
sb.subBandMatrix(int i1, int i2, int j1, int j2, int newnlo, int newhi)
sb.subBandMatrix(int i1, int i2, int j1, int j2, int newnlo, int newhi,
int istep, int jstep)
sb.diagRange(int k1, int k2)
sb.upperBand()
sb.lowerBand()
sb.upperBandOff()
sb.lowerBandOff()
```

These work the same as for a `BandMatrix` (See 6.2), except that the entire subvector or submatrix must be completely within the upper or lower band.

```
sb.subSymMatrix(int i1, int i2)
sb.subSymMatrix(int i1, int i2, int istep)
sb.subSymBandMatrix(int i1, int i2, int newnlo=m.nlo())
sb.subSymBandMatrix(int i1, int i2, int newnlo, int istep)
```

These return a view of the `SymMatrix` or `SymBandMatrix` which runs from `i1` to `i2` along the diagonal with an optional step, and includes the off-diagonals in the same rows/cols. For the first two, the `SymMatrix` must be completely within the band.

```
sb.symDiagRange(int newnlo)
```

Since `diagRange` returns a regular `BandMatrixView`, it must be completely within either the upper or lower band. This routine returns a `SymBandMatrixView` which straddles the diagonal with `newnlo` super- and sub-diagonals.

```
sb.transpose()
sb.conjugate()
sb.adjoint()
sb.view()
sb.cView()
sb.fView()
sb.realPart()
sb.imagPart()
```

Note that the imaginary part of a complex hermitian band matrix is skew-symmetric, so `sb.imagPart()` is illegal for a `HermBandMatrix`. If you need to manipulate the imaginary part of a `HermMatrix`, you could use `sb.upperBandOff().imagPart()` (since all the diagonal elements are real).

8.3 Functions

```
RT sb.norm1() = Norm1(sb)
RT sb.norm2() = Norm2(sb) = sb.doNorm2()
RT sb.normInf() = NormInf(sb)
RT sb.normF() = NormF(sb) = sb.norm() = Norm(sb)
RT sb.normSq() = NormSq(sb)
RT sb.normSq(RT scale)
RT sb.maxAbsElement() = MaxAbsElement(sb)
RT sb.maxAbs2Element() = MaxAbs2Element(sb)
T sb.trace() = Trace(sb)
T sb.sumElements() = SumElements(sb)
RT sb.sumAbsElements() = SumAbsElements(sb)
T sb.det() = Det(sb)
RT sb.logDet(T* sign=0) = LogDet(sb)
sinv = sb.inverse() = Inverse(sb)
bool sb.isSingular
RT sb.condition()
RT sb.doCondition()
sb.makeInverse(Matrix<T>& minv)
sb.makeInverse(SymMatrix<T>& sinv)
sb.makeInverseATA(Matrix<T>& cov)
```

The inverse of a `SymBandMatrix` is not (in general) banded. However, it is symmetric (or hermitian). So `sb.inverse()` may be assigned to either a `Matrix` or a `SymMatrix`.

```
sb.setZero()
sb.setAllTo(T x)
sb.clip(RT thresh)
sb.setToIdentity(T x = 1)
sb.conjugateSelf()
```

```
sb.transposeSelf()
Swap(sb1, sb2)
```

8.4 Arithmetic

In addition to `x`, `v`, `m`, `b` and `s` from before, we now add `sb` for a `SymBandMatrix`.

```
sb2 = -sb1
sb2 = x * sb1
sb2 = sb1 [*/] x
sb3 = sb1 [+-] sb2
m2 = m1 [+-] sb
m2 = sb [+-] m1
b2 = b1 [+-] sb
b2 = sb [+-] b1
s2 = s1 [+-] sb
s2 = sb [+-] s1
sb [*/]= x
sb2 [+-]= sb1
m [+-]= sb
b [+-]= sb
s [+-]= sb
v2 = sb * v1
v2 = v1 * sb
v *= sb
b = sb1 * sb2
m2 = sb * m1
m2 = m1 * sb
m *= sb
b2 = sb * b1
b2 = b1 * sb
b *= sb
m2 = sb * s1
m2 = s1 * sb
sb2 = sb1 [+-] x
sb2 = x [+-] sb1
sb [+-]= x
sb = x
```

8.5 Division

The division operations are:

```
v2 = v1 [/ %] sb
m2 = m1 [/ %] sb
m2 = sb [/ %] m1
m = sb1 [/ %] sb2
s = x [/ %] sb
v [/ %]= sb
m [/ %]= sb
```

`SymBandMatrix` has three possible choices for the division decomposition:

1. `m.divideUsing(tmv::LU)` actually does the `BandMatrix` version of LU, rather than a Bunch-Kaufman algorithm like for `SymMatrix`. The reason is that the pivots in the Bunch-Kaufman algorithm can arbitrarily expand the band width required to hold the information. The generic banded LU algorithm is limited to $3*nlo+1$ bands.

To access this decomposition, use:

```
bool sb.lud().IsTrans()
tmv::LowerTriMatrix<T,UnitDiag> sb.lud().getL()
tmv::ConstBandMatrixView<T> sb.lud().getU()
Permutation sb.lud().getP()
```

The following should result in a matrix numerically very close to `sb`.

```
tmv::Matrix<T> m2(sb.nrows(),sb.ncols);
tmv::MatrixView<T> m2v =
    sb.lud().IsTrans() ? m2.transpose() : m2.view();
m2v = sb.lud().getP() * sb.lud().getL() * sb.lud().getU();
```

2. `sb.divideUsing(tmv::CH)` will perform a Cholesky decomposition. `sb` must be hermitian (or real symmetric) to use CH, since that is the only kind of matrix that has a Cholesky decomposition.

As with a regular `SymMatrix`, the only real advantage of Cholesky over LU decomposition is speed. If you know your matrix is positive-definite, the Cholesky decomposition is the fastest way to do division.

If `sb` is tri-diagonal (i.e. `nlo = 1`), then we use a slightly different algorithm, which avoids the square roots required for a Cholesky decomposition. Namely, we form the decomposition $sb = LDL^\dagger$, where L is a unit-diagonal lower banded matrix with 1 sub-diagonal, and D is diagonal.

If `sb` has `nlo > 1`, then we just use a normal Cholesky algorithm where $sb = LL^\dagger$ and L is lower banded with the same `nlo` as `sb`.

Both versions of the algorithm are accessed with the same methods:

```
BandMatrix<T> sb.chd().getL()
DiagMatrix<T> sb.chd().getD()
```

with L being made unit-diagonal or D being set to the identity matrix as appropriate. (Obviously, `getL()` contains all of the information for the non-tridiagonal version.)

The following should result in a matrix numerically very close to `sb`.

```
Matrix<T> m2 = sb.chd().getL() * sb.chd().getD() *
    sb.chd().getL().adjoint()
```

3. `sb.divideUsing(tmv::SV)` will perform either an eigenvalue decomposition (for hermitian band and real symmetric band matrices) or a regular singular value decomposition (for complex symmetric band matrices).

To access this decomposition, use:

```
ConstMatrixView<T> sb.svd().getU()
DiagMatrix<RT> sb.svd().getS()
Matrix<T> sb.svd().getV()
```

(As for `SymMatrix`, a complex symmetric matrix needs to use the accessor `symsvd()` instead, whose `getS` and `getV` methods return Views rather than instantiated matrices.)

The following should result in a matrix numerically very close to `sb`.


```
Matrix<T> m2 = sb.svd().getU() * sb.svd().getS() * sb.svd().getV()
```

Both versions also have the same control and access routines as a regular SVD: (See 3.7.7):

```
sb.svd().thresh(RT thresh)
sb.svd().top(int nsing)
RT sb.svd().condition()
int sb.svd().getKMax()
```

(Likewise for `sb.symsvd()`.)

The routines

```
sb.saveDiv()
sb.setDiv()
sb.resetDiv()
sb.unsetDiv()
bool sb.divIsSet()
sb.divideInPlace()
```

work the same as for regular `Matrixes`. (See 3.7.5.)

8.6 Input/Output

The simplest output is the usual:

```
os << sb
```

where `os` is any `std::ostream`. The output format is the same as for a `Matrix`. (See 3.8.)

There is also a compact format for a `BandMatrix`:

```
sb.writeCompact(os)
\index{SymBandMatrix!methods!writeCompact}
```

outputs in the format:

```
hB/sB n nlo
( sb(0,0) )
( sb(1,0) sb(1,1) )
...
( sb(nlo,0) sb(nlo,1) ... sb(nlo,nlo) )
( sb(nlo+1,1) sb(nlo+1,2) ... sb(nlo+1,nlo+1) )
...
( sb(n-1,n-nlo-1) ... sb(n-1,n-1) )
```

where `hB/sB` means either `hB` or `sB`, which indicates whether the matrix is hermitian or symmetric.

The same compact format can be read back in the usual two ways:

```
tmv::SymBandMatrix<T> sb(n,nlo);
tmv::HermBandMatrix<T> hb(n,nlo);
is >> sb >> hb;
std::auto_ptr<tmv::SymBandMatrix<T> > sbptr;
std::auto_ptr<tmv::HermBandMatrix<T> > hbptr;
is >> sbptr >> hbptr;
```

One can write small values as 0 with

```
sb.write(std::ostream& os, RT thresh)
sb.writeCompact(std::ostream& os, RT thresh)
```

9 Permutations

The `Permutation` class is our permutation matrix class. A permutation matrix is square matrix with exactly one 1 in each row and column, and all the rest of the elements equal to zero.

However, internally we do not store a permutation this way. Instead, we treat a permutation as a series of pair-wise interchanges. This seems to be the fastest way to apply a permutation to a matrix or vector, rather than using an index-based method.

Also, I didn't bother to have a `PermutationView` class. Instead, the `Permutation` object keeps track of whether it owns its data or is just referencing values kept somewhere else. Whenever you perform a mutable action on the object, it copies the values if necessary. So you cannot indirectly modify another `Permutation` the way you can with `MatrixView`.

9.1 Constructors

- `tmv::Permutation p(size_t n)`

Makes an $n \times n$ `Permutation` set initially to the identity matrix.

- `tmv::Permutation p(size_t n, const int* pp, bool isinv, int det)`

Makes an $n \times n$ `Permutation` using the provided values as the list of interchanges. The meaning of `pp` is that $v = p * v$ is equivalent to

```
if (isinv) v.reversePermute(pp);  
else v.permute(pp);
```

The last two parameters may be omitted. If `det` is omitted, it will be calculated from the input values. And if `isinv` is omitted, it is taken to be `false`.

9.2 Access

```
p.nrows() = p.ncols() = p.colsize() = p.rowsize() = p.size()  
p(i, j)
```

Note: Because of the way the permutation is stored, `p(i, j)` is not terribly efficient. It takes $O(N)$ time to calculate. Also, there is no mutable version like there is for most matrices.

```
p.transpose() = p.inverse()
```

These are the same, and they do not create new storage. So statements like `v = p.transpose() * v` are efficient.

```
const int* p.getValues()
```

Get the indices of the interchanges. These are equivalent to the `pp` values described above for the constructor.

```
bool p.isInverse()
```

Returns true the the interchange values are taken in the reverse order (last to first) or false if not.

9.3 Functions

Most of these functions aren't very interesting, since most of them have trivial values like 1 or n . But we provide them all for consistency with the functions that other matrices provide.

```

int p.norm1() = Norm1(p)
int p.norm2() = Norm2(p) = p.doNorm2()
int p.normInf() = NormInf(p)
int p.maxAbsElement() = MaxAbsElement(p)
int p.maxAbs2Element() = MaxAbs2Element(p)
double p.normF() = NormF(p) = p.norm() = Norm(p)
int p.normSq() = NormSq(p)
int p.normSq(RT scale)
int p.trace() = Trace(p)
int p.sumElements() = SumElements(p)
int p.sumAbsElements() = SumAbsElements(p)
int p.det() = Det(p)
int p.logDet(int* sign=0) = LogDet(p)
bool p.isSingular()
int p.condition()
int p.doCondition()
pinv = p.inverse() = Inverse(p)
p.makeInverse(Matrix<T>& minv)
p.makeInverseATA(Matrix<T>& cov)

p.setToIdentity()
p.transposeSelf()
p.invertSelf()
Swap(p1,p2)

```

9.4 Arithmetic

```

v2 = p * v1
v2 = v1 * p
v2 = v1 / p
v2 = v1 % p
v *= p
v /= p
v %= p
m2 = p * m1
m2 = m1 * p
m2 = m1 / p
m2 = m1 % p
m *= p
m /= p
m %= p
p1 == p2
p1 != p2

```

9.5 Input/Output

The simplest output is the usual:

```
os << p
```

where `os` is any `std::ostream`. The output format is the same as for a `Matrix`, including all the 0's. (See 3.8.)

There is also a compact format:

```
p.writeCompact(os)
```

which outputs in the format:

```
P n isinv ( pp[0] pp[1] ... pp[n-1] )
```

where `isinv` is a boolean (1 or 0), and the `pp` values are the indices of the interchanges. In other words, the permutation could be created as

```
Permutation p(n,pp,isinv);
```

The same (compact, that is) format can be read back in the usual two ways:

```
tmv::Permutation p(n);
```

```
is >> p;
```

```
std::auto_ptr<tmv::Permutation> pptr;
```

```
is >> pptr;
```

where the first gives an error if `p` is the wrong size and the second allocates a new `Permutation` that is the correct size.

10 Errors and Exceptions

There are two kinds of errors that the TMV library looks for. The first are coding errors. Some examples are:

- Trying to access elements outside the range of a `Vector` or `Matrix`.
- Trying to add to `Vectors` or `Matrixes` that are different sizes.
- Trying to multiply a `Matrix` by a `Vector` where the number of columns in the `Matrix` doesn't match the size of the `Vector`.
- Viewing a `Matrix` as a `HermMatrix` when the diagonal isn't real.
- Calling `m.lud()`, `m.qrd()`, etc. for a `Matrix` that does not have that decomposition set (and saved).

I check for all of these (and similar) errors using assert statements. If these asserts fail, it should mean that the programmer made a mistake in the code. (Unless I've made a mistake in the TMV code, of course.)

Once the code is working, you can make the code slightly faster by compiling with `-DNDEBUG`. I say slightly, since most of these checks are pretty innocuous. And most of the computing time is usually in the depths of the various algorithms, not in these $O(1)$ time checks of the dimensions and such. There are a few checks which take $O(N)$ time, but these are only done if the code is compiled with the `-DXTTEST` flag.

The other kind of error checked for by the code is where the data don't behave in the way the programmer expected. Here is a (complete) list of these errors:

- A singular matrix is encountered in a division routine that cannot handle it.
- An input file has the wrong format.
- A Cholesky decomposition is attempted for a hermitian matrix that isn't positive definite.

These errors are always checked for even if `-DNDEBUG` is used. That's because they are not problems in the code per se, but rather are problems with the data or files used by the code. So they could still happen even after the code has been thoroughly tested.

All errors in the TMV library are indicated by throwing an object of type `tmv::Error`. If you decide to catch it, you can determine what went wrong by printing it:

```
catch (tmv::Error& e) {
    std::cerr << e << std::endl;
}
```

If you catch the error by value (i.e. without the `&`), it will print out a single line description. If you catch it by reference (as above), it may print out more information about the problem.

Also, `tmv::Error` derives from `std::exception` and overrides the `what()` method, so any program that catches these will catch `tmv::Error` as well.

If you want to be more specific, there are a number of classes that derive from `Error`:

10.1 FailedAssert

`tmv::FailedAssert` indicates that one of the assert statements failed. Since these are coding errors, if you catch this one, you'll probably just want to print out the error and abort the program so you can fix the bug. In addition to printing the text of the assert statement that failed (if you catch by reference), it will also indicate the file and line number as normal assert macros do. Unfortunately, it gives the line number in the TMV code, rather than in your own code, but hopefully seeing which function in TMV found the problem will help you figure out which line in your own code was incorrect.

If you believe that the assert failed due to a bug in the TMV code rather than your own code, please post a bug report at <http://code.google.com/p/tmv-cpp/issues>.

10.2 Singular

`tmv::Singular` indicates that you tried to invert or divide by a matrix that is (numerically) singular. This may be useful to catch specifically, since you may want to do something different when you encounter a singular matrix. Note however, that this only detects exactly singular matrices. If a matrix is numerically close to singular, but no actual zeros are found, then no error will be thrown, and your results will just be unreliable.

10.3 ReadError

`tmv::ReadError` indicates that there was some problem reading in a matrix or vector from an `istream` input. If you catch this by reference and write it, it will give you a fairly specific description of what the problem was as well as writing the part of the matrix or vector that was read in successfully.

10.4 NonPosDef

`tmv::NonPosDef` indicates that you tried to do some operation that requires a matrix to be positive definite, and it turned out not to be positive definite. The most common example would be performing a Cholesky decomposition on a hermitian matrix. I suspect that this is the most useful exception to catch specifically, as opposed to just via `tmv::Error` base class.

For example, the fastest algorithm for determining whether a matrix is (at least numerically) positive definite is to try the Cholesky decomposition and catch this exception. To wit:

```
bool IsPosDef(const tmv::GenSymMatrix<T>& m)
{
    assert(m.isherm());
    try {
        HermMatrix<T> m2 = m;
        CH_Decompose(m2.view());
    }
    catch (tmv::NonPosDef) {
        return false;
    }
    return true;
}
```

Or you might want to use Cholesky for division when possible and Bunch-Kaufman otherwise:

```
try {
    m.divideUsing(tmv::CH);
    m.setDiv();
}
catch (tmv::NonPosDef) {
    m.divideUsing(tmv::LU);
    m.setDiv();
}
x = b/m;
```

Note, however, that the speed difference between the two algorithms is only about 20% or so for typical matrices. So if a significant fraction of your matrices are not positive definite, you are probably better off always using the LU algorithm. Code like that given above would probably be most useful when all of your matrices should be positive definite in exact arithmetic, but you want to guard against one failing the Cholesky decomposition due to round-off errors.

It is also worth mentioning that the routine `QR_Downdate` described in §11.4 below will also throw the exception `NonPosDef` when it fails.

10.5 Warnings

There are also a few things that are not really errors, but where the user might be doing something wrong, or where something unexpected happened, but TMV was able to handle it. Here is a complete list of these situations:

- The function `makeInverseATA()` was called for a matrix with more columns than rows. The result in this case is really $(AA^\dagger)^{-1}$ instead of $(A^\dagger A)^{-1}$.
- The function `norm2()` or `condition()` was called without the SV decomposition being previously calculated (or saved) for the matrix. This is often inefficient, but if it is really what you want, then `doNorm2()` or `doCondition()` will bypass the warning.
- The LAPACK function `dsteqr` (or `ssteqr`) had an error, and the function `dstedc` (or `sstedc`) was called instead. The `dsteqr` calculation was wasted, although I don't think it is knowable a priori that this might happen.
- A LAPACK function requested more workspace than was provided. This shouldn't happen anymore, since I switched to using the LAPACK workspace queries. But the code to check for this is still there.

The default way of handling these situations is to print a message to `std::cout`. This is often appropriate, since it is probably something the programmer wants to know about. The first two items, in fact, are probably coding errors.

However, this behavior is clearly not always acceptable. So we provide a function to turn off the warnings:

```
void tmv::NoWarnings();
```

If you want to keep the warnings active, but don't like them being written to `std::cout`, then you can use:

```
std::ostream* tmv::WriteWarningsTo(std::ostream* newos)
```

This lets you redirect the warnings to some other `ostream`. It could be some log file that you look at later. Or it could even be a `stringstream` so you can examine the text of the warnings within your code and take appropriate action.

The function `WriteWarningsTo` returns a pointer to the old warning stream in case you only want to change the warning output stream temporarily. Also if you give `WriteWarningsTo` a null pointer, rather than an actual `ostream`, then this will turn off the warnings. In fact the above `NoWarnings` function is equivalent to `(void) WriteWarningsTo(0)`. These two features in combination let you turn off the warnings temporarily. For example:

```
tmv::Matrix<double> m(5,20);
// m = ....
tmv::Matrix<double> inv_mmt(5,5);
std::ostream* tempos = tmv::WriteWarningsTo(0);
m.makeInverseATA(int_mmt); // actually calculates (m*mt)^-1
tmv::WriteWarnignsTo(tempos);
```

11 Advanced Usage

11.1 Eigenvalues and eigenvectors

The eigenvalues of a matrix are important quantities in many matrix applications. A number, λ , is an eigenvalue of a square matrix, A , if for some non-zero vector v ,

$$Av = \lambda v$$

in which case v is called an eigenvector corresponding to the eigenvalue λ . Since any arbitrary multiple of v also satisfies this equation, it is common practice to scale the eigenvectors so that $\|v\|_2 = 1$. If v_1 and v_2 are eigenvectors whose eigenvalues are $\lambda_1 \neq \lambda_2$, then v_1 and v_2 are linearly independent.

The above equation implies that

$$\begin{aligned} Av - \lambda v &= 0 \\ (A - \lambda I)v &= 0 \\ \det(A - \lambda I) &= 0 \quad (\text{or } v = 0) \end{aligned}$$

If A is an $N \times N$ matrix, then the last expression is called the characteristic equation of A and the left hand side is a polynomial of degree N . Thus, it has potentially N solutions. Note that, even for real matrices, the solution may yield complex eigenvalues, in which case, the corresponding eigenvectors will also be complex.

If there are solutions to the characteristic equation which are multiple roots, then these eigenvalues are said to have a multiplicity greater than 1. These eigenvalues may have multiple corresponding eigenvectors. That is, different values of v (which are not just a multiple of each other) may satisfy the equation $Av = \lambda v$.

The number of independent eigenvectors corresponding to an eigenvalue with multiplicity > 1 may be less than that multiplicity⁸. Such eigenvalues are called “defective”, and any matrix with defective eigenvalues is likewise called defective.

If 0 is an eigenvalue, then the matrix A is singular. And conversely, singular matrices necessarily have 0 as one of their eigenvalues.

If we define Λ to be a diagonal matrix with the values of λ along the diagonal, then we have (for non-defective matrices)

$$AV = V\Lambda$$

where the columns of V are the eigenvectors. If A is defective, we can construct a V that satisfies this equation too, but some of the columns will have to be all zeros. There will be one such column for each missing eigenvector, and the other columns will be the eigenvectors.

If A is not defective, then all of the columns of V are linearly independent, which implies that V is not singular (i.e. V is “invertible”). Then,

$$\begin{aligned} A &= V\Lambda V^{-1} \\ \Lambda &= V^{-1}AV \end{aligned}$$

This is known as “diagonalizing” the matrix A . The determinant and trace are preserved by this procedure, which implies two more properties of eigenvalues:

$$\begin{aligned} \det(A) &= \prod_{k=1}^N \lambda_k \\ \text{tr}(A) &= \sum_{k=1}^N \lambda_k \end{aligned}$$

⁸ The multiplicity of the eigenvalue is generally referred to as its algebraic multiplicity. The number of corresponding eigenvectors is referred to as its geometric multiplicity. So $1 \leq \text{geometric multiplicity} \leq \text{algebraic multiplicity}$.

If A is a “normal” matrix – which means that A commutes with its adjoint, $AA^\dagger = A^\dagger A$ – then the matrix V is unitary, and A cannot be defective. The most common example of a normal matrix is a hermitian matrix (where $A^\dagger = A$), which has the additional property that all of the eigenvalues are real⁹.

So far, the TMV library can only find the eigenvalues and eigenvectors of hermitian matrices. The routines to do so are

```
void Eigen(const GenSymMatrix<T>& A,
           const MatrixView<T>& V, const VectorView<RT>& lambda);

void Eigen(const GenSymBandMatrix<T>& A,
           const MatrixView<T>& V, const VectorView<RT>& lambda);
```

where $V.col(i)$ is the eigenvector corresponding to each eigenvalue $lambda(i)$. The original matrix A can be obtained from

```
A = V * DiagMatrixViewOf(lambda) * V.adjoint();
```

There are also routines which only find the eigenvalues, which are faster, since they do not perform the calculations to determine the eigenvectors:

```
void Eigen(const SymMatrixView<T>& A, const VectorView<RT>& lambda);

void Eigen(const GenSymBandMatrix<T>& A, const VectorView<RT>& lambda);
```

Note that the first one uses the input matrix A as workspace and destroys the input matrix in the process.

11.2 Matrix decompositions

While many matrix decompositions are primarily useful for performing matrix division (or least-squares pseudo-division), one sometimes wants to perform the decompositions for their own sake. It is possible to get at the underlying decomposition with the various divider accessor routines like `m.lud()`, `m.qrd()`, etc. However, this is somewhat roundabout, and at times inefficient. So we provide direct ways to perform all of the various matrix decompositions that are implemented by the TMV code.

In the routines below, the matrix being decomposed is input as A , and we list the routines for all of the allowed types for A for each kind of decomposition. If A is listed as a “Gen” type, such as `GenBandMatrix<T>`, then that means the input matrix is not changed by the decomposition routine. If A is listed as a “View” type, such as `BandMatrixView<T>`, then that means the input matrix is changed.

In some cases where A is a View type, one of the decomposition components is returned in the location of A , or some part of it, overwriting the input matrix. In these cases, there will be a line indicating this after the function (e.g. `L = A.lowerTri()`). In other cases, the input matrix is just used as workspace, and it is junk on output, (in which case, there is no such line following the function).

Sometimes, only certain parts of a decomposition are wanted. For example, you might want to know the singular values of a matrix, but not care about the U and V matrices. For cases such as this, there are versions of the decomposition routines which omit certain output parameters. These routines are generally faster than the versions which include all output parameters, since they can omit some of the calculations.

Finally, a word about the permutations. In TMV, permutations are defined as a series of row or column swaps. I haven’t made a `Permutation` class yet to make it easy to use these permutations. But the code snippets which show how to recreate the input matrices from the decompositions should be sufficient to describe how to use the permutations as given.

- LU Decomposition

⁹ Other examples of normal matrices are unitary matrices ($A^\dagger A = AA^\dagger = I$) and skew-hermitian matrices ($A^\dagger = -A$). However, normal matrices do not have to be one of these special types.

(Matrix, BandMatrix)

$A \rightarrow PLU$ where L is lower triangular, U is upper triangular, and P is a permutation.

```
void LU_Decompose(const MatrixView<T>& A, Permutation& P);  
L = A.unitLowerTri();  
U = A.upperTri();
```

```
void LU_Decompose(const GenBandMatrix<T>& A,  
                  const LowerTriMatrixView<T>& L,  
                  const BandMatrixView<T>& U, Permutation& P);
```

In the second case, U must have $U.nhi() = A.nlo() + A.nhi()$, and L should be UnitDiag. In both cases, P must have $A.ncols()$ elements of memory allocated.

The original matrix A can be obtained from:

```
A = P * L * U;
```

- Cholesky Decomposition

(HermMatrix, HermBandMatrix)

$A \rightarrow LL^\dagger$, where L is lower triangular, and A is hermitian.

```
void CH_Decompose(const SymMatrixView<T>& A);  
L = A.lowerTri();
```

```
void CH_Decompose(const SymBandMatrixView<T>& A);  
L = A.lowerBand();
```

The original matrix A is very simply

```
A = L * L.adjoint();
```

- Bunch-Kaufman Decomposition

(HermMatrix, SymMatrix)

If A is hermitian, $A \rightarrow PLDL^\dagger P^T$, and if A is symmetric, $A \rightarrow PLDL^T P^T$, where P is a permutation, L is lower triangular, and D is hermitian or symmetric tridiagonal (respectively). In fact, D is even more special than that: it is block diagonal with 1×1 and 2×2 blocks, which means that there are no two consecutive non-zero elements along the off-diagonal.

```
void LDL_Decompose(const SymMatrixView<T>& A,  
                   const SymBandMatrixView<T>& D, Permutation& P);  
L = A.unitLowerTri();
```

P must have $A.size()$ elements of memory allocated.

The original matrix A can be obtained from:

```
A = P * L * D * (A.isherm() ? L.adjoint() : L.transpose())  
                  P.transpose();
```

- Tridiagonal LDL^\dagger Decomposition

(HermBandMatrix, SymBandMatrix with `nlo = 1`)

This decomposition for symmetric or hermitian tri-diagonal matrices is similar to the Bunch-Kaufman decomposition: $A \rightarrow LDL^\dagger$ or $A \rightarrow LDL^T$ where this time D is a regular diagonal matrix and L is a lower band matrix with a single subdiagonal and all 1's on the diagonal.

It turns out that the Bunch-Kaufman algorithm on banded matrices tends to expand the band structure without limit because of the pivoting involved, so it is not practical. However, with tridiagonal matrices, it is often possible to perform the decomposition without pivoting. Thus, this one does not have any growth of the band structure, but it is not as stable for singular or nearly singular matrices. If an exact zero is found on the diagonal along the way `tmv::NonPosDef` is thrown.¹⁰

```
void LDL_Decompose(const SymBandMatrixView<T>& A);
L = A.lowerBand();
L.diag().setAllTo(T(1));
D = DiagMatrixViewOf(A.diag());
```

The original matrix A can be obtained from:

```
A = L * D * (A.isherm() ? L.adjoint() : L.transpose());
```

- QR Decomposition

(Matrix, BandMatrix)

$A \rightarrow QR$, where Q is column-unitary (i.e. $Q^\dagger Q = I$), R is upper triangular, and A is either square or has more rows than columns.

```
void QR_Decompose(const MatrixView<T>& A,
                  const UpperTriMatrixView<T>& R);
Q = A;

void QR_Decompose(const GenBandMatrix<T>& A,
                  const MatrixView<T>& Q, const BandMatrixView<T>& R);
```

In the second case, R must have $R.nhi() \geq A.nlo() + A.nhi()$.

If you only need R , the following versions are faster, since they do not fully calculate Q .

```
void QR_Decompose(const MatrixView<T>& A);
R = A.upperTri();

void QR_Decompose(const GenBandMatrix<T>& A,
                  const BandMatrixView<T>& R);
```

- QRP Decomposition

(Matrix)

$A \rightarrow QRP$, where Q is column-unitary (i.e. $Q^\dagger Q = I$), R is upper triangular, P is a permutation, and A is either square or has more rows than columns.

```
void QRP_Decompose(const MatrixView<T>& A,
                  const UpperTriMatrixView<T>& R, Permutation& P,
                  bool strict=false);
Q = A;
```

¹⁰ Note, however, that if A is complex, symmetric - i.e. not hermitian - then this doesn't actually mean that A is not positive definite (since such a quality is only defined for hermitian matrices). Furthermore, hermitian matrices that are not positive definite will probably be decomposed successfully without throwing, resulting in D having negative values.

Also, the LAPACK implementation throws an exception for matrices that the native code successfully decomposes. It throws for hermitian matrices whenever they are not positive definite, whereas the native code succeeds for many indefinite matrices.

P must have `A.ncols()` elements of memory allocated.

As discussed in §3.7.3, there are two slightly different algorithms for doing a QRP decomposition. If `strict = true`, then the diagonal elements of R be strictly decreasing (in absolute value) from upper-left to lower-right¹¹.

If `strict` is false however (the default), then the diagonal elements of R will not be strictly decreasing. Rather, there will be no diagonal element of R below and to the right of one which is more than a factor of $\epsilon^{1/4}$ smaller in absolute value, where ϵ is the machine precision. This restriction is almost always sufficient to make the decomposition useful for singular or nearly singular matrices, and it is much faster than the strict algorithm.

The original matrix A is obtained from:

```
A = Q * R * P;
```

If you only need R , the following versions is faster, since it does not fully calculate Q .

```
void QRP_Decompose(const MatrixView<T>& A, bool strict=false);
R = A.upperTri();
```

- Singular Value Decomposition

(Matrix, SymMatrix, HermMatrix, BandMatrix, SymBandMatrix, HermBandMatrix)

$A \rightarrow USV$, where U is column-unitary (i.e. $U^\dagger U = I$), S is real diagonal, V is square unitary, and A is either square or has more rows than columns.¹²

```
void SV_Decompose(const MatrixView<T>& A,
                  const DiagMatrixView<RT>& S, const MatrixView<T>& V);
U = A;

void SV_Decompose(const GenSymMatrix<T>& A,
                  const MatrixView<T>& U, const DiagMatrixView<RT>& S,
                  const MatrixView<T>& V);

void SV_Decompose(const GenBandMatrix<T>& A,
                  const MatrixView<T>& U, const DiagMatrixView<RT>& S,
                  const MatrixView<T>& V);

void SV_Decompose(const GenSymBandMatrix<T>& A,
                  const MatrixView<T>& U, const DiagMatrixView<RT>& S,
                  const MatrixView<T>& V);
```

The input A matrix must not have more columns than rows. If you want to calculate the SVD of such a matrix, you should decompose A^T instead:

¹¹ If you are using a LAPACK library, you might find that the output R diagonal is not always strictly decreasing, although it will usually be close. If strictly monotonic diagonal elements are important for you, you can use the native TMV algorithm instead by compiling with the flag `-DNOGEQP3`.

¹² The singular value decomposition is more commonly written as $A \rightarrow USV^T$. As far as I can tell, this seems to be a holdover from the days of Fortran programming. In Fortran, matrices are stored in column-major format. Considering that the rows of what we call V are the singular vectors, also known as principal components, of A , it made more sense for Fortran programmers to use the transpose of V which has the principal components in the columns. This complication is unnecessary in TMV. If you want the principal components stored contiguously, just make V row-major. On the other hand, decomposition with column-major storage of V is usually a bit faster, so you need to make a choice appropriate for your particular program.

```

tmv::Matrix<double> A(nrows,ncols); // ncols > nrows
[ A = ... ]
tmv::Matrix<double> V = A;
tmv::DiagMatrix<double> S(nrows);
tmv::Matrix<double> U(nrows,nrows);
SV_Decompose(V.transpose(),S.view(),U.transpose());
// Now A = U * S * V

```

If you only need S , or S and V , or S and U , the following versions are faster, since they do not fully calculate the omitted matrices.

```

void SV_Decompose(const MatrixView<T>& A,
                  const DiagMatrixView<RT>& S, const MatrixView<T>& V, false);
// U != A

void SV_Decompose(const MatrixView<T>& A,
                  const DiagMatrixView<RT>& S, bool StoreU);
if (StoreU) U = A;

void SV_Decompose(const SymMatrixView<T>& A,
                  const DiagMatrixView<RT>& S);

void SV_Decompose(const GenSymMatrix<T>& A,
                  const DiagMatrixView<RT>& S, const MatrixView<T>& V);

void SV_Decompose(const GenSymMatrix<T>& A,
                  const MatrixView<T>& U, const DiagMatrixView<RT>& S);

void SV_Decompose(const GenBandMatrix<T>& A,
                  const DiagMatrixView<RT>& S);

void SV_Decompose(const GenBandMatrix<T>& A,
                  const DiagMatrixView<RT>& S, const MatrixView<T>& V);

void SV_Decompose(const GenBandMatrix<T>& A,
                  const MatrixView<T>& U, const DiagMatrixView<RT>& S);

void SV_Decompose(const GenSymBandMatrix<T>& A,
                  const DiagMatrixView<RT>& S);

void SV_Decompose(const GenSymBandMatrix<T>& A,
                  const DiagMatrixView<RT>& S, const MatrixView<T>& V);

void SV_Decompose(const GenSymBandMatrix<T>& A,
                  const MatrixView<T>& U, const DiagMatrixView<RT>& S);

```

- Polar Decomposition

(Matrix, BandMatrix)

$A \rightarrow UP$ where U is unitary and P is positive definite hermitian.

This is similar to polar form of a complex number: $z = re^{i\theta}$. In the matrix version, P acts as r , being in some sense the “magnitude” of the matrix. And U acts as $e^{i\theta}$, being a generalized rotation.

```
void Polar_Decompose(const MatrixView<T>& A,
    const SymMatrixView<T>& P);
U = A;

void Polar_Decompose(const GenBandMatrix<T>& A,
    const MatrixView<T>& U, const SymMatrixView<T>& P);
```

- Matrix Square Root

(HermMatrix, HermBandMatrix)

$A \rightarrow SS$, where A and S are each positive definite hermitian matrices.

```
void SquareRoot(const SymMatrixView<T>& A);
S = A;

void SquareRoot(const GenSymBandMatrix<T>& A,
    const SymMatrixView<T>& S);
```

If A is found to be not positive definite, a NonPosDef exception is thrown.

11.3 Update a QR decomposition

One reason that it can be useful to create and deal with the QR decomposition directly, rather than just relying on the division routines is the possibility of updating or “downdating” the resulting R matrix.

If you are doing a least-square fit to a large number of linear equations, you can write the system as a matrix equation: $Ax = b$, where A is a matrix with more rows than columns, and you are seeking, not an exact solution for x , but rather the value of x which minimizes $\|b - Ax\|_2$. See §3.7.2 for a more in depth discussion of this topic.

It may be the case that you have more rows (i.e. constraints) than would allow the entire matrix to fit in memory. In this case it may be tempting to use the so-called normal equation instead:

$$A^\dagger Ax = A^\dagger b$$

$$x = (A^\dagger A)^{-1} A^\dagger b$$

This equation theoretically gives the same solution as using the QR decomposition on the original design matrix. However, it can be shown that the condition of $A^\dagger A$ is the square of the condition of A . Since larger condition values lead to larger numerical instabilities and round-off problems, a mildly ill-conditioned matrix is made much worse by this procedure.

When all of A fits in memory, the better solution is to use the QR decomposition, $A = QR$, to calculate x .

$$QRx = b$$

$$x = R^{-1}Q^\dagger b$$

In fact, this is the usual behind-the-scenes procedure when you write $x = b/A$ in TMV. But if A is too large to fit in memory, then so is Q .

A compromise solution, which is not quite as good as doing the full QR decomposition, but is better than using the normal equation, is to just calculate the R of the QR decomposition, and not Q . Then:

$$A^\dagger Ax = A^\dagger b$$

$$R^\dagger Q^\dagger QRx = R^\dagger Rx = A^\dagger b$$

$$x = R^{-1}(R^\dagger)^{-1} A^\dagger b$$

Calculating R directly from A is numerically much more stable than calculating it through, say, a Cholesky decomposition of $A^\dagger A$. So this method produces a more accurate answer for x than the normal equation does.

But how can R be calculated if we cannot fit all of A into memory at once?

First, we point out a characteristic of unitary matrices that the product of two or more of them is also unitary. This implies that if we can calculate something like: $A = Q_0 Q_1 Q_2 \dots Q_n R$, then this is the R that we want.

So, consider breaking A into a submatrix, A_0 , which can fit into memory, plus the remainder, A_1 .

$$A = \begin{pmatrix} A_0 \\ A_1 \end{pmatrix}$$

First perform a QR decomposition of $A_0 = Q_0 R_0$. Then we have:

$$\begin{aligned} A &= \begin{pmatrix} Q_0 R_0 \\ A_1 \end{pmatrix} \\ &= \begin{pmatrix} Q_0 & 0 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} R_0 \\ A_1 \end{pmatrix} \\ &\equiv Q'_0 A'_1 \end{aligned}$$

Assuming that A_0 has more rows than columns, then A'_1 has fewer rows than the original matrix A . So we can iterate this process until the resulting matrix can fit in memory, and we can perform the final QR update to get the final value of R .

For the numerical reasons mentioned above, the fewer such iterations you do, the better. So you should try to include as many rows of the matrix A as possible in each step, given the amount of memory available.

The solution equation, written above, also needs the quantity $A^\dagger b$, which can be accumulated in the same blocks:

$$A^\dagger b = A_1^\dagger b_1 + A_2^\dagger b_2 + \dots$$

This, combined with the calculation of R , allows us to determine x using the above formula.

The TMV library includes a command which does the update step of the above procedure directly, which is slightly more efficient than explicitly forming the A'_k matrices. The commands is

```
void QR_Update(const UpperTriMatrixView<T>& R, const MatrixView<T>& M)
```

which updates the value of R such that $R_{\text{out}}^\dagger R_{\text{out}} = R_{\text{in}}^\dagger R_{\text{in}} + X^\dagger X$. (The input matrix X is destroyed in the process.) This is equivalent to the QR definition of the update described above.

So the entire process might be coded using TMV as:

```
int n_full = nrows_for_full_A_matrix;
int n_mem = nrows_that_fit_in_memory;
assert(n_mem <= n_full);
assert(n_mem > ncols);

tmv::Matrix<double> A(n_mem, ncols);
tmv::Vector<double> b(n_mem);

// Import_Ab sets A to the first n_mem rows of the full matrix,
// and also sets b to the same components of the full rhs vector.
// Maybe it reads from a file, or performs a calculation, etc.
Import_Ab(0, n_mem, A, b);

// x will be the solution to A_full x = b_full when we are done
// But for now, it is accumulating A_full.transpose() * b_full.
tmv::Vector<double> x = A.transpose() * b;
```

```

// Do the initial QR Decomposition:
QR_Decompose(A.view());
tmv::UpperTriMatrix<double> R = A.upperTri();

// Iterate until we have done all the rows
int n1 = 0, n2 = n_mem;
while (n2 < n_full)
{
    n1 = n2; n2 += n_mem;
    if (n2 > n_full) n2 = n_full;

    // Import the next bit:
    Import_Ab(n1, n2, A, b);

    // (Usually, A1==A, b1==b, but not the last time through the loop.)
    tmv::MatrixView<double> A1 = A.rowRange(0, n2-n1);
    tmv::VectorView<double> b1 = b.subVector(0, n2-n1);

    // Update, x, R:
    x += A1.transpose() * b1;
    QR_Update(R.view(), A1);
}

// Finish the solution:
x /= R.transpose();
x /= R;

```

11.4 Downdate a QR decomposition

When performing a least-square fit of some data to a model, it is common to do some kind of outlier rejection to remove data that seem not to be applicable to the model - things like spurious measurements and such. For this, we basically want the opposite of a QR update - instead we want to find the QR decomposition that results from removing a few rows from A . This is called a QR “downdate”, and is performed using the subroutine:

```
void QR_Downdate(const UpperTriMatrixView<T>& R, const GenMatrix<T>& X)
```

X represents the rows from the original matrix to remove from the QR decomposition.

It is possible for the downdate to fail (and throw an exception) if the matrix X does not represent rows of the matrix that was originally used to create R . Furthermore, with round-off errors, the error may still result with actual rows from the original A if R gets too close to singular. In this case, `QR_Downdate` throws a `NonPosDef` exception. This might seem like a strange choice, but the logic is that $R^\dagger R$ is the Cholesky decomposition of $A^\dagger A$, and `QR_Downdate(R, X)` basically updates R to be the Cholesky decomposition of $A^\dagger A - X^\dagger X$. The procedure fails (and throws) when this latter matrix is found not to be positive definite.

It is worth pointing out that the algorithm used in TMV is a new one developed by Mike Jarvis. Most of the texts and online resources that discuss the QR downdate algorithm only explain how to do one row at a time, using a modification of the QR update using Givens rotations. If you are doing many rows, it is common that roundoff errors in such a procedure accumulate sufficiently for the routine to fail. The TMV algorithm instead downdates all of the rows together using a modification of the Householder reflection algorithm for updates. This algorithm seems to be much more stable than ones that use Givens rotations.

The only references to a similar algorithm that I could find in the literature is described in the paper, "Stability Analysis of a Householder-based Algorithm for Downdating the Cholesky Factorization", Bojanczyk and Steinhardt, 1991, Siam J. Sci. Stat. Comput. 12, 6, 1255¹³. This paper describes a similar algorithm to compute the downdated R matrix using Householder matrices. However, the details of the computation are somewhat different from the TMV algorithm. Also, they only consider real matrices, and they do not include the block-householder techniques in their description to employ more so-called "level-3" matrix operations.

Therefore, I will describe the TMV downdate algorithm here. I think it is clearer to begin by describing the update algorithm in §11.4.1, since it is quite similar to the algorithm we use for downdating, but is a bit easier to understand. Then the downdate algorithm is described in §11.4.2.

11.4.1 The update algorithm

First lets look at the Householder algorithm for QR update:

Given the initial decomposition $A_0 = Q_0 R_0$, we want to find R such that

$$\begin{aligned} A_1 &= \begin{pmatrix} A_0 \\ X \end{pmatrix} = Q_1 R_1 \\ &\begin{pmatrix} Q_0 R_0 \\ X \end{pmatrix} = Q_1 R_1 \\ \begin{pmatrix} Q_0 & 0 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} R_0 \\ X \end{pmatrix} &= Q_1 R_1 \end{aligned}$$

So if we perform a QR decomposition:

$$S \equiv \begin{pmatrix} R_0 \\ X \end{pmatrix} = Q_S R_S$$

Then this is the R we want: $R_1 = R_S$, and

$$Q_1 = \begin{pmatrix} Q_0 & 0 \\ 0 & 1 \end{pmatrix} Q_S$$

For the following discussion, let N be the number of rows (and columns) in R_0 , and M be the number of rows in X .

To perform the decomposition, we multiply S by a series of Householder reflections on the left to zero out each column of X one at a time. Householder reflections are unitary, so their product in the reverse order is Q_S :

$$\begin{aligned} \begin{pmatrix} R_1 \\ 0 \end{pmatrix} &= H_N H_{N-1} \dots H_2 H_1 \begin{pmatrix} R_0 \\ A_1 \end{pmatrix} \\ Q_S &= H_1^\dagger H_2^\dagger \dots H_{N-1}^\dagger H_N^\dagger \end{aligned}$$

Householder reflections are defined as $H = I - \beta(x - ye_1)(x - ye_1)^\dagger$ where x is a (column) vector, y is a scalar with $|y| = \|x\|_2$, e_1 is the basis vector whose only non-zero element is the first: $e_1(1) = 1$, and $\beta = (\|x\|_2^2 - y^* x(1))^{-1}$. They have the useful properties that $Hx = ye_1$ and they are unitary: $H^\dagger H = HH^\dagger = I$. Furthermore, if β is real, they are also hermitian: $H = H^\dagger$.

H_1 is defined for the vector $x = (R_0(1, 1), 0, 0, \dots, 0, 0, X(1, 1), X(2, 1), \dots, X(M, 1))$ where the stretch of 0's includes a total of $(N - 1)$ 0's. This value of x completely determines the Householder matrix H_1 up to an arbitrary sign on either y or β (or in general an arbitrary factor $e^{i\theta}$) which is chosen to minimize rounding errors.

¹³ It seems that this paper has become a bit forgotten. A recent paper, "Efficient Algorithms for Block Downdating of Least Squares Solutions", Yanev and Kontoghiorghes, 2004, Applied Numerical Mathematics, 49, 3, evaluates five algorithms for doing the downdate. However, all of them are block versions of the Given matrix approach. They do not consider any algorithms that use Householder matrices to do the downdate, and do not reference the above paper by Bojanczyk and Steinhardt. In addition, none of the papers that do cite the Bojanczyk and Steinhardt paper seem to be about the general problem of QR (or Cholesky) downdating.

The optimal choice is to choose $y = -||x||_2 x(1)/|x(1)|$, which makes β real. However, the LAPACK choice is $y = -||x||_2 \text{sign}(\text{real}(x(1)))$, which means β is complex, and H is not Hermitian¹⁴.

The product $H_1 S$ “reflects” the first column of X into the first diagonal element of R_0 . Because of all the 0’s, most of R_0 is unaffected – only the first row of R_0 and the rest of X are changed. The subsequent Householder reflections are defined similarly, each zeroing out a column of X , and modifying the corresponding row of R_0 and the remaining elements of X .

At the end of this procedure, the matrix R_0 will be changed into the matrix R_1 . If desired, Q_S (and then Q_1) may also be calculated in the process, but the TMV implementation of the QR update does not calculate Q_1 . If there is a demand for such a routine, it would not be hard to add it, but I think most applications of the update do not use the Q matrix explicitly.

11.4.2 The downdate algorithm

Given the initial decomposition

$$A_1 = \begin{pmatrix} A_0 \\ X \end{pmatrix} = Q_1 R_1$$

we want to find R_0 such that $A_0 = Q_0 R_0$.

The TMV algorithm to do this essentially performs the same steps as in the update algorithm above, but instead removes the effect of each H from R_1 . This is easy to do if we can determine what each H is, since $H^{-1} = H^\dagger$, so we just apply H^\dagger to update each row of R_1 . The X update takes the regular H matrix, since we need to replicate the steps that we would do for an update to keep finding the correct values for the remaining columns of X .

All of the values in the vector x needed to define H_1 are given, except for the first, $R_0(0, 0)$. But this is easy to calculate, since

$$|R_1(0, 0)|^2 = |R_0(0, 0)|^2 + ||X(1 : M, 0)||_2^2$$

This determines the x vector, which in turn defines H_1 (modulo an arbitrary sign, which again is chosen to minimize rounding errors). Thus, we can calculate H_1 and apply it as described above. Each subsequent Householder matrix is created and applied similarly for each column of X . When we have finished this process, we are left with R_0 in the place of R_1 .

If at any point in the process, we find the calculated $|R_0(k, k)|^2 < 0$, then the algorithm fails. In the TMV implementation, a `NonPosDef` exception is thrown.

In practice, for both of these algorithms, we actually use a blocked implementation for updating the R and X matrices. We accumulate the effect of the Householder matrices until there are sufficiently many (e.g. 64), at which point we update the appropriate rows of the R matrix and the rest of X . Implementing this correctly is mostly a matter of keeping track of which elements have been updated yet, making sure that whenever an element is used, it is already updated, while delaying as much of the calculation as possible in order to make maximum use of the so-called “level-3” matrix functions, which are the most efficient on modern computers. We also make the additional improvement of using a recursive algorithm within each block, which gains some additional level-3 operations, for a bit more efficiency.

11.5 Other SymMatrix operations

There are three more arithmetic routines that we provide for `SymMatrix`, which do not have any corresponding shorthand with the usual arithmetic operators.

The first two are:

¹⁴ This choice complicates a lot of the calling routines which use Householder matrices, since you need to keep track of conjugation of the β values. Since TMV is designed to be able to call LAPACK when possible, it is forced to follow the same convention.

In fact, it could be argued that the LAPACK convention is even “wrong” in the sense that their Householder matrices are not actually “reflections”. A reflection is a unitary matrix whose determinant is -1 . The determinant of a Householder matrix as defined here is $-\beta^2/|\beta|^2$ which is -1 for real β , but not for complex β . But we are stuck with their choice, so we allow β to be complex in this discussion.

```

tmv::Rank2Update<bool add>(T x, const GenVector<T1>& v1,
    const GenVector<T2>& v2, const SymMatrixView<T>& s)
tmv::Rank2KUpdate<bool add>(T x, const GenMatrix<T1>& m1,
    const GenMatrix<T2>& m2, const SymMatrixView<T>& s)

```

They are similar to the Rank1Update and RankKUpdate routines, which are implemented in TMV with the expressions $s += x * v \wedge v$ and $s += x * m * m.transpose()$.

A rank-2 update calculates

```

s (+=) x * ((v1 ^ v2) + (v2 ^ v1))
s (+=) x * (v1 ^ v2.conjugate()) + conj(x) * (v2 ^ v1.conjugate())

```

for a symmetric or hermitian s respectively, where “(+=)” means “+=” if `add` is `true` and “=” if `add` is `false`. Likewise, a rank-2k update calculates:

```

s (+=) x * (m1 * m2.transpose() + m2 * m1.transpose())
s (+=) x * m1 * m2.adjoint() + conj(x) * m2 * m1.adjoint()

```

for a symmetric or hermitian s respectively.

We don’t have an arithmetic operator shorthand for these, because, as you can see, the operator overloading required would be quite complicated. And since they are pretty rare, I decided to just let the programmer call the routines explicitly.

The other routine is:

```

tmv::SymMultMM<bool add>(T x, const GenMatrix<T>& m1,
    const GenMatrix<T>& m2, const SymMatrixView<T>& s)

```

This calculates the usual generalized matrix product: $s (+=) x * m1 * m2$, but it basically asserts that the product $m1 * m2$ is symmetric (or hermitian as appropriate).

Since a matrix product is not in general symmetric, I decided not to allow this operation with just the usual operators to prevent the user from doing this accidentally. However, there are times when the programmer can know that the product should be (at least numerically close to) symmetric and that this calculation is ok. Therefore it is provided as a subroutine.

11.6 Element-by-element product

The two usual kinds of multiplication for vectors are the inner product and the outer product, which result in a scalar and a matrix respectively. However there is also a third kind of multiplication that is sometimes needed where each element in a vector is multiplied by the corresponding element in another vector: $v(i) = v(i) * w(i)$.

There are two functions that should provide all of this kind of functionality for you:

```

ElementProd(T x, const GenVector<T1>& v1, const VectorView<T>& v2);
AddElementProd(T x, const GenVector<T1>& v1, const GenVector<T2>& v2,
    const VectorView<T>& v3)

```

The first performs $v_2(i) = x * v_1(i) * v_2(i)$, and the second performs $v_3(i) = v_3(i) + x * v_1(i) * v_2(i)$ for $i = 0 \dots (N - 1)$ (where N is the size of the vectors).

There is no operator overloading for Vectors that would be equivalent to these expressions. But they are actually equivalent to the following:

```

v2 *= x * DiagMatrixViewOf(v1);
v3 += x * DiagMatrixViewOf(v1) * v2;

```

respectively. In fact, these statements inline to the above function calls automatically. Depending on you preference and the meanings of your vectors, these statements may or may not be clearer as to what you are doing.

There are also corresponding functions for Matrix and for each of the special matrix types:

```

ElementProd(T x, const GenMatrix<T1>& m1, const MatrixView<T>& m2);
AddElementProd(T x, const GenMatrix<T1>& m1, const GenMatrix<T2>& m2,
               const MatrixView<T>& m3);

```

Likewise for the other special matrix classes. The first performs $m_2(i, j) = x * m_1(i, j) * m_2(i, j)$, and the second performs $m_3(i, j) = m_3(i, j) + x * m_1(i, j) * m_2(i, j)$ for every i, j in the matrix.

These don't have any `DiagMatrixViewOf` version, since the corresponding concept would require a four-dimensional tensor, and the TMV library just deals with one- and two-dimensional objects.

The matrices all have to be the same size and shape, but can have any (i.e. not necessarily the same) storage method. However, the routines are faster if the matrices use the same storage.

11.7 BaseMatrix views

If you are dealing with objects that are only known to be `BaseMatrix`s (i.e. they could be a `Matrix` or a `DiagMatrix` or a `SymMatrix`, etc.), then methods like `m.transpose()`, `m.view()`, and such can't know what kind of object to return. So these methods can't be defined for a `BaseMatrix`.

Instead, we have the following virtual methods, which are available to a `BaseMatrix` object and are defined in each specific kind of matrix to return a pointer to the right kind of object:

```

std::auto_ptr<tmv::BaseMatrix<T> > m.newCopy()
std::auto_ptr<tmv::BaseMatrix<T> > m.newView()
std::auto_ptr<tmv::BaseMatrix<T> > m.newTranspose()
std::auto_ptr<tmv::BaseMatrix<T> > m.newConjugate()
std::auto_ptr<tmv::BaseMatrix<T> > m.newAdjoint()
std::auto_ptr<tmv::BaseMatrix<T> > m.newInverse()

```

`newCopy` and `newInverse` create new storage to store a copy of the matrix or its inverse, respectively. The other four just return views of the current matrix.

11.8 Iterators

We mentioned that the iterators through a `Vector` are:

```

typename tmv::Vector<T>::iterator
typename tmv::Vector<T>::const_iterator
typename tmv::Vector<T>::reverse_iterator
typename tmv::Vector<T>::const_reverse_iterator

```

just like for standard library containers. The specific types to which these typedefs refer are:

```

tmv::VIt<T, tmv::Unit, tmv::NonConj>
tmv::CVIt<T, tmv::Unit, tmv::NonConj>
tmv::VIt<T, tmv::Step, tmv::NonConj>
tmv::CVIt<T, tmv::Step, tmv::NonConj>

```

respectively.

`VIt` is a mutable-iterator, and `CVIt` is a const-iterator. `Unit` indicates that the step size is 1, while `Step` allows for any step size between successive elements (and is therefore slower). For the reverse iterators, the step size is -1.

This can be worth knowing if you are going to be optimizing code that uses iterators of `VectorViews`. This is because their iterators are instead:

```

tmv::VIter<T>
tmv::CVIter<T>

```

which always check the step size (rather than assuming unit steps) and always keep track of a possible conjugation.

If you know that you are dealing with a view that is not conjugated, you can convert your iterator into one of the above `VIt` or `CVIt` types, which will be faster, since they won't check the conjugation bit each time.

Likewise, if you know that it *is* conjugated, then you can use `tmv::Conj` for the third template parameter above. This indicates that the vector view really refers to the conjugates of the values stored in the actual memory locations.

Also, if you know that your view has unit steps between elements, converting to an iterator with `tmv::Unit` will iterate faster. It is often faster to check the step size once at the beginning of the routine and convert to a unit-step iterator if possible.

All of these conversions can be done with a simple cast or constructor, such as:

```
if (v.step() == 1) {
    for(VIt<float,Unit,NonConj> it = v.begin(); it != v.end(); it++)
        (*it) = sqrt(*it);
} else {
    for(VIt<float,Step,NonConj> it = v.begin(); it != v.end(); it++)
        (*it) = sqrt(*it);
}
```

Regular Vectors are always `Unit` and `NonConj`, so those iterators are already fast without using the specific `VIt` names. That is, you can just use `Vector<T>::iterator` rather than `VIt<T,Unit,NonConj>` without any performance drop.

11.9 Direct memory access

We provide methods for accessing the memory of a matrix or vector directly. This is especially useful for meshing the TMV objects with other libraries (such as BLAS or LAPACK). But it can also be useful for writing some optimized code for a particular function.

The pointer to the start of the memory for a vector can be obtained by:

```
T* v.ptr()
const T* v.cptr() const
```

Using the direct memory access requires that you know the spacing of the elements in memory and (for views) whether the view is conjugated or not. So we also provide:

```
int v.step() const
bool v.isconj() const
```

For matrices, the corresponding routines return the upper-left element of the matrix. Note that for some matrices, (e.g. `BandMatrix<T,DiagMatrix>`) this is not necessarily the first element in memory. We also need to know the step size in both directions:

```
T* m.ptr()
const T* m.cptr() const
int m.stepi() const
int m.stepj() const
bool m.isconj() const
bool m.isrm() const
bool m.iscm() const
```

The step in the “down” direction along a column is `stepi`, and the step to the “right” along a row is `stepj`. The last two check if a matrix is `RowMajor` or `ColMajor` respectively.

For band matrices, there are also:

```
int m.diagstep() const
bool m.isdm() const
```

which return the step along the diagonal and whether the matrix is `DiagMajor`.

For symmetric/hermitian matrices, there are some more methods:

```
bool m.isherm()
bool m.issym()
bool m.isupper()
```

The first two both return `true` for real symmetric matrices, but differentiate between hermitian and symmetric varieties for complex types. The last one tells you whether the actual elements to be accessed are stored in the upper triangle half of the matrix (true) or the lower (false).

11.10 “Linear” views

Our matrices generally store the data contiguously in memory with all of the methods like `row` and `col` returning the appropriate slice through the data. Occasionally, though, it can be useful to treat the whole matrix as a single vector of elements. We use this internally for implementing routines like `setAllTo` and matrix addition, among others. These are faster than accessing the data in ways that use the actual matrix structure.

This kind of access may be useful for some users of the library, so the following methods are available:

```
tmv::VectorView<T> m.linearView()
tmv::ConstVectorView<T> m.constLinearView()
bool m.canLinearize()
```

These return a view to the elements of a `Matrix` as a single vector. It is always allowed for an actual `Matrix`. For a `MatrixView` (or `ConstMatrixView`), it is only allowed if all of the elements in the view are in one contiguous block of memory. The helper function `m.canLinearize()` returns whether or not the first two methods are legal.

The same methods are also defined for `BandMatrix` (and corresponding views). In this case, there are a few elements in memory that are not necessarily defined, since they lie outside of the actual band structure, so some care should be used depending on the application of the returned vector views. (For example, one cannot compute things like the minimum or maximum element this way, since the undefined elements may have very large or small values which would corrupt this calculation.)

The triangular and symmetric matrices have too much memory that is not actually used by the matrix for these to be very useful, so we do not provide them. When we eventually implement the packed storage varieties, these methods will be provided for those.

Along the same lines is another method for a `Vector`:

```
tmv::VectorView<RT> v.flatten()
```

This returns a real view to the real and imaginary elements of a complex `Vector`. The initial `Vector` is required to have unit step. The returned view has twice the length of `v` and also has unit step.

This probably isn’t very useful for most users either, but it is useful internally, since it allows code such as:

```
tmv::Vector<complex<double> > v(500);
[...]
v *= 2.3;
```

to call the BLAS routine `dscal` with `x=2.3`, rather than `zscal` with `x=complex<double>(2.3,0.0)`, which would be slower.

11.11 Getting the Version of TMV

At times it can be useful to be able to access what version of TMV is installed on a particular machine, either to log the information as part of the meta-data about a run of a program, or even to modify the code depending on which features of TMV are available. To address this, we provide three ways to access this information.

First, there is a function you can call from within your program:

```
std::string TMV_Version();
```

For this release, this function returns the string “0.64”. This is useful for inserting into a log file or anywhere else that you want to record the version somewhere.

Second we also provide three C preprocessor definitions:

```
TMV_MAJOR_VERSION
TMV_MINOR_VERSION
TMV_VERSION_AT_LEAST(major, minor)
```

The first two are defined to be 0 and 64 for this release. The third can be used with an `#if` directive to change what code you want to compile according to the version of TMV that is installed. For example, in version 0.64 we added the `m.unitUpperTri()` feature. So you could write¹⁵:

```
#if TMV_VERSION_AT_LEAST(0, 64)
U = m.unitUpperTri();
#else
U = m.upperTri(tmv::UnitDiag);
#endif
```

And finally, we also include a bash script called `tmv-version` that will output the version number, so you can run this directly from the command line. For this release, this script produces the output:

```
$ tmv-version
0.64
$
```

This can be useful as part of a larger script if you want to log the TMV version from that rather than from within a C++ program.

Hopefully these three methods will satisfy any potential manner in which you might want to access the version of TMV that is on your system. If there is something I missed, and none of these work for you, please let me know, and I’ll probably be happy to provide another mechanism in future releases.

¹⁵ Well, technically, this wouldn’t be very useful, since the `TMV_VERSION_AT_LEAST` macro wasn’t introduced until version 0.64 either...

12 Obtaining and Compiling the Library

This code is licensed using the Gnu General Public License. See §16 below for more details.

12.1 Obtaining the source code

1. Go to <http://code.google.com/p/tmv-cpp/> for a link to a tarball with all of the source code (on the right hand side as a “Featured Download”, or also in the “Downloads” tab), and copy it to the directory where you want to put the TMV library.¹⁶

2. Unpack the tarball:

```
gunzip tmv0.64.tar.gz
tar xf tmv0.64.tar
```

This will make a directory called `tmv0.64` with the subdirectories: `doc`, `examples`, `include`, `lib`, `src` and `test` along with the files `README`, `INSTALL` and others in the top directory.

At this point there are several different choices for how to compile the library. The preferred method is to use SCons. To use this method, continue onto §12.2. However, we also provide instructions for using standard make (§12.3), CMake (§12.5), or Visual C++ (§12.6).

12.2 Installing using SCons

With the version 0.62 release of TMV, I introduced a new installation method using SCons. For that release it was considered an alternative method, since it was brand new. However, by now, I have tried it on quite a few platforms and it seems to be working very well. It usually finds the BLAS and LAPACK libraries on your system automatically, which is nicer than having to specify them by hand.

If you have trouble installing with SCons, you can try the Makefile installation instead (§12.3). And please let me know about the problem you are having either at <mailto:tmv-discuss@googlegroups.com> or <http://code.google.com/p/tmv-cpp/issues>.

1. Download and unpack the tarball as described above (§12.1).
2. Make sure you have SCons installed on your system. It is available for free from <http://www.scons.org/>. (It is a very quick installation if you have Python installed.)

3. Type
`scons`

This will make the libraries `libtmv.a` and `libtmv_symband.a` and put them into the directory `lib`. Like with `make`, you can add the flag `-j4` to use 4 (or whatever number of) compilers simultaneously. Also, the command `scons -h` will print some help information, and `scons --help` will print information about the options specific to TMV.

There are a number of command-line options that you might need (but try it with no flags first – it can often find everything automatically). The options are listed with their default value. You change them simply by specifying a different value on the command line. For example:

```
scons CXX=icpc INST_LONGDOUBLE=true
```

If you need to run SCons multiple times (for example to compile the test suite or install the libraries as described below), you only need to specify the new parameter values the first time you run the SCons. The program automatically saves your options and continues to use them until you change a value again.

- `CXX=g++` specifies which C++ compiler to use.

¹⁶ If this is unavailable for some reason, I also post the source code to <https://sourceforge.net/projects/tmv-cpp/>.

- `FLAGS=' '` specifies the basic flags to pass to the compiler. The default behavior is to automatically choose good flags to use according to which kind of compiler you are using. It has defaults for `g++`, `icpc` and `pgCC`. If you are using a different compiler or don't like the default, then you can specify this by hand. Remember to put the flags in quotes, so the whitespace doesn't confuse the parser. e.g. `scons FLAGS='-O3 -g'`
- `DEBUG=false` specifies whether to keep the debugging assert statements in the compiled library code.
- `PREFIX=/usr/local` specifies where to install the library when running `scons install` (see below).
- `INST_FLOAT=true` specifies whether to instantiate the `<float>` templates.
- `INST_DOUBLE=true` specifies whether to instantiate the `<double>` templates.
- `INST_LONGDOUBLE=false` specifies whether to instantiate the `<long double>` templates.
- `INST_INT=false` specifies whether to instantiate the `<int>` templates.
- `WITH_OPENMP=true` specifies whether to use OpenMP to parallelize some parts of the code.

The next flags set up the paths that SCons will use to try to find your BLAS and LAPACK libraries.

- `IMPORT_ENV=false` specifies whether to import the entire environment from the calling shell. The default is to start with a clean environment to be less susceptible to a particular user having an unusual set up. But sometimes sysadmins set things up in non-standard ways and use the environment variables to make everything work. If this is the case, then `IMPORT_ENV` should do the trick. It imports the environment, but does not add any `-I` or `-L` flags when compiling.
- `EXTRA_PATH=' '` specifies directories in which to search for executables (notably the compiler, although you can also just give the full path in the `CXX` parameter) in addition to the standard locations such as `/usr/bin` and `/usr/local/bin`. If you are giving multiple directories, they should be separated by colons.
- `EXTRA_INCLUDE_PATH=' '` specifies directories in which to search for header files (such as the BLAS or LAPACK header files) in addition to the standard locations such as `/usr/include` and `/usr/local/include`. These directories are specified as `-I` flags to the compiler. If you are giving multiple directories, they should be separated by colons.
- `EXTRA_LIB_PATH=' '` specifies directories in which to search for libraries (such as the BLAS or LAPACK libraries) in addition to the standard locations such as `/usr/lib` and `/usr/local/lib`. These directories are specified as `-L` flags to the linker. If you are giving multiple directories, they should be separated by colons.
- `IMPORT_PATHS=false` specifies whether to import extra path directories from the environment variables: `PATH`, `C_INCLUDE_PATH`, `LD_LIBRARY_PATH` and `LIBRARY_PATH`.

The next options can be used to specify what BLAS and/or LAPACK libraries to use (if any), overriding the default of using whatever libraries SCons chooses from searching through your path and trying to link the libraries that it finds. The `FORCE` options can be useful if SCons finds a library before trying the one that you want, or if SCons fails in the linking step even though the library should link successfully (I'm still not sure why this happens sometimes), or if you want to compile for a library that requires different linking instructions than the ones that SCons tries¹⁷. The `FORCE` options will try to test linking with the library requested, but if it fails, then it will just give a warning message.

¹⁷ If you have a case of needing different linking instructions, and your BLAS or LAPACK is a standard installation on your machine (not some goofy personal installation that no one else will duplicate), then let me know and I'll add it to the SConstruct file for the next release.

- `WITH_BLAS=true` specifies whether to look for and try to use a BLAS library.
- `WITH_LAPACK=false` specifies whether to look for and try to use a LAPACK library. Note: the default here used to be `true`. But my experience has been that the TMV code is more stable than typical LAPACK distributions, especially with regard to overflow and underflow. And it is generally equally fast, and sometimes faster. The only exception is finding Eigenvectors for extremely large hermitian matrices¹⁸. And even then, TMV may still be faster if you use a machine with multiple cores, since TMV seems to have better parallelization of this algorithm (if `WITH_OPENMP=true`) than many LAPACK libraries.
- `FORCE_MKL=false` forces the use of the Intel Math Kernel library. It requires the header file `"mkl.h"` to be found in your path.
- `FORCE_ACML=false` forces the use of the AMD Core Math library. It requires the header file `"acml.h"` to be found in your path.
- `FORCE_GOTO=false` forces the use of the GotoBlas library.
- `FORCE_ATLAS=false` forces the use of the ATLAS library (for BLAS). It requires the header file `"cblas.h"` to be found in your path.
- `FORCE_CBLAS=false` forces the use of a CBLAS library.
- `FORCE_FBLAS=false` forces the use of a Fortran BLAS library.
- `FORCE_CLAPACK=false` forces the use of the CLAPACK library. It requires the CLAPACK version of the header file `"clapack.h"` to be found in your path.
- `FORCE_ATLAS_LAPACK=false` forces the use of the LAPACK portion of the ATLAS Library. It requires the ATLAS version of the header file `"clapack.h"` to be found in your path.
- `FORCE_FLAPACK=false` forces the use of a Fortran LAPACK library.
- `LIBS=' '` directly specifies the library flags to use for linking if the automatic methods aren't working for you. Because of the way SCons works, these should omit the `-l` part of the flag, since SCons will add this to what is provided. For example, to specify an alternate name for the CLAPACK library, use `scons LIBS=lapack_LINUX`. Multiple libraries here should be separated by whitespace and enclosed in quotes.

Finally, some miscellaneous options that you are less likely to need:

- `STATIC=false` specifies whether to use static linkage. Some systems have trouble with dynamic linkage of libraries. This usually indicates that something is installed incorrectly, but it can be easier to just use static linkage when you compile as a workaround. This flag does this for the test suite executables.
- `WITH_SSE=true` specifies whether to use the `-msse2` flag with `icpc` compilations. It seems like most machines that have `icpc` are able to use SSE commands. However, I couldn't figure out a compiler flag that would turn on SSE *if and only if* the machine supports it. So the default is to use the flag `-msse2`, but you can disable it by setting `WITH_SSE=false` if your machine doesn't have SSE support.
- `XTEST=0` specifies whether to include extra tests in the test suite. `XTEST` is treated as a bit set, with each non-zero bit turning on particular tests. Type `"scons -h"` for more information.
- `MEM_TEST=false` specifies whether to include extra memory tests in the library and test suite.

¹⁸ LAPACK implements an algorithm called Relatively Robust Representation (RRR) to find the eigenvectors, rather than the divide and conquer algorithm used by TMV. RRR is faster, but only when the matrix size is sufficiently large. Including this algorithm in TMV is #6 on my To-do list (§14).

- `USE_STEGR=true` specifies whether to use the LAPACK algorithm called `dsteqr` (or `ssteqr` for `<float>`) for symmetric eigenvector calculation. If it is false, the divide-and-conquer algorithm, named `dstedc` (or `sstedc`) will be used instead.
- `USE_GEQP3=true` specifies whether to use the LAPACK algorithm called `dgeqp3` (or its variants) for the strict QRP decomposition. If it is false, the native TMV code will be used instead.
- `SMALL_TESTS=false` specifies whether to make the smaller test suite programs: `tmvtest1a`, `tmvtest1b`, etc.
- `WARN=false` specifies whether to have the compiler report warning (with `-Wall` or something similar).
- `NOMIX_SMALL=false` specifies whether to avoid mixing Small and regular arithmetic. (This is automatically selected if `GotoBLAS` is being used.)
- `CACHE_LIB=true` specifies whether to cache the results of the library checks. Scons usually detects correctly when something salient has changed to require a recompilation. However, sometimes it doesn't, in which case setting this to false will force recompilation of everything each time you type `scons`.
- `WITH_UPS=false` specifies whether to install TMV information for ups into the `PREFIX/ups` directory.
- `TEST_DEPRECATED=false` specifies whether to use the old deprecated method names in the test suite.

When SCons starts up, it will look through the standard paths, along with any extra paths you have specified with the above options, to find BLAS and LAPACK libraries. This can sometimes require a few iterations to get working correctly. You should look at the initial output from SCons to make sure it finds the correct BLAS and LAPACK libraries that you think it should find. Here is a sample output:¹⁹

```
$ scons
scons: Reading SConscript files ...

Using compiler: g++-4
compiler version: 4.3.2
Debugging turned off
Checking for MKL... no
Checking for ACML... no
Checking for GotoBLAS... no
Checking for CBLAS... yes
Using CBLAS
Checking for CLAPACK... no
Checking for Fortran LAPACK... yes
Using Fortran LAPACK
scons: done reading SConscript files.
scons: Building targets ...
```

(Starts the actual compiling)

If a “Checking for...” line ends with `no`, even though you think that library is installed on your computer, then it probably means that you need to tell SCons which directories to search, in addition to the standard locations. The most straightforward way to do this is with the parameters `EXTRA_INCLUDE_PATH` and `EXTRA_LIB_PATH`. These are described in detail above. See also `IMPORT_ENV` and `IMPORT_PATHS`.

¹⁹ This is the exact output that I get with the default options on my MacBook.

4. (Optional) Type

```
scons test
```

This will make three executables called `tmvtest1`, `tmvtest2` and `tmvtest3` in the `test` directory.

Then you should run the three test suites. They should output a bunch of lines reading *[Something]* passed all tests. If one of them ends in a line that starts with `Error`, then please post a bug report at <http://code.google.com/p/tmv-cpp/issues> about the problem including what compiler you are using, some details about your system, and what (if any) BLAS and LAPACK libraries you are linking to.

If you specify `SMALL_TESTS=true`, then the smaller test executables `tmvtest1a-d`, `tmvtest2a-c`, and `tmvtest3a-e` (where `a-d` means four files with each of `a`, `b`, `c` and `d`) will be made instead. These perform the same tests as the larger test executables, but can be easier for some linkers.

5. Type

```
scons install
```

(or possibly `sudo scons install` if you are installing into `/usr/local` or somewhere similar).

This will install the necessary header files into the directory `/usr/local/include` and the libraries into `/usr/local/lib`. As mentioned above, you can also specify a different prefix with the command line option `PREFIX=install-dir`. A common choice for users without `sudo` privileges is `PREFIX=~` which will install the library in `~/include` and `~/lib`.

At the end of the installation process, you should see a message similar to:

```
The TMV library was successfully installed.
To link your code against the TMV library, you should use the
link flags:
```

```
-ltmv -llapack -lblas -lpthread -fopenmp
```

Or if you are using `Band`, `Sym` or `SymBand` matrices, use:

```
-ltmv_symband -ltmv -llapack -lblas -lpthread -fopenmp
```

These flags (except for the optional `-ltmv_symband`) have been saved in the file:

```
/usr/local/share/tmv-link
```

so you can automatically use the correct flags in a makefile (for example) by using lines such as:

```
TMVLINK := $(shell cat /usr/local/share/tmv-link)
LIBS = $(TMVLINK) [... other libs ...]
```

```
scons: done building targets.
```

(Again, this is the actual output on my laptop when I type `sudo scons install`.) These instructions tell you what BLAS and LAPACK libraries were found on your system and are needed for proper linkage of your code. The linkage flags are stored in a file on your computer so that you can automatically get the

linkage correct according to what options you decide to use when installing TMV. This is especially useful on systems where a system administrator installs the library, which is then used by many users.

You can have your makefiles read this file as described above. Or if you are using SCons to build your software, you can similarly add the contents of the file into `env[LIBS]`. The `examples` directory has a makefile that uses the above lines (using the local `share` directory rather than the installed location).

12.3 Installing using make

This is a more standard installation method if you don't want to bother with SCons, or if you are just more familiar with the make program and don't want to deal with SCons. It requires a bit more user input in terms of directly specifying the BLAS and/or LAPACK libraries, but because of that it is also a bit easier to exactly customize the installation if you want to do something non-standard.

1. Download and unpack the tarball as described above (§12.1).
2. Edit the makefile.

The start of the makefile lists 5 things to specify: the compiler, the include directories, the other flags to send to the compiler, any necessary BLAS/LAPACK linkage flags, and the installation directory. The default setup is:

```
CC= g++
INCLUDE= -Iinclude
CFLAGS= $(INCLUDE) -O2 -DNOBLAS -DNDEBUG
BLASLIBS=
PREFIX=/usr/local
```

but you will probably want to change this²⁰.

This default setup will compile using g++ without any BLAS or LAPACK library and with debugging turned off. This setup should work on any system with gcc, although it almost certainly won't be as fast as using an optimized BLAS library and/or a LAPACK library.

You should edit these so that:

- `CC` specifies the C++ compiler to use.
- `INCLUDE` specifies the directories for any BLAS and LAPACK header files you want to include. (You should leave `-Iinclude` there as well.)
- `CFLAGS` contains any compiler flags you want. (See below for TMV-specific flags to consider.)
- `BLASLIBS` specifies the libraries required for linking your BLAS and LAPACK libraries.
- `PREFIX` specifies the directory for `make install` to copy the files to.

After these lines, there are several commented-out examples for different systems using various BLAS and LAPACK versions, showcasing some of the compiler options described below, and giving examples of what you need for several common (or at least representative) systems. If you have a system similar to one of these, then it should be a good starting point for you to figure out what you want to use.

See the next section below for a complete list of compiler flags that control how the TMV library is built.

²⁰ The original version of the `Makefile` is copied in the file `Makefile.original` in case your copy gets messed up, and you want to go back to the original.

3. (Advanced usage) Edit the Inst and Blas files.

By default, the library will include instantiations of all classes and functions that use either `double` or `float` (including complex versions of each). There are a few flags, such as `-DNO_INST_FLOAT` and `-DINST_LONGDOUBLE`, that change this as described below. But if you want to compile routines for some other class, such as a user-defined `MyQuadPrecisionType` class, then you will need to modify the file `src/TMV_Inst.h`. You simply need to add the lines:

```
#define T MyQuadPrecisionType
#include InstFile
#undef T
```

to the end of the file before compiling. (Obviously, you should replace `MyQuadPrecisionType` with whatever type you want to be instantiated in the library.)

Also, the file `src/TMV_Blas.h` sets up all the BLAS and LAPACK calling structures, as well as the necessary `#include` statements. So if the BLAS or LAPACK options aren't working for your system, you may need to edit these files as well. This is especially true if your BLAS or LAPACK versions are not one of ATLAS, CLAPACK, FLAPACK, MKL, or ACML. The comments at the beginning of `TMV_Blas.h` gives instructions on how to set up the file for other installations.

4. Type:

```
make libs
```

This will make the TMV libraries, `libtmv.a` and `libtmv.symband.a`, which will be located in the directory `lib`.

5. (Optional) Next type:

```
make test
```

This will make three executables called `tmvtest1`, `tmvtest2` and `tmvtest3` in the `bin` directory.

Then you should run the three test suites. They should output a bunch of lines reading *[Something]* passed all tests. If one of them ends in a line that starts with `Error`, then please post a bug report at <http://code.google.com/p/tmv-cpp/issues> about the problem including what compiler you are using, some details about your system, and what (if any) BLAS and LAPACK libraries you are linking to.

6. Next type:

```
make install
```

(or possibly `sudo make install` if you are installing into `/usr/local` or somewhere similar).

This will copy the header files and the libraries to `PREFIX/include` and `PREFIX/lib`. The default `PREFIX` is `/usr/local`, but you can change this in the `Makefile` to some other installation directory if you want.

7. (Optional) You may also want to make the example programs by typing:

```
make examples
```

This will make five executables called `vector`, `matrix`, `division`, `bandmatrix`, and `symmatrix`. These programs, along with their corresponding source code in the `examples` directory, give concrete examples of some of the common things you might want to do with the TMV library. They don't really try to be comprehensive, but they do give a pretty good overview of the main features, so looking at them may be a useful way to get started programming with TMV.

Also, they are a lot quicker to compile than the test suite, so they can provide a quick way to make sure that you have specified the correct linking flags for any BLAS or LAPACK libraries that you might need (according to the `-D` flags you specified for building the TMV library).

8. Compile your program

Each .cpp file that uses TMV will need to have

```
#include "TMV.h"
```

at the top. If you are using any of the special matrix types (other than diagonal and triangular), then you also need to include their particular header files as well. (e.g. `TMV_Band.h`, `TMV_Small.h`, etc.)

If you did not install the header files (with `make install` above) to a directory that is in the standard path for header files, then you will need to use the compile flag `-I[tmvdir]/include` when making the object file to tell the compiler where the TMV header files are.

For the linking step, you need to compile with the flag `-ltmv`. If you are using band, symmetric or symmetric band matrices in your code, then you will need to link with the flags `-ltmv_symband` `-ltmv`. And if you did not install the libraries to a directory in your linking path, then you need to include the flag `-L[tmvdir]/lib` to tell the linker where the libraries are.

If you are using BLAS and/or LAPACK calls from the TMV code, then you will also need to link with their libraries. For example, for my version of Intel's Math Kernel Library, I use `-lmkl_lapack` `-lmkl_ia32` `-lguide` `-lpthread`. For ATLAS, I use `-llapack` `-lcblas` `-latlas`. For your specific installation, you may need the same thing, or something slightly different, including possibly `-L` flags to indicate where the BLAS or LAPACK libraries are located.

12.4 Compiler flags

Here are some compiler define flags to consider using:

- `-DNDEBUG` will turn off debugging. My recommendation is to leave debugging on in your own code that is calling TMV routines, since it doesn't slow things down too much. Then when you decide to export a version of the code that needs to be as fast as possible, recompile with this flag. On the other hand, the internal TMV code should be pretty well debugged when you get it, so for compiling the library, you should go ahead and use this flag.

The rest of the flags are only relevant when compiling the TMV library, so you won't need any of them when compiling your own code that uses TMV routines. And if you install with SCons, it will determine which flags to set according to the somewhat simpler list of SCons options.

- `-DNO_INST_FLOAT` will not instantiate any `<float>` classes or routines.
- `-DNO_INST_DOUBLE` will not instantiate any `<double>` classes or routines.
- `-DNO_INST_COMPLEX` will not instantiate any complex classes or routines.
- `-DINST_LONGDOUBLE` will instantiate `<long double>` classes and routines.
- `-DINST_INT` will instantiate `<int>` classes and routines.

The next flags specify what BLAS and/or LAPACK libraries to use (if any):

- `-DNOBLAS` will not call any external BLAS or LAPACK routines
- `-DNOLAP` will not call any external LAPACK routines
- `-DATLAS` will set up the BLAS calls as defined by ATLAS. And (if `-DNOLAP` is not specified), it will also call the several LAPACK routines provided by ATLAS.
- `-DCBLAS` will set up the BLAS calls using the notation provided by the CBLAS interface to BLAS. Many systems have a library like this installed as `/usr/local/lib/libcblas.a`, or some similar location, so it is worth looking to see if you have this before trying to install something new.

- `-DFBLAS` will set up the BLAS calls using a library that was designed to be called from Fortran programs rather than C. Most BLAS installations also install the C interface, but if your BLAS does not have this (e.g. GotoBLAS), then this flag should work. Since there is often no header file for a Fortran BLAS library, we include the file `fblas.h` in the `TMV include` directory that should work.
- `-DCLAPACK` will set up the LAPACK calls for the CLAPACK distribution. I find this version easier to get installed than the Fortran LAPACK distribution, so I would recommend using this if you don't already have a version of LAPACK installed somewhere on your system. Defining both `ATLAS` and `CLAPACK` will use the CLAPACK version for all LAPACK routines, including the ones also provided by ATLAS. That is, ATLAS will only be used for its BLAS routines. If you want the ATLAS versions of its few LAPACK routines instead, the ATLAS installation instructions describe a way to get them into the CLAPACK library. Also, you should make sure the `-I` directories lists the directory with the CLAPACK version of `clapack.h` before the ATLAS version.
- `-DFLAPACK` will set up the LAPACK calls for a distribution that was designed to be called from Fortran programs rather than C. Since such distributions do not provide a header file, we provide `flapack.h` in the `TMV include` directory.
- `-DMKL` will call all the external BLAS and LAPACK routines as defined by the Intel Math Kernel Library. You should specify the directory with `mk1.h` with a `-I` flag if it is not installed in your include path.
- `-DACML` will call all the external BLAS and LAPACK routines as defined by the AMD Core Math Library. You should specify the directory with `acml.h` with a `-I` flag if it is not installed in your include path.

The next set of compiler defines are not usually necessary. But if you have problems, these might be useful:

- `-DXTEST=[num]` will do extra testing in the test suite, as well as add a few $O(N)$ time assert statements. (Most of the assert statements that are normally run only take $O(1)$ time.) I always do tests with `XTEST=127` to turn on all of the extra tests before releasing a version, but the executable gets quite large, as do many of the `TMV_Test*.o` files. Plus it takes a lot longer to compile. So the default is to not use it.
- `-DMEMTEST` will do extra testing of the memory usage to look for memory leaks and other memory errors. Again, there is probably no need for you to use this flag. But if you think there is a memory leak in TMV, this could help you find it, in which case please let me know.
- `-DNOSTEGR` specifies that the LAPACK algorithm called `dstegr` (or `sstegr` for `<float>`) should not be used for symmetric eigenvector calculation. In this case, the divide-and-conquer algorithm, named `dstedc` (or `sstedc`) will be used instead.

I used to have problems with the `stegr` algorithm producing incorrent results, but I figured out how to detect the problem and now the TMV code starts by calling `stegr`, and then if it fails it calls the slower, but more robust, `stedc` routine. If you want to avoid this behavior and always use the `stedc` algorithm, you can compile with this flag.

- `-DNOGEQP3` specifies that the LAPACK algorithm called `dgeqp3` (or its variants) should not be used for the strict QRP decomposition. In this case, the native TMV code will be used instead.

I have found that the LAPACK code for `?geqp3` does not always produce an R matrix with strictly decreasing elements along the diagonal. So if this is important for you, then you should use this flag.

- `-DWORKQUERY` specifies that your LAPACK library does not support work size queries. There are a few LAPACK routines that require workspace whose optimal size depends on details of the L1 cache size of your machine. Normally, you can pass the routine a work size of `-1` to indicate that you are doing a workspace query, in which case the routine returns the optimal value to use. This is the normal mode of operation for TMV.

However, some older LAPACK distributions do not support this process. The example I found was the MKL with icc version 9.0. So if you get errors similar to

TMV Error: info < 0 returned by LAPACK function dormqr

then this flag will compile without the workspace queries, instead just using a good guess for the optimal size.

- `-DNOMIX_SMALL` will avoid all arithmetic that mixes `SmallMatrix` and `SmallVector` with regular `Matrix` or `Vector` in the test program `tmvtest3`. Some BLAS libraries – specifically `GotoBLAS`, possibly others – don’t work with this kind of mixing, so this flag allows you to test the rest of the `Small` arithmetic if the mixing is causing problems for your system.
- `-DNOSTL` uses some workarounds for segments of code that use the STL library, but which didn’t work for one of my compilers. I’m pretty sure it is because the compiler wasn’t installed correctly, so I don’t think you should really ever need to use this flag. But in any case, it will use a median-of-three quicksort algorithm for sorting rather than the standard library’s sort. And it manually reads strings using `character reads`, rather than using the `>>` operator.
- `-DXDEBUG` will do different extra (much slower) debugging. This one checks for incorrect results from the various algorithms by doing things the simple slow way and comparing the results to the fast blocked or recursive or in-place version to make sure the answer isn’t (significantly) different. I use this one a lot when debugging new algorithms, usually on a file-by-file basis. Again, you shouldn’t need this for an official release version. But if you do get wrong answers for something, you could use this to try to find the problem.
- `-DTMV_BLOCKSIZE=NN` will change the block size used by some routines. The current value is 64, which is good for many computers. The optimal value will depend on the size of your CPU’s L1 cache. So if you want to try to tune the algorithms, you can modify this value to something more appropriate for your computer. However, you are probably better off just using an optimized BLAS and LAPACK library and let TMV call that instead.

12.5 Installing using CMake

CMake is another automated build tool. I am not very well versed in it, but a user of the TMV library, Andy Molloy, sent me the `CMakeLists.txt` files that he used to compile TMV, so I am distributing them for anyone else who wants to use them. It is an easy way to build a basic TMV library, but I am not sure how hard it is to get it working with BLAS or LAPACK libraries²¹. The basic usage is as follows:

1. Download and unpack the tarball as described above (§12.1).
2. Download CMake from <http://www.cmake.org/> if you do not already have it.
3. In the TMV directory, type

```
cmake -i .
or
ccmake .
```

for Unix, or run `CMakeSetup.exe` in Windows.
Answer the configuration questions and/or set the options listed.
4. Type:

```
make tmv
make tmv_syband
```

to make the libraries.

²¹ I was about to start learning more about it when a colleague, Erin Sheldon, introduced me to SCons, which I like a lot, so I haven’t been motivated to go back to CMake and really flesh it out like I did with the SCons installation script.

5. (Optional) Type:

```
make tmvtest1
make tmvtest2
make tmvtest3
to make the test suites.
```

I haven't incorporated most of the options that I added to the SCons installation above. But I think all of that is possible with CMake, so if someone who is an expert in CMake wants to do this and send me the CMake files for it, I'll include them in future releases.

12.6 Installation using Microsoft Visual C++

Using TMV with Microsoft Visual C++ is a bit different from the other compilers, since it has a Windows framework for building programs, rather than a command line²².

1. Download and unpack the tarball as described above (§12.1).

There are many Windows utilities that can unpack the tarball. With IZArc (<http://www.izarc.org/>), for example, you right click on the file `tmv0.64.tar.gz` in Windows Explorer, select IZArc, then select Extract Here. This should make a directory called `tmv0.64` which has all of the source code and other files for the TMV library.

2. Start Microsoft Visual C++.

I have Visual C++ 2008 Express Edition, so all instructions below about menus and such refer that that edition. I would assume that other editions have their menus arranged similarly, but there may be some differences.

3. Open the "Solution" file `tmvtest1.sln`.

Each test suite needs to be made individually, so I'll give detailed instructions for the first test program. The other two are made the same way.

Go to the File menu. Select Open. Then select Project/Solution... In the `tmv0.64` directory, look for the file `tmvtest1.sln` and select it.

This includes a project for the first test program, `tmvtest1.vcproj`, and a project for the main TMV library, `tmv.vcproj`.

4. Select Debug or Release mode.

There are two modes for building programs in Visual C++: Debug and Release. You can choose either one. Debug mode will compile faster and Release mode will execute faster.

When writing your own code, you will probably want to start with Debug mode and then switch to Release when everything is working, so you will end up compiling the TMV library in both modes anyway.

You can select the mode in a pull down menu in the top button bar. With my setup (which I think is the default), it is directly to the right of a green arrow.

5. Build `tmvtest1`.

Go to the Build menu. Select Build Solution.

²² There is a way to do compile on the command line, but I suspect that will not be the usual way that people will want to use the TMV library. If you are interested in compiling this on the command line, I did have success installing SCons and then, within the Visual Studio 2008 Command Prompt, using the command: `C:\Python26\Scripts\scons CXX=cl IMPORT_ENV=true`.

6. Run `tmvtest1`.

Go to the Debug menu. Select Start Without Debugging. (Of course, you can instead choose Start Debugging if you'd rather, but that will make it run slower.)

A console window should open up to show the output of the program. When it is done, you should see the message, "Press any key to continue . . ."

The previous output lines should all read *[Something]* passed all tests. If the last line starts with `Error`, then please post an issue at <http://code.google.com/p/tmv-cpp/issues> about the problem.

7. Repeat for `tmvtest2` and `tmvtest3`.

Note: the solution file for `tmvtest2` includes the project for the TMV library with symmetric and banded matrices, `tmv_symband.vcproj`.

8. Include `tmv.vcproj` (and `tmv_symband.vcproj` if necessary) in the solution for your own project.

Any program that uses the TMV library needs to include the project `tmv.vcproj`. To do this, Go to the File, select Add, Existing Project... In the `tmv0.64` directory, look for the file `tmv.vcproj` and select it. If you are going to be using symmetric and/or banded matrices, then you should also select `tmv_symaband.vcproj`.

Then, in the Solution Explorer window, select your project. Then go to the Project menu and select Properties. Select Common Properties, then Framework and References. Click the Add New Reference... button. Select the TMV library (or libraries), and press OK.

Next you need to tell your program where to look for the TMV header files. So on the Property Page, select All Configurations at the top. Then go to Configuration Properties, C/C++, General. At the entry for Additional Include Directories, click the [...] button at the right. Click on the folder icon to add a new line, and then click the [...] button. Browse to the `tmv0.64` directory, and select the include directory. Press OK three times.

Now you should be all set to use the TMV library in your program.

12.7 Known compiler issues

I have tested the code using the following compilers:

GNU's g++ – versions 3.4.6, 4.1.2, 4.3.3, 4.4.2, and 4.4.4

Apple's g++ – version 4.0.1, 4.2.1

Intel's icpc – versions 9.1, 10.1, 11.1

Portland's pgCC – version 6.1

Microsoft's cl – Visual C++ 2008 Express Edition

It should work with any ansi-compliant compiler, but no guarantees if you use one other than these²³. So if you do try to compile on a different compiler, I would appreciate it if you could let me know whether you were successful. Please email the TMV discussion group at <mailto:tmv-discuss@googlegroups.com> with your experience (good or bad) using compilers other than these.

There are a few issues that I have discovered when compiling with various versions of compilers, and I have usually come up with a work-around for the problems. So if you have a problem, check this list to see if a solution is given for you.

²³ It does seem to be the case that every time I try the code on a new compiler, there is some issue that needs to be addressed. Either because the compiler fails to support some aspect of the C++ standard, or they enforce an aspect that I have failed to strictly conform to.

- **g++ version 4.4 on Snow Leopard:** g++ 4.4 (at least 4.4.2 through 4.4.4) has a pretty egregious bug in its exception handling that can cause the code to abort rather than allow a thrown value to be caught. This is a reported bug (42159), and according to the bug report it seems to only occur in conjunction with MacOS 10.6 (Snow Leopard), however it is certainly possible that it may show up in other systems too.

However, I have discovered a workaround that seems to fix the problem. Link with `-lpthread` even if you are not using OpenMP or have any other reason to use that library. Something about linking with that library fixes the exception handling problems. In a SCons installation, this is done automatically for g++ versions 4.4 and 4.5. If you are installing with make, then you'll have to add this yourself. Likewise, you'll need to link your own code with `-lpthread` as well.

- **g++ -O2, versions 4.1 and 4.2:** It seems that there is some problem with the -O2 optimization of g++ versions 4.1.2 and 4.2.2 when used with TMV debugging turned on. Everything compiles fine when I use `g++ -O or g++ -O2 -DNDEBUG`. But when I compile with `g++ -O2 (or -O3)` without `-DNDEBUG`, then the test suite fails, getting weird results for some arithmetic operations that look like uninitialized memory was used.

I distilled the code down to a small code snippet that still failed and sent it to Gnu as a bug report. They confirmed the bug and suggested using the flag `-fno-strict-aliasing`, which did fix the problems.

Another option, which might be a good idea anyway is to just use `-O` when you want a version that includes the TMV assert statements, and make sure to use `-DNDEBUG` when you want a more optimized version.

- **Apple g++:** Older versions of Apple's version of g++ that they shipped with the Tiger OS did not work for compilation of the TMV library. It was called version 4.0, but I do not remember the build number. They seem to have fixed the problem with later XCode update, but if you have an older Mac and want to compile TMV on it and the native g++ is giving you trouble, you should either upgrade to a newer Xcode distribution or download the real GNU gcc instead; I recommend using Fink (<http://fink.sourceforge.net/>).
- **pgCC:** I only have access to pgCC version 6.1, so these notes only refer to that version.

- Apparently pgCC does not by default support exceptions when compiled with openmp turned on. So if you want to use the parallel versions of the algorithms, you need to compile with the flags `-mp --exceptions`. This is a documented feature, but it's not very obvious, so I figured I'd point it out here.
- There was a bug in pgCC version 6.1 that was apparently fixed in version 7.0 where `long double` variables were not correctly written with `std::ostream`. The values were written as either 0 or -. So I have written a workaround in the code for pgCC versions before version 7.0²⁴, where the `long double` values are copied to `double` before writing. This works, but only for values that are so convertible. If you have values that are outside the range representable by a `double`, then you may experience overflow or underflow on output.

- **Borland's C++ Builder:** I tried to compile the library with Borland's C++ Builder for Microsoft Windows Version 10.0.2288.42451 Update 2, but it failed at fairly foundational aspects of the code, so I do not think it is possible to get the code to work. However, if somebody wants to try to get the code running with this compiler or some other Borland product, I welcome the effort and would love to hear about a successful compilation (at <mailto:tmv-discuss@googlegroups.com>).
- **Sun CC:** I also tried to compile the library with Sun's CC compiler, version 5.3, and it didn't understand some of the template syntax used in TMV. However, this version is pretty old (2001), so it is likely that newer versions are more compliant with the C++ standard (they are now up to version 5.9). Unfortunately, I do not

²⁴ Thanks to Dan Bonachea for making available his file `portable_platform.h`, which makes it particularly easy to test for particular compiler versions.

have access to a newer version of CC yet, so I have not been able try it. As usual, if you have managed to compile TMV with any version of Sun's CC, I would love to hear about it (at <mailto:tmv-discuss@googlegroups.com>tmv-discuss@googlegroups.com).

- **Memory requirements:** The library is pretty big, so it can take quite a lot of memory to compile. For most compilers, it seems that a minimum of around 512K is required. For compiling the test suite with the `-DXTEST` flag, more than 2GB of memory is recommended.
- **Linker choking:** Some linkers (e.g. the one on my old Mac G5) have trouble with the size of some of the test suite's executables, especially when compiled with `-DXTEST`. If you encounter this problem, you can instead compile the smaller test suites.

The tests in `tmvtest1` are split into `tmvtest1a`, `tmvtest1b` and `tmvtest1c`. Likewise `tmvtest2` has a, b and c versions, and `tmvtest3` has a, b, c and d versions. These are compiled by typing `make test1a`, `make test1b`, etc. So if you want to run the tests on a machine that can't link the full programs, these smaller versions can help.

With the SCons installation, you can type `scons test SMALL_TESTS=true`. Or you can make them one at a time by typing `scons test1a`, `scons test1b`, etc.

You might also try testing only one type at a time: First compile with `-DNO_INST_FLOAT`, and then with `-DNO_INST_DOUBLE` (or with SCons, use `INST_FLOAT=false` and then `INST_FLOAT=true` `INST_DOUBLE=false`). This cuts the size of the executables in half, which can also help if the above trick doesn't work. (I had to do this on one of my test systems for `test2c`.)

- **BLAS/LAPACK problems:** There are a number of possible errors that are related to particular BLAS or LAPACK distributions, or combinations thereof:
 - **Strict QRP decomposition fails:** Some versions of the LAPACK function `dgeqp3` do not produce the correct result for R . The diagonal elements of R are supposed to be monotonically decreasing along the diagonal, but sometimes this is not strictly true. This probably varies among implementations, so your version might always succeed.
But if this feature is important to you, then you can compile with the flag `-DNOGEQP3`, which will use the native TMV code for strict QRP decomposition for this rather than the LAPACK function (which is called `?geqp3`).
 - **Info < 0 errors:** If you get an error that looks something like:
TMV Error: info < 0 returned by LAPACK function `dormqr`
then this probably means that your LAPACK distribution does not support workspace size queries. The solution is to use the flag `-DNO_WORKQUERY`.
 - **Unable to link LAPACK on Mac:** I get a linking error on my Mac when combining the XCode LAPACK library (either `libclapack.dylib` or `liblapack.dylib`) with a non-XCode BLAS library. Specifically, it can't find the `?getri` functions. Basically, the reason is that the XCode LAPACK libraries are designed to be used with one of the BLAS libraries, `libcblas.dylib` or `libblas.dylib`, in the same directory, and that BLAS library has the `getri` functions, even though they are properly LAPACK functions. Anyway, it's basically a bug in the Apple XCode distribution of these files, and the result is that if you use a different BLAS library, the linking fails.
The solution is either to use one of the XCode BLAS library or to install your own CLAPACK library. If you do the latter, you will probably want to rename (or delete) these Mac library files, since they are in the `/usr/lib` directory, and `-L` flags usually can't take precedence over `/usr/lib`.
 - **Errors in Matrix<float> calculations using BLAS on Mac** I have found that some XCode BLAS libraries seem to have errors in the calculations for large <float> matrices. I think it is an error in the `sgemm` function. Since very many of the other algorithms use this function, the error propagates to

lots of other routines as well. The errors seem to be only for large matrices with at least one dimension $N > 100$ or so. (I'm not sure of the exact crossover point from working to non-working code.)

I haven't surveyed which versions of XCode have this bug, but it seems to be in version 3.2.2 at least. I reported the bug to Apple (ID 8069848) fairly soon before this release, so they have understandably not fixed it yet. The best thing to do to see if your XCode installation has this problem is to install the test suite and make sure that it runs without errors. If it makes it through without errors, then you don't have to worry about it.

If you do get an error in the test suite, and you want to use `<float>` matrices, then the only workaround I can suggest is to either compile without using BLAS or download and make a different BLAS library. For example ATLAS or GOTO BLAS are free to download and are both pretty fast.

- **Segmentation faults with `SmallMatrix`:** My version of the GotoBLAS library seems to have trouble with arithmetic statements that mix `SmallMatrix` or `SmallVector` with `Matrix` or `Vector`. I think the problem stems from the fact that the `Small` objects allocate their memory on the stack, while the regular objects allocate their memory on the heap. One of the ATLAS installations I tested also had similar trouble.

If your BLAS distribution has problems running `tmvtest3` (which is the one that runs the tests of `SmallMatrix` and `SmallVector`), then you should try compiling with the flag `-DNOMIX_SMALL` to tell the code to skip the tests that mix `Small` and regular objects. And of course, you will also want to remember to watch out for this in your own code if you use `SmallMatrix` and `SmallVector`. The SCons installer automatically uses this flag when using GotoBLAS.

If you want to test the mixed arithmetic specifically, you can use `scons testmix` to compile the test program `tmvtestmix`. I am planning on having `scons` compile and run this program before compiling `tmvtest3` to see whether it needs the flag, but I haven't figured out how to do that yet.

- **Overflow and underflow when using LAPACK:** The normal distribution of LAPACK algorithms are not as careful as the TMV code when it comes to guarding against overflow and underflow. As a result, you may find matrix decompositions resulting in values with `nan` or `inf` when using TMV with LAPACK support enabled. The solution is to compile TMV without the LAPACK library (e.g. by using `WITH_LAPACK=false` in the SCons installation method). This does not generally result in slower code, since the native TMV code is almost always as fast (usually using the same algorithm) as the LAPACK distribution, but has better guards against overflow and underflow.

- **Standard Template Library:** In a few places, the TMV code uses standard template library algorithms. Normally these work fine, but a couple of the compilers I tested the code on didn't seem to have the STL correctly installed.

On one computer, the linker complained that it couldn't find the code for `sort`, even though `sort` should be completely inlined, so the code should already have been in `TMV_Vector.o`, where I use it. Similarly, it had trouble linking the string read commands needed for reading in `SymBandMatrixes`. Again, this should be included in the object file where I use it, `TMV_SymBandMatrix.o`.

The problem is probably due to something not being installed correctly on that computer, or maybe I just did not include the correct linkages or something. But rather than trying to get the sysadmin for the computer to find and fix the problem, I just added an option to compile with a simple median-of-three quicksort algorithm, rather than the STL `sort` command, and to read the strings in character by character. You can use this option by compiling with the flag `-DNOSTL`.

- **Non-portable `IsNaN()`:** The LAPACK function `sstegr` sometimes produces `nan` values on output. Apparently there is a bug in at least some distributions of this function. Anyway, TMV checks for this and calls the alternative (but slower) function `sstedc` instead whenever a `nan` is found. The problem is that there is no C++ standard way to check for a `nan` value.

The usual way is to use a macro `isnan`, which is usually defined in the file `<math.h>`. However, this is technically a C99 extension, not standard C++. So if this macro is not defined, then TMV tries two other tests that usually detect `nan` correctly. But if this doesn't work correctly for you, then you may need to edit the file `src/TMV_IsNaN.cpp` to work with your system²⁵.

Alternatively, you can compile with `-DNOSTEGR` which will always use `sstedc` instead of `sstegr` at the expense of a bit of speed. Since this is the only place we need to use the `IsNaN` function, that should fix the problem.

²⁵ However, the provided code did work successfully on all the compilers I tested it on, so technically this is not a “known” compiler issue, just a potential issue.

13 Example Code

Here are five complete programs that use the TMV library. They are intended to showcase some of the features of the library, but they are certainly not comprehensive. Hopefully, they will be useful for someone who is just getting started with the TMV library as a kind of tutorial of the basics.

The code includes the output as comments that start with `//!`, so you can more easily see what is going on. And each file listed here is also included with the TMV distribution in the `examples` directory, so you can easily compile and run them yourself.

13.1 Vector

```
#include "TMV.h"
#include <iostream>

int main() try
{
    // Several ways to create and initialize vectors:

    // Create with uninitialized values
    tmv::Vector<double> v1(6);
    // In debug mode, all are initialized to 888 to make it easier to
    // notice when you fail to initialize correctly.
    std::cout<<"v1 = "<<v1<<std::endl;
    //! v1 = 6 ( 888 888 888 888 888 888 )

    // Initialize with STL-compliant iterator
    std::generate(v1.begin(), v1.end(), rand);
    v1 /= double(RAND_MAX);
    std::cout<<"v1 = "<<v1<<std::endl;
    //! v1 = 6 ( 0.840188 0.394383 0.783099 0.79844 0.911647 0.197551 )

    // Create with all 2's.
    tmv::Vector<double> v2(6, 2.);
    std::cout<<"v2 = "<<v2<<std::endl;
    //! v2 = 6 ( 2 2 2 2 2 2 )

    // Create with given elements from C-array
    double vv[6] = {1.1, 8., -15., 2.5, 6.3, -12.};
    tmv::Vector<double> v3(6, vv);
    std::cout<<"v3 = "<<v3<<std::endl;
    //! v3 = 6 ( 1.1 8 -15 2.5 6.3 -12 )

    // Initialize with direct access of each element
    tmv::Vector<double> v4(6);
    for(size_t i=0; i<v4.size(); i++) {
        v4(i) = 2*i+10.; // Could also use v4[i] instead of v4(i)
    }
    std::cout<<"v4 = "<<v4<<std::endl;
    //! v4 = 6 ( 10 12 14 16 18 20 )

    // Initialize with comma-delimited list
    v4 << 1.2, 9., 12, 2.5, -7.4, 14;
    std::cout<<"v4 = "<<v4<<std::endl;
    //! v4 = 6 ( 1.2 9 12 2.5 -7.4 14 )
    // If the list is the wrong size, a run time error occurs:
    try {
        v4 << 1.2, 9., 12, 2.5;
    } catch (tmv::Error& e) {
        std::cout<<"Caught e = "<<e;
        //! Caught e = TMV Read Error: Reading from list initialization.
        //! Reached end of list, but expecting 2 more elements.
    }
    try {
        v4 << 1.2, 9., 12, 2.5, -7.4, 14, 99;
    } catch (tmv::Error& e) {
```

```

std::cout<<"Caught e = "<<e;
    /// Caught e = TMV Read Error: Reading from list initialization.
    /// List has more elements than expected.
}

// Norms, etc.
std::cout<<"Norm1(v3) = "<<Norm1(v3);
std::cout<<" = "<<v3.sumAbsElements()<<std::endl;
/// Norm1(v3) = 44.9 = 44.9

std::cout<<"Norm2(v3) = "<<Norm2(v3);
std::cout<<" = "<<v3.norm()<<std::endl;
/// Norm2(v3) = 21.9123 = 21.9123

std::cout<<"NormInf(v3) = "<<NormInf(v3);
std::cout<<" = "<<v3.maxAbsElement()<<std::endl;
/// NormInf(v3) = 15 = 15

std::cout<<"SumElements(v3) = "<<SumElements(v3)<<std::endl;
/// SumElements(v3) = -9.1

// Min/Max elements:
int i1,i2,i3,i4;
double x1 = v3.minAbsElement(&i1);
double x2 = v3.maxAbsElement(&i2);
double x3 = v3.minElement(&i3);
double x4 = v3.maxElement(&i4);
std::cout<<"|v3("<<i1<<")| = "<<x1<<" is the minimum absolute value\n";
/// |v3(0)| = 1.1 is the minimum absolute value
std::cout<<"|v3("<<i2<<")| = "<<x2<<" is the maximum absolute value\n";
/// |v3(5)| = 15 is the maximum absolute value
std::cout<<"v3("<<i3<<") = "<<x3<<" is the minimum value\n";
/// v3(4) = -15 is the minimum value
std::cout<<"v3("<<i4<<") = "<<x4<<" is the maximum value\n";
/// v3(5) = 8 is the maximum value

// Modifications:

std::cout<<"v1 = "<<v1<<std::endl;
/// v1 = 6 ( 1 2 4 8 16 32 )

v1.addToAll(5.);
std::cout<<"v1.addToAll(5.) = "<<v1<<std::endl;
/// v1.addToAll(5.) = 6 ( 6 7 9 13 21 37 )

v1.reverseSelf();
std::cout<<"v1.reverseSelf() = "<<v1<<std::endl;
/// v1.reverseSelf() = 6 ( 37 21 13 9 7 6 )

v1.setZero();
std::cout<<"v1.setZero() = "<<v1<<std::endl;
/// v1.setZero() = 6 ( 0 0 0 0 0 0 )

v1.setAllTo(20.);
std::cout<<"v1.setAllTo(20.) = "<<v1<<std::endl;

```

```

!!! v1.setAllTo(20.) = 6 ( 20 20 20 20 20 20 )

v1.makeBasis(2);
std::cout<<"v1.makeBasis(2) = "<<v1<<std::endl;
!!! v1.makeBasis(2) = 6 ( 0 0 1 0 0 0 )

// Views:

std::cout<<"v3 = "<<v3<<std::endl;
!!! v3 = 6 ( 1.1 8 -15 2.5 6.3 -12 )
std::cout<<"v3.subVector(0,3) = "<<v3.subVector(0,3)<<std::endl;
!!! v3.subVector(0,3) = 3 ( 1.1 8 -15 )
std::cout<<"v3.subVector(0,6,2) = "<<v3.subVector(0,6,2)<<std::endl;
!!! v3.subVector(0,6,2) = 3 ( 1.1 -15 6.3 )
std::cout<<"v3.reverse() = "<<v3.reverse()<<std::endl;
// v3.reverse() = 6 ( -12 6.3 2.5 -15 8 1.1 )

// Views can be initialized with << too.
v3.reverse() << 1.2, 9., 12, 2.5, -7.4, 14;
std::cout<<"v3 = "<<v3<<std::endl;
!!! v3 = 6 ( 14 -7.4 2.5 12 9 1.2 )
v3.subVector(0,6,2) << 18, 22, 33;
std::cout<<"v3 = "<<v3<<std::endl;
!!! v3 = 6 ( 18 -7.4 22 12 33 1.2 )

// Can use the views within expressions
v3.reverse() += v4;
std::cout<<"v3.reverse() += v4 => v3 = "<<v3<<std::endl;
!!! v3.reverse() += v4 => v3 = 6 ( 32 -14.8 24.5 24 42 2.4 )

v3.subVector(0,3) *= 2.;
std::cout<<"v3.subVector(0,3) *= 2 => v3 = "<<v3<<std::endl;
!!! v3.subVector(0,3) *= 2 => v3 = 6 ( 64 -29.6 49 24 42 2.4 )

// Fortran Indexing:

tmv::Vector<double,tmv::FortranStyle> fv3 = v3;

std::cout<<"fv3 = v3 = "<<fv3<<std::endl;
!!! fv3 = v3 = 6 ( 64 -29.6 49 24 42 2.4 )
std::cout<<"fv3(1) = "<<fv3(1)<<std::endl;
!!! fv3(1) = 64
std::cout<<"fv3(6) = "<<fv3(6)<<std::endl;
!!! fv3(6) = 2.4
std::cout<<"fv3.subVector(1,3) = "<<fv3.subVector(1,3)<<std::endl;
!!! fv3.subVector(1,3) = 3 ( 64 -29.6 49 )
std::cout<<"fv3.makeBasis(3) = "<<fv3.makeBasis(3)<<std::endl;
!!! fv3.makeBasis(2) = 6 ( 0 0 1 0 0 0 )

// Vector arithmetic:

tmv::Vector<double> v3pv4 = v3 + v4;
std::cout<<"v3 + v4 = "<<v3pv4<<std::endl;

```

```

!!! v3 + v4 = 6 ( 65.2 -20.6 61 26.5 34.6 16.4 )

// Inner product
double v3v4 = v3 * v4;
std::cout<<"v3 * v4 = "<<v3v4<<std::endl;
!!! v3 * v4 = 181.2

v3 *= 2.;
std::cout<<"v3 *= 2 = "<<v3<<std::endl;
!!! v3 *= 2 = 6 ( 128 -59.2 98 48 84 4.8 )

v3 += v4;
std::cout<<"v3 += v4 = "<<v3<<std::endl;
!!! v3 += v4 = 6 ( 129.2 -50.2 110 50.5 76.6 18.8 )

// Get as complicated as you want:
std::cout<<"(v1*v2) * v3 + ((-v4 + 4.*v1)*v2) * v2/20. = \n"<<
    (v1*v2) * v3 + ((-v4 + 4.*v1)*v2) * v2/20.<<std::endl;
!!! (v1*v2) * v3 + ((-v4 + 4.*v1)*v2) * v2/20. =
!!! 6 ( 252.94 -105.86 214.54 95.54 147.74 32.14 )

// Automatically checks for aliases:
v3 = v4 - 3.*v3;
std::cout<<"v3 = v4-3.*v3 => v3 = "<<v3<<std::endl;
!!! v3 = v4-3.*v3 => v3 = 6 ( -386.4 159.6 -318 -149 -237.2 -42.4 )

// Complex vectors:

tmv::Vector<std::complex<double>> cv4 = v4 * std::complex<double>(1,2);
std::cout<<"cv4 = v4 * (1+2i) = \n"<<cv4<<std::endl;
!!! cv4 = v4 * (1+2i) =
!!! 6 ( (1.2,2.4) (9,18) (12,24) (2.5,5) (-7.4,-14.8) (14,28) )

std::cout<<"cv4.conjugate() = \n"<<cv4.conjugate()<<std::endl;
!!! cv4.conjugate() =
!!! 6 ( (1.2,-2.4) (9,-18) (12,-24) (2.5,-5) (-7.4,14.8) (14,-28) )
std::cout<<"cv4.realPart() = "<<cv4.realPart()<<std::endl;
!!! cv4.realPart() = 6 ( 1.2 9 12 2.5 -7.4 14 )
std::cout<<"cv4.imagPart() = "<<cv4.imagPart()<<std::endl;
!!! cv4.imagPart() = 6 ( 2.4 18 24 5 -14.8 28 )
std::cout<<"Norm(cv4) = "<<Norm(cv4)<<std::endl;
!!! Norm(cv4) = 49.1655
std::cout<<"sqrt(cv4*cv4.conjugate()) = "<<
    sqrt(cv4*cv4.conjugate())<<std::endl;
!!! sqrt(cv4*cv4.conjugate()) = (49.1655,0)
std::cout<<"cv4.maxAbsElement() = "<<cv4.maxAbsElement()<<std::endl;
!!! cv4.maxAbsElement() = 31.305

// Can mix real and complex in any combination
std::cout<<"cv4 - v4 = "<<cv4 - v4<<std::endl;
!!! cv4 - v4 = 6 ( (0,2.4) (0,18) (0,24) (0,5) (0,-14.8) (0,28) )
std::cout<<"cv4 * v4 * (1-2i) = "<<
    cv4 * v4 * std::complex<double>(1,-2)<<std::endl;
!!! cv4 * v4 * (1-2i) = (2417.25,0)

```

```

// Sorting:

v4 << 2, 5.3, -1.5, -7, 0.5, -2.8;
std::cout<<"v4 = "<<v4<<std::endl;
/// v4 = 6 ( 2 5.3 -1.5 -7 0.5 -2.8 )
int p[6];
v4.sort(p);
std::cout<<"Sorted: v4 = "<<v4<<std::endl;
/// Sorted: v4 = 6 ( -7 -2.8 -1.5 0.5 2 5.3 )
v4.reversePermute(p);
std::cout<<"Sort undone: v4 = "<<v4<<std::endl;
/// Sort undone: v4 = 6 ( 2 5.3 -1.5 -7 0.5 -2.8 )
v4.sort(); // Don't necessarily need p.
std::cout<<"Resorted: v4 = "<<v4<<std::endl;
/// Resorted: v4 = 6 ( -7 -2.8 -1.5 0.5 2 5.3 )

// Can sort by other criteria:
// (Note: the 0 here is p. If p=0, then it is not set.)
std::cout<<"v4.sort(0,Descend) = "<<v4.sort(0,tmv::Descend)<<std::endl;
/// v4.sort(0,Descend) = 6 ( 5.3 2 0.5 -1.5 -2.8 -7 )

cv4.realPart() << -3, 1, -2, -1, 7, 3;
cv4.imagPart() << 4, -1, 0, -6, 5, -1;
// (I find this complex initialization to be more readable than a list
// filled with values that look like complex<double>(-3,4), ...)
std::cout<<"cv4 = "<<cv4<<std::endl;
/// cv4 = 6 ( (-3,4) (1,-1) (-2,0) (-1,-6) (7,5) (3,-1) )

std::cout<<"cv4.sort(0,Descend,RealComp) = \n"<<
cv4.sort(0,tmv::Descend,tmv::RealComp)<<std::endl;
/// cv4.sort(0,Descend,RealComp) =
/// 6 ( (7,5) (3,-1) (1,-1) (-1,-6) (-2,0) (-3,4) )
std::cout<<"cv4.sort(0,Ascend,ImagComp) = \n"<<
cv4.sort(0,tmv::Ascend,tmv::ImagComp)<<std::endl;
/// cv4.sort(0,Ascend,ImagComp) =
/// 6 ( (-1,-6) (3,-1) (1,-1) (-2,0) (-3,4) (7,5) )
std::cout<<"cv4.sort(0,Ascend,AbsComp) = \n"<<
cv4.sort(0,tmv::Ascend,tmv::AbsComp)<<std::endl;
/// cv4.sort(0,Ascend,AbsComp) =
/// 6 ( (1,-1) (-2,0) (3,-1) (-3,4) (-1,-6) (7,5) )
std::cout<<"cv4.sort(0,Ascend,ArgComp) = \n"<<
cv4.sort(0,tmv::Ascend,tmv::ArgComp)<<std::endl;
/// cv4.sort(0,Ascend,ArgComp) =
/// 6 ( (-1,-6) (1,-1) (3,-1) (7,5) (-3,4) (-2,0) )

// The default component for complex vectors is RealComp:
std::cout<<"cv4.sort() = \n"<< cv4.sort()<<std::endl;
/// cv4.sort() =
/// 6 ( (-3,4) (-2,0) (-1,-6) (1,-1) (3,-1) (7,5) )

return 0;
} catch (tmv::Error& e) {
std::cerr<<e<<std::endl;
return 1;
}

```

13.2 Matrix

```
#include "TMV.h"
#include <iostream>

int main() try
{
    // Several ways to create and initialize matrices:

    // Create with uninitialized values
    tmv::Matrix<double> m1(4,3);
    // In debugging mode (the default), these are all initialized to 888
    // to help you more easily notice when you have failed to correctly
    // initialize a matrix.
    std::cout<<"m1 = \n"<<m1<<std::endl;
    /// m1 =
    /// 4 3
    /// ( 888 888 888 )
    /// ( 888 888 888 )
    /// ( 888 888 888 )
    /// ( 888 888 888 )

    // Initialize with direct element access:
    for(size_t i=0;i<m1.nrows();i++)
        for(size_t j=0;j<m1.ncols();j++)
            m1(i,j) = 2.*i-3.*j+10.;
    std::cout<<"m1 = \n"<<m1<<std::endl;
    /// m1 =
    /// 4 3
    /// ( 10 7 4 )
    /// ( 12 9 6 )
    /// ( 14 11 8 )
    /// ( 16 13 10 )

    // Create with all 2's.
    tmv::Matrix<double> m2(4,3,2.);
    std::cout<<"m2 = \n"<<m2<<std::endl;
    /// m2 =
    /// 4 3
    /// ( 2 2 2 )
    /// ( 2 2 2 )
    /// ( 2 2 2 )
    /// ( 2 2 2 )

    // Create from given elements in a C array;
    double mm[12] = {1,2,3,4, 5,6,7,8, 9,10,11,12};
    tmv::Matrix<double> m3(4,3,mm); // Default order is ColMajor
    std::cout<<"m3 (ColMajor) = \n"<<m3<<std::endl;
    /// m3 (ColMajor) =
    /// 4 3
    /// ( 1 5 9 )
    /// ( 2 6 10 )
    /// ( 3 7 11 )
    /// ( 4 8 12 )

    tmv::Matrix<double, tmv::RowMajor> m4(4,3,mm);
```

```

std::cout<<"m4 (RowMajor) = \n"<<m4<<std::endl;
///! m4 (RowMajor) =
///! 4 3
///! ( 1 2 3 )
///! ( 4 5 6 )
///! ( 7 8 9 )
///! ( 10 11 12 )

// Initialize with comma-delimited list.
// Note that this is more intuitive with RowMajor Storage
m3 <<
    2, -5, 1, -3,
    -4, 8, -7, 1,
    -1, 2, 4, 0;
std::cout<<"m3 (ColMajor) => \n"<<m3<<std::endl;
///! m3 (ColMajor) =>
///! 4 3
///! ( 2 -4 -1 )
///! ( -5 8 2 )
///! ( 1 -7 4 )
///! ( -3 1 0 )

m4 <<
    2, -4, 1,
    -5, 8, 2,
    1, -7, 4,
    -3, 1, 0;
std::cout<<"m4 (RowMajor) => \n"<<m4<<std::endl;
///! m4 (RowMajor) =>
///! 4 3
///! ( 2 -4 1 )
///! ( -5 8 2 )
///! ( 1 -7 4 )
///! ( -3 1 0 )

// Create from STL vector of vectors
std::vector<std::vector<double> > mm2(3, std::vector<double>(3));
for( size_t i=0;i<3;i++)
    for( size_t j=0;j<3;j++)
        mm2[i][j] = 2.*i+j-3.*i*j;
tmv::Matrix<double> m5(mm2);
std::cout<<"m5 = \n"<<m5<<std::endl;
///! m5 =
///! 3 3
///! ( 0 1 2 )
///! ( 2 0 -2 )
///! ( 4 -1 -6 )

// Norms, etc.

std::cout<<"Norm1(m1) = "<<Norm1(m1)<<std::endl;
///! Norm1(m1) = 52
std::cout<<"Norm2(m1) = "<<Norm2(m1)<<std::endl;
///! Norm2(m1) = Warning:
///! Calling Matrix::Norm2 without previously calling DivideUsing(SV)

```

```

!!! 36.452
// This is potentially inefficient, especially for large matrices.
// To avoid the warning, you can use:
std::cout<<"m1.doNorm2() = "<<m1.doNorm2()<<std::endl;
!!! Norm2(m1) = 36.452
std::cout<<"NormInf(m1) = "<<NormInf(m1)<<std::endl;
!!! NormInf(m1) = 39
std::cout<<"NormF(m1) = "<<NormF(m1);
std::cout<<" = "<<Norm(m1)<<std::endl;
!!! NormF(m1) = 36.4966 = 36.4966
std::cout<<"MaxAbsElement(m1) = "<<MaxAbsElement(m1)<<std::endl;
!!! MaxAbsElement(m1) = 16
std::cout<<"Trace(m5) = "<<Trace(m5)<<std::endl;
!!! Trace(m5) = -6
std::cout<<"Det(m5) = "<<Det(m5)<<std::endl;
!!! Det(m5) = 0

// Views:

std::cout<<"m1 = \n"<<m1<<std::endl;
!!! m1 =
!!! 4 3
!!! ( 10 7 4 )
!!! ( 12 9 6 )
!!! ( 14 11 8 )
!!! ( 16 13 10 )
std::cout<<"m1.row(1) = "<<m1.row(1)<<std::endl;
!!! m1.row(1) = 3 ( 12 9 6 )
std::cout<<"m1.col(2) = "<<m1.col(2)<<std::endl;
!!! m1.col(2) = 4 ( 4 6 8 10 )
std::cout<<"m1.diag() = "<<m1.diag()<<std::endl;
!!! m1.diag() = 3 ( 10 9 8 )
std::cout<<"m1.diag(1) = "<<m1.diag(1)<<std::endl;
!!! m1.diag(1) = 2 ( 7 6 )
std::cout<<"m1.diag(-1) = "<<m1.diag(-1)<<std::endl;
!!! m1.diag(-1) = 3 ( 12 11 10 )
std::cout<<"m1.subMatrix(2,4,0,2) = \n"<<m1.subMatrix(2,4,0,2)<<std::endl;
!!! m1.subMatrix(2,4,0,2) =
!!! 2 2
!!! ( 14 11 )
!!! ( 16 13 )
std::cout<<"m1.subMatrix(0,4,1,3,2,1) = \n"<<
    m1.subMatrix(0,4,1,3,2,1)<<std::endl;
!!! m1.subMatrix(0,4,1,3,2,1) =
!!! 2 2
!!! ( 7 4 )
!!! ( 11 8 )
std::cout<<"m1.transpose() = \n"<<m1.transpose()<<std::endl;
!!! m1.transpose() =
!!! 3 4
!!! ( 10 12 14 16 )
!!! ( 7 9 11 13 )
!!! ( 4 6 8 10 )
std::cout<<"m1.colRange(1,3) = \n"<<m1.colRange(1,3)<<std::endl;
!!! m1.colRange(1,3) =

```



```

    ///! 4 2
    ///! ( 7 4 )
    ///! ( 9 6 )
    ///! ( 11 8 )
    ///! ( 13 10 )
    std::cout<<"m1.rowPair(3,0) = \n"<<m1.rowPair(3,0)<<std::endl;
    ///! m1.rowPair(3,0) =
    ///! 2 3
    ///! ( 16 13 10 )
    ///! ( 10 7 4 )
    // Can use the views within expressions
    m1.rowRange(0,3) += m5.transpose();
    std::cout<<"m1.rowRange(0,3) += m5.transpose() => m1 = \n"<<m1<<std::endl;
    ///! m1.rowRange(0,3) += m5.transpose() => m1 =
    ///! 4 3
    ///! ( 10 9 8 )
    ///! ( 13 9 5 )
    ///! ( 16 9 2 )
    ///! ( 16 13 10 )
    m1.row(0) *= 2.;
    std::cout<<"m1.row(0) *= 2 => m1 = \n"<<m1<<std::endl;
    ///! m1.row(0) *= 2 => m1 =
    ///! 4 3
    ///! ( 20 18 16 )
    ///! ( 13 9 5 )
    ///! ( 16 9 2 )
    ///! ( 16 13 10 )

    // Fortran Indexing:

    tmv::Matrix<double,tmv::ColMajor,tmv::FortranStyle> fm1 = m1;
    std::cout<<"fm1 = m1 = \n"<<fm1<<std::endl;
    ///! fm1 = m1 =
    ///! 4 3
    ///! ( 20 18 16 )
    ///! ( 13 9 5 )
    ///! ( 16 9 2 )
    ///! ( 16 13 10 )
    std::cout<<"fm1(1,1) = "<<fm1(1,1)<<std::endl;
    ///! fm1(1,1) = 20
    std::cout<<"fm1(4,3) = "<<fm1(4,3)<<std::endl;
    ///! fm1(4,3) = 10
    std::cout<<"fm1.row(1) = "<<fm1.row(1)<<std::endl;
    ///! fm1.row(1) = 3 ( 20 18 16 )
    std::cout<<"fm1.col(3) = "<<fm1.col(3)<<std::endl;
    ///! fm1.col(3) = 4 ( 16 5 2 10 )
    std::cout<<"fm1.subMatrix(2,3,1,2) = \n"<<fm1.subMatrix(2,3,1,2)<<std::endl;
    ///! fm1.subMatrix(2,3,1,2) =
    ///! 2 2
    ///! ( 13 9 )
    ///! ( 16 9 )
    std::cout<<"fm1.subMatrix(1,3,2,3,2,1) = \n"<<
        fm1.subMatrix(1,3,2,3,2,1)<<std::endl;
    ///! fm1.subMatrix(1,3,2,3,2,1) =
    ///! 2 2

```

```

///! ( 18 16 )
///! ( 9 2 )
std::cout<<"fml.colRange(1,2) = \n"<<fml.colRange(1,2)<<std::endl;
///! fml.colRange(1,2) =
///! 4 2
///! ( 20 18 )
///! ( 13 9 )
///! ( 16 9 )
///! ( 16 13 )
std::cout<<"fml.rowPair(4,1) = \n"<<fml.rowPair(4,1)<<std::endl;
///! fml.rowPair(4,1) =
///! 2 3
///! ( 16 13 10 )
///! ( 20 18 16 )

```

// Matrix arithmetic:

```

tmv::Matrix<double> m1pm3 = m1 + m3;
std::cout<<"m1 + m3 = \n"<<m1pm3<<std::endl;
///! m1 + m3 =
///! 4 3
///! ( 22 14 15 )
///! ( 8 17 7 )
///! ( 17 2 6 )
///! ( 13 14 10 )
// Works correctly even if matrices are stored in different order:
tmv::Matrix<double> m3pm4 = m3 + m4;
std::cout<<"m3 + m4 = \n"<<m3pm4<<std::endl;
///! m3 + m4 =
///! 4 3
///! ( 4 -8 0 )
///! ( -10 16 4 )
///! ( 2 -14 8 )
///! ( -6 2 0 )

```

```

m1 *= 2.;
std::cout<<"m1 *= 2 = \n"<<m1<<std::endl;
///! m1 *= 2 =
///! 4 3
///! ( 40 36 32 )
///! ( 26 18 10 )
///! ( 32 18 4 )
///! ( 32 26 20 )

```

```

m1 += m4;
std::cout<<"m1 += m4 = \n"<<m1<<std::endl;
///! m1 += m4 =
///! 4 3
///! ( 42 32 33 )
///! ( 21 26 12 )
///! ( 33 11 8 )
///! ( 29 27 20 )

```

// Vector outer product

```

tmv::Vector<double> v1 = m4.col(0);

```

```

tmv::Vector<double> v2 = m4.row(1);
tmv::Matrix<double> v1v2 = v1^v2;
std::cout<<"v1 = "<<v1<<std::endl;
///! v1 = 4 ( 2 -5 1 -3 )
std::cout<<"v2 = "<<v2<<std::endl;
///! v2 = 3 ( -5 8 2 )
std::cout<<"v1^v2 = \n"<<v1v2<<std::endl;
///! v1^v2 =
///! 4 3
///! ( -10 16 4 )
///! ( 25 -40 -10 )
///! ( -5 8 2 )
///! ( 15 -24 -6 )
std::cout<<"ColVectorViewOf(v1)*RowVectorViewOf(v2) = \n"<<
    ColVectorViewOf(v1)*RowVectorViewOf(v2);
///! ColVectorViewOf(v1)*RowVectorViewOf(v2) =
///! 4 3
///! ( -10 16 4 )
///! ( 25 -40 -10 )
///! ( -5 8 2 )
///! ( 15 -24 -6 )

// Matrix * Vector product
std::cout<<"m1 * v2 = "<<m1*v2<<std::endl;
///! m1 * v2 = 4 ( 112 127 -61 111 )
std::cout<<"v1 * m1 = "<<v1*m1<<std::endl;
///! v1 * m1 = 3 ( -75 -136 -46 )

// Matrix * Matrix product
tmv::Matrix<double> m1m5 = m1 * m5;
std::cout<<"m1 * m5 = \n"<<m1m5<<std::endl;
///! m1 * m5 =
///! 4 3
///! ( 196 9 -178 )
///! ( 100 9 -82 )
///! ( 54 25 -4 )
///! ( 134 9 -116 )
std::cout<<"m1.row(0) * m5.col(2) = "<<m1.row(0)*m5.col(2)<<std::endl;
///! m1.row(0) * m5.col(2) = -178
std::cout<<"(m1 * m5)(0,2) = "<<m1m5(0,2)<<std::endl;
///! (m1 * m5)(0,2) = -178

// Can handle aliases:
// No alias problem here:
std::cout<<"m1 + 3*m1-m2 = \n"<<m1+3.*m1-m2<<std::endl;
///! m1 + 3*m1-m2 =
///! 4 3
///! ( 166 126 130 )
///! ( 82 102 46 )
///! ( 130 42 30 )
///! ( 114 106 78 )
// But this would be a problem for a naive implementation:
m1 += 3.*m1-m2;
std::cout<<"m1 += 3*m1-m2 = \n"<<m1<<std::endl;
///! m1 += 3*m1-m2 =
///! 4 3

```

```

!!! ( 166 126 130 )
!!! ( 82 102 46 )
!!! ( 130 42 30 )
!!! ( 114 106 78 )

// Again: here is the correct answer:
std::cout<<"m5 * m5 = \n"<<m5*m5<<std::endl;
!!! m5 * m5 =
!!! 3 3
!!! ( 10 -2 -14 )
!!! ( -8 4 16 )
!!! ( -26 10 46 )
// And here, the alias is dealt with correctly:
m5 *= m5;
std::cout<<"m5 *= m5 = \n"<<m5<<std::endl;
!!! m5 *= m5 =
!!! 3 3
!!! ( 10 -2 -14 )
!!! ( -8 4 16 )
!!! ( -26 10 46 )

// Scalars can be treated as a multiple of identity matrix
m5 += 32.;
std::cout<<"m5 += 32 = \n"<<m5<<std::endl;
!!! m5 += 32 =
!!! 3 3
!!! ( 42 -2 -14 )
!!! ( -8 36 16 )
!!! ( -26 10 78 )

// Complex matrices:

tmv::Matrix<std::complex<double>> cm4 = m4 * std::complex<double>(1,2);
std::cout<<"cm4 = m4 * (1+2i) = \n"<<cm4<<std::endl;
!!! cm4 = m4 * (1+2i) =
!!! 4 3
!!! ( (2,4) (-4,-8) (1,2) )
!!! ( (-5,-10) (8,16) (2,4) )
!!! ( (1,2) (-7,-14) (4,8) )
!!! ( (-3,-6) (1,2) (0,0) )
std::cout<<"cm4.conjugate() = \n"<<cm4.conjugate()<<std::endl;
!!! cm4.conjugate() =
!!! 4 3
!!! ( (2,-4) (-4,8) (1,-2) )
!!! ( (-5,10) (8,-16) (2,-4) )
!!! ( (1,-2) (-7,14) (4,-8) )
!!! ( (-3,6) (1,-2) (0,-0) )
std::cout<<"cm4.transpose() = \n"<<cm4.transpose()<<std::endl;
!!! cm4.transpose() =
!!! 3 4
!!! ( (2,4) (-5,-10) (1,2) (-3,-6) )
!!! ( (-4,-8) (8,16) (-7,-14) (1,2) )
!!! ( (1,2) (2,4) (4,8) (0,0) )
std::cout<<"cm4.adjoint() = \n"<<cm4.adjoint()<<std::endl;
!!! cm4.adjoint() =

```

```

    ///! 3 4
    ///! ( (2,-4) (-5,10) (1,-2) (-3,6) )
    ///! ( (-4,8) (8,-16) (-7,14) (1,-2) )
    ///! ( (1,-2) (2,-4) (4,-8) (0,-0) )
    std::cout<<"cm4.realPart() = \n"<<cm4.realPart()<<std::endl;
    ///! cm4.realPart() =
    ///! 4 3
    ///! ( 2 -4 1 )
    ///! ( -5 8 2 )
    ///! ( 1 -7 4 )
    ///! ( -3 1 0 )
    std::cout<<"cm4.imagPart() = \n"<<cm4.imagPart()<<std::endl;
    ///! cm4.imagPart() =
    ///! 4 3
    ///! ( 4 -8 2 )
    ///! ( -10 16 4 )
    ///! ( 2 -14 8 )
    ///! ( -6 2 0 )
    std::cout<<"Norm(cm4) = "<<Norm(cm4)<<std::endl;
    ///! Norm(cm4) = 30.8221
    std::cout<<"cm4*cm4.adjoint() = \n"<<cm4*cm4.adjoint()<<std::endl;
    ///! cm4*cm4.adjoint() =
    ///! 4 4
    ///! ( (105,0) (-200,0) (170,0) (-50,0) )
    ///! ( (-200,0) (465,0) (-265,0) (115,0) )
    ///! ( (170,0) (-265,0) (330,0) (-50,0) )
    ///! ( (-50,0) (115,0) (-50,0) (50,0) )

    // Can mix real and complex in any combination
    std::cout<<"cm4 - m4 = \n"<<cm4 - m4<<std::endl;
    ///! cm4 - m4 =
    ///! 4 3
    ///! ( (0,4) (0,-8) (0,2) )
    ///! ( (0,-10) (0,16) (0,4) )
    ///! ( (0,2) (0,-14) (0,8) )
    ///! ( (0,-6) (0,2) (0,0) )
    std::cout<<"cm4 * m4.transpose() * (1-2i) = \n"<<
        cm4 * m4.transpose() * std::complex<double>(1,-2)<<std::endl;
    ///! cm4 * m4.transpose() * (1-2i) =
    ///! 4 4
    ///! ( (105,0) (-200,0) (170,0) (-50,0) )
    ///! ( (-200,0) (465,0) (-265,0) (115,0) )
    ///! ( (170,0) (-265,0) (330,0) (-50,0) )
    ///! ( (-50,0) (115,0) (-50,0) (50,0) )

    return 0;
} catch (tmv::Error& e) {
    std::cerr<<e<<std::endl;
    return 1;
}

```

13.3 Division

```
#include "TMV.h"
#include "TMV_Diag.h"
#include <iostream>

int main() try
{
    tmv::Matrix<double> A(4,4);
    for( size_t i=0;i<A.nrows();i++)
        for( size_t j=0;j<A.ncols();j++)
            A(i,j) = 6.*i-2*j*j+2.;
    A.diag().addToAll(5.);

    tmv::Vector<double> b(4);
    for( size_t i=0;i<b.size();i++)
        b(i) = 3.+2.*i;

    // Basic Ax=b solution:
    std::cout<<"A = \n"<<A<<std::endl;
    /// A =
    /// 4 4
    /// ( 7 0 -6 -16 )
    /// ( 8 11 0 -10 )
    /// ( 14 12 11 -4 )
    /// ( 20 18 12 7 )
    std::cout<<"b = "<<b<<std::endl;
    /// b = 4 ( 3 5 7 9 )

    // Solve: Ax = b
    tmv::Vector<double> x = b/A; // Default: use LU decomposition
    std::cout<<"x = b/A = "<<x<<std::endl;
    /// x = b/A = 4 ( 0.294545 0.170909 0.0472727 -0.0763636 )
    std::cout<<"Check: A*x = "<<A*x<<std::endl;
    /// Check: A*x = 4 ( 3 5 7 9 )

    // Can update A and then re-solve:
    A(0,0) = 20.;
    x = b/A;
    std::cout<<"Now x = b/A = "<<x<<std::endl;
    /// Now x = b/A = 4 ( 0.133458 0.29877 0.117091 -0.064587 )
    std::cout<<"A*x = "<<A*x<<std::endl;
    /// A*x = 4 ( 3 5 7 9 )

    // If the matrix won't change, but you want to solve with multiple
    // vectors, then it's faster to let TMV know this with saveDiv().
    tmv::Matrix<double> A2 = A;
    A2.saveDiv();
    x = b/A2;
    std::cout<<"x1 = "<<x<<std::endl;
    /// x1 = 4 ( 0.133458 0.29877 0.117091 -0.064587 )
    std::cout<<"A*x = "<<A2*x<<std::endl;
    /// A*x = 4 ( 3 5 7 9 )
    x = b.reverse()/A2; // Fast, since doesn't recalculate LU decomposition.
    std::cout<<"x2 = "<<x<<std::endl;
    /// x2 = 4 ( 0.136753 0.207381 -0.0775483 -0.362478 )
}
```

```

std::cout<<"A*x2 = "<<A2*x<<std::endl;
/// A*x2 = 4 ( 9 7 5 3 )
A2.row(0) *= 2.;
x = b/A2; // Wrong, since doesn't recalculate LU decomposition.
std::cout<<"Wrong x = "<<x<<std::endl;
/// Wrong x = 4 ( 0.133458 0.29877 0.117091 -0.064587 )
std::cout<<"A*x = "<<A2*x<<std::endl;
/// A*x = 4 ( 6 5 7 9 )
// If the matrix does change when saveDiv() is set,
// you can manually recalculate the decomposition:
A2.resetDiv();
x = b/A2; // Now it is correct.
std::cout<<"x = "<<x<<std::endl;
/// x = 4 ( 0.0703537 0.348858 0.144442 -0.0599736 )
std::cout<<"A*x = "<<A2*x<<std::endl;
/// A*x = 4 ( 3 5 7 9 )

// Matrix inverse:
tmv::Matrix<double> A2inv = A2.inverse();
std::cout<<"Ainv = \n"<<A2inv<<std::endl;
/// Ainv =
/// 4 4
/// ( 0.0210347 -0.0395431 -0.012522 0.0325132 )
/// ( -0.016696 0.0966608 -0.0527241 0.0316344 )
/// ( -0.00911687 -0.0643234 0.139631 -0.0537786 )
/// ( -0.00153779 -0.0253076 -0.0680141 0.0608084 )
std::cout<<"Ainv*A = \n"<<A2inv*A2<<std::endl;
/// Ainv*A =
/// 4 4
/// ( 1 0 0 1.66533e-16 )
/// ( -2.22045e-16 1 -5.55112e-17 -2.77556e-17 )
/// ( 0 0 1 3.88578e-16 )
/// ( 2.22045e-16 -2.22045e-16 -2.22045e-16 1 )
// This is a case where it can be useful to see the
// matrix elements that are larger than some threshold value:
std::cout<<"Ainv*A = \n";
(A2inv*A2).write(std::cout,1.e-8);
/// Ainv*A =
/// 4 4
/// ( 1 0 0 0 )
/// ( 0 1 0 0 )
/// ( 0 0 1 0 )
/// ( 0 0 0 1 )

// 1/x notation is treated as arithmetic:
// (But the 1 has to be the same type as the elements of the matrix.
/// In this case, double. With a float matrix, use 1.F instead.)
std::cout<<"1./A = \n"<<1./A2<<std::endl;
/// 1./A =
/// 4 4
/// ( 0.0210347 -0.0395431 -0.012522 0.0325132 )
/// ( -0.016696 0.0966608 -0.0527241 0.0316344 )
/// ( -0.00911687 -0.0643234 0.139631 -0.0537786 )
/// ( -0.00153779 -0.0253076 -0.0680141 0.0608084 )
std::cout<<"5./A = \n"<<5./A2<<std::endl;
/// 5./A =

```

```

//! 4 4
//! ( 0.105174 -0.197715 -0.0626098 0.162566 )
//! ( -0.0834798 0.483304 -0.26362 0.158172 )
//! ( -0.0455844 -0.321617 0.698155 -0.268893 )
//! ( -0.00768893 -0.126538 -0.34007 0.304042 )

// Can also use inverse() notation instead of /
// x = b/A2 inlines to exactly the same thing as x = A2.inverse() * b;
// ie. accurate back-substitution methods are used rather than
// actually computing the inverse and then multiplying:
x = A2.inverse() * b;
std::cout<<"x = A.inverse() * b = "<<x<<std::endl;
//! x = A.inverse() * b = 4 ( 0.0703537 0.348858 0.144442 -0.0599736 )

// Division from the other side can either be done with this
// inverse() notation or with the % operator:
// This is the solution to the equation x A = b, rather than A x = b.
x = b * A2.inverse();
std::cout<<"x = b * A.inverse() = "<<x<<std::endl;
//! x = b * A.inverse() = 4 ( -0.0980338 -0.313357 0.0641037 0.426538 )
x = b % A2;
std::cout<<"x = b % A = "<<x<<std::endl;
//! x = b % A = 4 ( -0.0980338 -0.313357 0.0641037 0.426538 )
std::cout<<"Check: x*A = "<<x*A2<<std::endl;
//! Check: x*A = 4 ( 3 5 7 9 )

tmv::Matrix<double> B(4,3);
for(size_t i=0;i<B.nrows();i++)
    for(size_t j=0;j<B.ncols();j++)
        B(i,j) = 1.+2.*i+j*j;
// Multiple right hand sides may be calculated at once if
// B is a matrix, rather than a vector:
std::cout<<"B = \n"<<B<<std::endl;
//! B =
//! 4 3
//! ( 1 2 5 )
//! ( 3 4 7 )
//! ( 5 6 9 )
//! ( 7 8 11 )
tmv::Matrix<double> X = B/A2;
std::cout<<"X = B/A = \n"<<X<<std::endl;
//! X = B/A =
//! 4 3
//! ( 0.067388 0.0688708 0.0733194 )
//! ( 0.231107 0.289982 0.466608 )
//! ( 0.119618 0.13203 0.169266 )
//! ( 0.0081283 -0.0259227 -0.128076 )
std::cout<<"AX = "<<A2*X<<std::endl;
//! AX = 4 3
//! ( 1 2 5 )
//! ( 3 4 7 )
//! ( 5 6 9 )
//! ( 7 8 11 )

// And as always, you can mix complex and real objects:
tmv::Vector<std::complex<double>> cb = b * std::complex<double>(3,-2);

```



```

cb(1) = std::complex<double>(-1,8);
cb(2) = std::complex<double>(1,1);
std::cout<<"cb = "<<cb<<std::endl;
/// cb = 4 ( (9,-6) (-1,8) (1,1) (27,-18) )
tmv::Vector<std::complex<double>> cx = cb/A;
std::cout<<"cx = cb/A = \n"<<cx<<std::endl;
/// cx = cb/A =
/// 4 ( (1.2835,-1.16652) (0.404218,0.351494) (-1.41217,0.70246)
/// (1.57144,-1.34657) )
/// A*cx = 4 ( (9,-6) (-1,8) (1,1) (27,-18) )
tmv::Matrix<std::complex<double>> CA = A * std::complex<double>(5,-2);
CA(1,1) = std::complex<double>(1,6);
CA(2,3) = std::complex<double>(4,-1);
std::cout<<"CA = "<<CA<<std::endl;
/// CA = 4 4
/// ( (100,-40) (0,0) (-30,12) (-80,32) )
/// ( (40,-16) (1,6) (0,0) (-50,20) )
/// ( (70,-28) (60,-24) (55,-22) (4,-1) )
/// ( (100,-40) (90,-36) (60,-24) (35,-14) )
std::cout<<"cx = b/CA = \n"<<(cx=b/CA)<<std::endl;
/// cx = b/CA =
/// 4 ( (-0.0858093,0.219988) (0.270489,-0.423628) (-0.0665359,0.214597)
/// (-0.114638,0.18158) )
std::cout<<"CA*cx = "<<CA*cx<<std::endl;
/// 4 ( (3,3.55271e-15) (5,0) (7,-3.55271e-15) (9,-3.55271e-15) )
std::cout<<"cb/CA = \n"<<(cx=cb/CA)<<std::endl;
/// cx = cb/CA =
/// 4 ( (0.0374442,0.698211) (0.792852,-1.67098) (-0.916414,0.915387)
/// (0.267616,0.555355) )
std::cout<<"CA*cx = "<<CA*cx<<std::endl;
/// 4 ( (9,-6) (-1,8) (1,1) (27,-18) )

// Least-squares solutions:
// If A in the matrix equation A x = b has more rows than columns,
// then there is, in general, no solution to the equation.
// Instead, one is generally looking for the x that comes closest
// to satisfying the equation, in a least-squares sense.
// Specifically, the x for which Norm2(b-Ax) is minimized.
// This is the solution produced by TMV for such matrices.
//
// Here I model a theoretical system for which each observation
// is 5 + 6i - 3i^2 in the absense of measurement errors.
// We observe the actual values which are not quite equal to
// that because of noise. And the goal is to determine the
// coefficients of 1,i,i^2 (5,6,-3) from noisy observations:
tmv::Vector<double> b3(10);
tmv::Matrix<double> A3(10,3);
double errors[10] = {0.01,-0.02,0.02,-0.02,0.00,-0.03,0.01,-0.02,0.03,0.02};
for(int i=0;i<10;i++) {
    b3(i) = 5. + 6.*i - 3.*i*i + errors[i]; // Model of measurements
    A3(i,0) = 1.; // Parameterization of the model...
    A3(i,1) = i;
    A3(i,2) = i*i;
}
double sigma = 0.02; // sigma = estimate of rms errors

```

```

A3 /= sigma;
b3 /= sigma;

tmv::Vector<double> x3 = b3/A3; // Uses QR decomposition by default
std::cout<<"x = "<<x3<<std::endl;
/// x = 3 ( 5.00773  5.98989  -2.99867 )
std::cout<<"A*x => \n"<<A3*x3<<std::endl;
/// A*x =>
/// 10 ( 250.386  399.947  249.64  -200.534  -950.576  -2000.48  -3350.26
/// -4999.91  -6949.42  -9198.8 )
std::cout<<"chisq = NormSq(A*x-b) = "<<NormSq(A3*x3-b3)<<std::endl;
/// chisq = NormSq(A*x-b) = 6.99811
// The expected value for this is 10 observations minus 3 degrees
// of freedom = 7.

// The covariance matrix for x is (A.Transpose() * A)^-1
// This combination is easy to calculate from the QR decomposition
// that TMV has used to do the division. Therefore, we provide
// it as an explicit function:
tmv::Matrix<double> cov(3,3);
A3.makeInverseATA(cov);
std::cout<<"Cov(x) = \n"<<cov<<std::endl;
/// Cov(x) =
/// 3 3
/// ( 0.000247273  -0.000103636  9.09091e-06 )
/// ( -0.000103636  6.62121e-05  -6.81818e-06 )
/// ( 9.09091e-06  -6.81818e-06  7.57576e-07 )

// The singular value decomposition can detect ill-conditioned matrices
// and correct for them.
// For example, if you model the above observations with the
// 1,i,i^2 components as before, but add as well 6*i-5 as a component,
// then that would be degenerate with 1 and i.
// SVD is able to detect this defect and deal with it appropriately:
tmv::Matrix<double> A4(10,4);
A4.colRange(0,3) = A3*sigma;
for(int i=0;i<10;i++) A4(i,3) = 6.*i-5.;
std::cout<<"Now A*sigma = \n"<<A4<<std::endl;
/// Now A*sigma =
/// 10 4
/// ( 1 0 0 -5 )
/// ( 1 1 1 1 )
/// ( 1 2 4 7 )
/// ( 1 3 9 13 )
/// ( 1 4 16 19 )
/// ( 1 5 25 25 )
/// ( 1 6 36 31 )
/// ( 1 7 49 37 )
/// ( 1 8 64 43 )
/// ( 1 9 81 49 )
A4 /= sigma;
try {
    // This may or may not succeed, but if it does, the results will be
    // unusable, typically with values around 1.e13 and such.
    tmv::Vector<double> x4 = b3/A4;
    std::cout<<"Unstable x = b/A = \n"<<x4<<std::endl;

```

```

    /// Unstable x = b/A =
    /// 4 ( 2.8506e+13 -3.42073e+13 -2.9984 5.70121e+12 )
    std::cout<<"A*x = \n"<<A4*x4<<std::endl;
    /// A*x =
    /// 10 ( 248.75 398.438 248.25 -202 -951 -2001 -3351 -5000 -6946
    /// -9200 )
} catch (tmv::Error& e) {
    std::cout<<"Tried x = b/A, but caught error: \n"<<e<<std::endl;
}

// So instead, tell TMV to use SVD for division rather than QR.
A4.divideUsing(tmv::SV);
std::cout<<"Singular values for A are "<<A4.svd().getS().diag()<<std::endl;
/// Singular values for A are 4 ( 7618.67 733.074 116.312 0 )
std::cout<<"Using only the first "<<A4.svd().getKMax()<<" components\n";
/// Using only the first 3 components
tmv::Vector<double> x4 = b3/A4;
std::cout<<"SVD division yields: x = "<<x4<<std::endl;
/// SVD division yields: x = 4 ( 5.88681 4.93498 -2.99867 0.175817 )
std::cout<<"chisq = "<<NormSq(A4*x4-b3);
std::cout<<"      Norm(x) = "<<Norm(x4)<<std::endl;
/// chisq = 6.99811      Norm(x) = 8.24813

// QRP can also give useful results, but isn't quite as flexible
// as SVD:
A4.divideUsing(tmv::QRP);
x4 = b3/A4;
std::cout<<"QRP division yields: x = "<<x4<<std::endl;
std::cout<<"chisq = "<<NormSq(A4*x4-b3);
std::cout<<"      Norm(x) = "<<Norm(x4)<<std::endl;
/// QRP division yields: x = 4 ( 5.00773 5.98989 -2.99867 0 )
/// chisq = 6.99811      Norm(x) = 8.3635

// Note that both methods give answers with an equal chisq, so
// Ax is equally close to b for each but they have different
// specific choices for the degeneracy.
// SVD will give the solution within this degeneracy freedom that
// has the minimum Norm(x), but QRP will be faster —
// significantly so for large matrices.

return 0;
}
catch (tmv::Error& e)
{
    std::cerr<<e<<std::endl;
    return 1;
}

```

13.4 BandMatrix

```
#include "TMV.h"
// Note: extra include file for BandMatrix
#include "TMV_Band.h"
// Also need to link with -ltmv_syband
#include <iostream>

int main() try
{
    // Several ways to create and initialize band matrices:

    // Create with uninitialized values
    tmv::BandMatrix<double> m1(6,6,1,2);
    for( size_t i=0; i<m1.nrows(); i++)
        for( size_t j=0; j<m1.ncols(); j++)
            if (i<=j+m1.nlo() && j<=i+m1.nhi())
                m1(i,j) = 3.*i-j*j+7.;
    std::cout<<"m1 = \n"<<m1<<std::endl;
    /*! m1 =
    /*! 6 6
    /*! ( 7 6 3 0 0 0 )
    /*! ( 10 9 6 1 0 0 )
    /*! ( 0 12 9 4 -3 0 )
    /*! ( 0 0 12 7 0 -9 )
    /*! ( 0 0 0 10 3 -6 )
    /*! ( 0 0 0 0 6 -3 )

    // Create with all 2's.
    tmv::BandMatrix<double> m2(6,6,1,3,2.);
    std::cout<<"m2 = \n"<<m2<<std::endl;
    /*! m2 =
    /*! 6 6
    /*! ( 2 2 2 2 0 0 )
    /*! ( 2 2 2 2 2 0 )
    /*! ( 0 2 2 2 2 2 )
    /*! ( 0 0 2 2 2 2 )
    /*! ( 0 0 0 2 2 2 )
    /*! ( 0 0 0 0 2 2 )

    // A BandMatrix can be non-square:
    tmv::BandMatrix<double> m3(6,8,1,3,2.);
    std::cout<<"m3 = \n"<<m3<<std::endl;
    /*! m3 =
    /*! 6 8
    /*! ( 2 2 2 2 0 0 0 0 )
    /*! ( 2 2 2 2 2 0 0 0 )
    /*! ( 0 2 2 2 2 2 0 0 )
    /*! ( 0 0 2 2 2 2 2 0 )
    /*! ( 0 0 0 2 2 2 2 2 )
    /*! ( 0 0 0 0 2 2 2 2 )

    // Create from given elements:
    double mm[21] = {1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20,21};
    tmv::BandMatrix<double> m4(6,6,2,1,mm); // Default order is ColMajor
    std::cout<<"m4 (ColMajor) = \n"<<m4<<std::endl;
```

```

//! m4 (ColMajor) =
//! 6 6
//! ( 1 4 0 0 0 0 )
//! ( 2 5 8 0 0 0 )
//! ( 3 6 9 12 0 0 )
//! ( 0 7 10 13 16 0 )
//! ( 0 0 11 14 17 20 )
//! ( 0 0 0 15 18 21 )
tmv::BandMatrix<double, tmv::RowMajor> m5(6,6,2,1,mm);
std::cout<<"m5 (RowMajor) = \n"<<m5<<std::endl;
//! m5 (RowMajor) =
//! 6 6
//! ( 1 2 0 0 0 0 )
//! ( 4 5 6 0 0 0 )
//! ( 7 8 9 10 0 0 )
//! ( 0 11 12 13 14 0 )
//! ( 0 0 15 16 17 18 )
//! ( 0 0 0 19 20 21 )
tmv::BandMatrix<double, tmv::DiagMajor> m6(6,6,2,1,mm);
std::cout<<"m6 (DiagMajor) = \n"<<m6<<std::endl;
//! m6 (DiagMajor) =
//! 6 6
//! ( 11 17 0 0 0 0 )
//! ( 6 12 18 0 0 0 )
//! ( 1 7 13 19 0 0 )
//! ( 0 2 8 14 20 0 )
//! ( 0 0 3 9 15 21 )
//! ( 0 0 0 4 10 16 )

// Can make from the banded portion of a regular Matrix:
tmv::Matrix<double> xm(6,6);
for(size_t i=0;i<xm.nrows();i++)
    for(size_t j=0;j<xm.ncols();j++)
        xm(i,j) = 5.*i-j*j+3.;
tmv::BandMatrix<double> m7(xm,3,2);
std::cout<<"m7 = \n"<<m7<<std::endl;
//! m7 =
//! 6 6
//! ( 3 2 -1 0 0 0 )
//! ( 8 7 4 -1 0 0 )
//! ( 13 12 9 4 -3 0 )
//! ( 18 17 14 9 2 -7 )
//! ( 0 22 19 14 7 -2 )
//! ( 0 0 24 19 12 3 )
// Or from a wider BandMatrix:
tmv::BandMatrix<double> m8(m7,3,0);
std::cout<<"m8 = \n"<<m8<<std::endl;
//! m8 =
//! 6 6
//! ( 3 0 0 0 0 0 )
//! ( 8 7 0 0 0 0 )
//! ( 13 12 9 0 0 0 )
//! ( 18 17 14 9 0 0 )
//! ( 0 22 19 14 7 0 )
//! ( 0 0 24 19 12 3 )

```

```

// Shortcuts to Bi- and Tri-diagonal matrices:
tmv::Vector<double> v1(5,1.);
tmv::Vector<double> v2(6,2.);
tmv::Vector<double> v3(5,3.);
tmv::BandMatrix<double> m9 = LowerBiDiagMatrix(v1,v2);
tmv::BandMatrix<double> m10 = UpperBiDiagMatrix(v2,v3);
tmv::BandMatrix<double> m11 = TriDiagMatrix(v1,v2,v3);
std::cout<<"LowerBiDiagMatrix(v1,v2) = \n"<<m9<<std::endl;
//! LowerBiDiagMatrix(v1,v2) =
//! 6 6
//! ( 2 0 0 0 0 0 )
//! ( 1 2 0 0 0 0 )
//! ( 0 1 2 0 0 0 )
//! ( 0 0 1 2 0 0 )
//! ( 0 0 0 1 2 0 )
//! ( 0 0 0 0 1 2 )
std::cout<<"UpperBiDiagMatrix(v2,v3) = \n"<<m10<<std::endl;
//! UpperBiDiagMatrix(v2,v3) =
//! 6 6
//! ( 2 3 0 0 0 0 )
//! ( 0 2 3 0 0 0 )
//! ( 0 0 2 3 0 0 )
//! ( 0 0 0 2 3 0 )
//! ( 0 0 0 0 2 3 )
//! ( 0 0 0 0 0 2 )
std::cout<<"TriDiagMatrix(v1,v2,v3) = \n"<<m11<<std::endl;
//! TriDiagMatrix(v1,v2,v3) =
//! 6 6
//! ( 2 3 0 0 0 0 )
//! ( 1 2 3 0 0 0 )
//! ( 0 1 2 3 0 0 )
//! ( 0 0 1 2 3 0 )
//! ( 0 0 0 1 2 3 )
//! ( 0 0 0 0 1 2 )

// Norms, etc.

std::cout<<"Norm1(m1) = "<<Norm1(m1)<<std::endl;
//! Norm1(m1) = 30
std::cout<<"Norm2(m1) = "<<Norm2(m1)<<std::endl;
//! Norm2(m1) = Warning:
//! Calling BandMatrix::Norm2 without previously calling DivideUsing(SV)
//! 24.0314
std::cout<<"m1.doNorm2() = "<<m1.doNorm2()<<std::endl;
//! m1.doNorm2() = 24.0314
std::cout<<"NormInf(m1) = "<<NormInf(m1)<<std::endl;
//! NormInf(m1) = 28
std::cout<<"NormF(m1) = "<<NormF(m1)<<" = "<<Norm(m1)<<std::endl;
//! NormF(m1) = 32.0312 = 32.0312
std::cout<<"MaxAbsElement(m1) = "<<MaxAbsElement(m1)<<std::endl;
//! MaxAbsElement(m1) = 12
std::cout<<"Trace(m1) = "<<Trace(m1)<<std::endl;
//! Trace(m1) = 32
std::cout<<"Det(m1) = "<<Det(m1)<<std::endl;
//! Det(m1) = 67635

```

```

// Views:

std::cout<<"m1 = \n"<<m1<<std::endl;
///! m1 =
///! 6 6
///! ( 7 6 3 0 0 0 )
///! ( 10 9 6 1 0 0 )
///! ( 0 12 9 4 -3 0 )
///! ( 0 0 12 7 0 -9 )
///! ( 0 0 0 10 3 -6 )
///! ( 0 0 0 0 6 -3 )
std::cout<<"m1.diag() = "<<m1.diag()<<std::endl;
///! m1.diag() = 6 ( 7 9 9 7 3 -3 )
std::cout<<"m1.diag(1) = "<<m1.diag(1)<<std::endl;
///! m1.diag(1) = 5 ( 6 6 4 0 -6 )
std::cout<<"m1.diag(-1) = "<<m1.diag(-1)<<std::endl;
///! m1.diag(-1) = 5 ( 10 12 12 10 6 )
std::cout<<"m1.subBandMatrix(0,3,0,3,1,1) = \n"<<
    m1.subBandMatrix(0,3,0,3,1,1)<<std::endl;
///! m1.subBandMatrix(0,3,0,3,1,1) =
///! 3 3
///! ( 7 6 0 )
///! ( 10 9 6 )
///! ( 0 12 9 )
std::cout<<"m1.transpose() = \n"<<m1.transpose()<<std::endl;
///! m1.transpose() =
///! 6 6
///! ( 7 10 0 0 0 0 )
///! ( 6 9 12 0 0 0 )
///! ( 3 6 9 12 0 0 )
///! ( 0 1 4 7 10 0 )
///! ( 0 0 -3 0 3 6 )
///! ( 0 0 0 -9 -6 -3 )

// rowRange, colRange shrink both dimensions of the matrix to include only
// the portions that are in those rows or columns:
std::cout<<"m1.rowRange(0,4) = \n"<<m1.rowRange(0,4)<<std::endl;
///! m1.rowRange(0,4) =
///! 4 6
///! ( 7 6 3 0 0 0 )
///! ( 10 9 6 1 0 0 )
///! ( 0 12 9 4 -3 0 )
///! ( 0 0 12 7 0 -9 )
std::cout<<"m1.colRange(1,4) = \n"<<m1.colRange(1,4)<<std::endl;
///! m1.colRange(1,4) =
///! 5 3
///! ( 6 3 0 )
///! ( 9 6 1 )
///! ( 12 9 4 )
///! ( 0 12 7 )
///! ( 0 0 10 )
std::cout<<"m1.diagRange(0,2) = \n"<<m1.diagRange(0,2)<<std::endl;
///! m1.diagRange(0,2) =
///! 6 6

```

```

///! ( 7 6 0 0 0 0 )
///! ( 0 9 6 0 0 0 )
///! ( 0 0 9 4 0 0 )
///! ( 0 0 0 7 0 0 )
///! ( 0 0 0 0 3 -6 )
///! ( 0 0 0 0 0 -3 )
std::cout<<"m1.diagRange(-1,1) = \n"<<m1.diagRange(-1,1)<<std::endl;
///! m1.diagRange(-1,1) =
///! 6 6
///! ( 7 0 0 0 0 0 )
///! ( 10 9 0 0 0 0 )
///! ( 0 12 9 0 0 0 )
///! ( 0 0 12 7 0 0 )
///! ( 0 0 0 10 3 0 )
///! ( 0 0 0 0 6 -3 )

```

// Fortran Indexing:

```

tmv::BandMatrix<double,tmv::ColMajor,tmv::FortranStyle> fml = m1;
std::cout<<"fml = m1 = \n"<<fml<<std::endl;
///! fml = m1 =
///! 6 6
///! ( 7 6 3 0 0 0 )
///! ( 10 9 6 1 0 0 )
///! ( 0 12 9 4 -3 0 )
///! ( 0 0 12 7 0 -9 )
///! ( 0 0 0 10 3 -6 )
///! ( 0 0 0 0 6 -3 )
std::cout<<"fml(1,1) = "<<fml(1,1)<<std::endl;
///! fml(1,1) = 7
std::cout<<"fml(4,3) = "<<fml(4,3)<<std::endl;
///! fml(4,3) = 12
std::cout<<"fml.subBandMatrix(1,3,1,3,1,1) = \n"<<
    fml.subBandMatrix(1,3,1,3,1,1)<<std::endl;
///! fml.subBandMatrix(1,3,1,3,1,1) =
///! 3 3
///! ( 7 6 0 )
///! ( 10 9 6 )
///! ( 0 12 9 )
std::cout<<"fml.rowRange(1,4) = \n"<<fml.rowRange(1,4)<<std::endl;
///! fml.rowRange(1,4) =
///! 4 6
///! ( 7 6 3 0 0 0 )
///! ( 10 9 6 1 0 0 )
///! ( 0 12 9 4 -3 0 )
///! ( 0 0 12 7 0 -9 )
std::cout<<"fml.colRange(2,4) = \n"<<fml.colRange(2,4)<<std::endl;
///! fml.colRange(2,4) =
///! 5 3
///! ( 6 3 0 )
///! ( 9 6 1 )
///! ( 12 9 4 )
///! ( 0 12 7 )
///! ( 0 0 10 )
std::cout<<"fml.diagRange(0,1) = \n"<<fml.diagRange(0,1)<<std::endl;

```



```

//! fm1.diagRange(0,1) =
//! 6 6
//! ( 7 6 0 0 0 0 )
//! ( 0 9 6 0 0 0 )
//! ( 0 0 9 4 0 0 )
//! ( 0 0 0 7 0 0 )
//! ( 0 0 0 0 3 -6 )
//! ( 0 0 0 0 0 -3 )
std::cout<<"fm1.diagRange(-1,0) = \n"<<fm1.diagRange(-1,0)<<std::endl;
//! fm1.diagRange(-1,0) =
//! 6 6
//! ( 7 0 0 0 0 0 )
//! ( 10 9 0 0 0 0 )
//! ( 0 12 9 0 0 0 )
//! ( 0 0 12 7 0 0 )
//! ( 0 0 0 10 3 0 )
//! ( 0 0 0 0 6 -3 )
//!

```

// Matrix arithmetic:

```

tmv::BandMatrix<double> m1pm2 = m1 + m2;
std::cout<<"m1 + m2 = \n"<<m1pm2<<std::endl;
//! m1 + m2 =
//! 6 6
//! ( 9 8 5 2 0 0 )
//! ( 12 11 8 3 2 0 )
//! ( 0 14 11 6 -1 2 )
//! ( 0 0 14 9 2 -7 )
//! ( 0 0 0 12 5 -4 )
//! ( 0 0 0 0 8 -1 )
// Works correctly even if matrices are stored in different order:
tmv::BandMatrix<double> m5pm6 = m5 + m6;
std::cout<<"m5 + m6 = \n"<<m5pm6<<std::endl;
//! m5 + m6 =
//! 6 6
//! ( 12 19 0 0 0 0 )
//! ( 10 17 24 0 0 0 )
//! ( 8 15 22 29 0 0 )
//! ( 0 13 20 27 34 0 )
//! ( 0 0 18 25 32 39 )
//! ( 0 0 0 23 30 37 )
// Also expands the number of off-diagonals appropriately as needed:
tmv::BandMatrix<double> m2pm4 = m2 + m4;
std::cout<<"m2 + m4 = \n"<<m2pm4<<std::endl;
//! m2 + m4 =
//! 6 6
//! ( 3 6 2 2 0 0 )
//! ( 4 7 10 2 2 0 )
//! ( 3 8 11 14 2 2 )
//! ( 0 7 12 15 18 2 )
//! ( 0 0 11 16 19 22 )
//! ( 0 0 0 15 20 23 )

```

m1 *= 2.;

```

std::cout<<"m1 *= 2 = \n"<<m1<<std::endl;
//!  m1 *= 2 =
//!  6 6
//!  ( 14 12 6 0 0 0 )
//!  ( 20 18 12 2 0 0 )
//!  ( 0 24 18 8 -6 0 )
//!  ( 0 0 24 14 0 -18 )
//!  ( 0 0 0 20 6 -12 )
//!  ( 0 0 0 0 12 -6 )

m2 += m1;
std::cout<<"m2 += m1 = \n"<<m2<<std::endl;
//!  m2 += m1 =
//!  6 6
//!  ( 16 14 8 2 0 0 )
//!  ( 22 20 14 4 2 0 )
//!  ( 0 26 20 10 -4 2 )
//!  ( 0 0 26 16 2 -16 )
//!  ( 0 0 0 22 8 -10 )
//!  ( 0 0 0 0 14 -4 )

tmv::Vector<double> v = xm.col(0);
std::cout<<"v = "<<v<<std::endl;
//!  v = 6 ( 3 8 13 18 23 28 )
std::cout<<"m1 * v = "<<m1*v<<std::endl;
//!  m1 * v = 6 ( 216 396 432 60 162 108 )
std::cout<<"v * m1 = "<<v*m1<<std::endl;
//!  v * m1 = 6 ( 202 492 780 832 396 -768 )

// Matrix * matrix product also expands bands appropriately:
tmv::BandMatrix<double> m1m2 = m1 * m2;
std::cout<<"m1 * m2 = \n"<<m1m2<<std::endl;
//!  m1 * m2 =
//!  6 6
//!  ( 488 592 400 136 0 12 )
//!  ( 716 952 704 264 -8 -8 )
//!  ( 528 948 904 272 -56 -32 )
//!  ( 0 624 844 464 -320 -104 )
//!  ( 0 0 520 452 -80 -332 )
//!  ( 0 0 0 264 12 -96 )

// Can mix BandMatrix with other kinds of matrices:
std::cout<<"xm * m1 = \n"<<xm*m1<<std::endl;
//!  xm * m1 =
//!  6 6
//!  ( 82 48 -120 -348 -336 396 )
//!  ( 252 318 180 -128 -276 216 )
//!  ( 422 588 480 92 -216 36 )
//!  ( 592 858 780 312 -156 -144 )
//!  ( 762 1128 1080 532 -96 -324 )
//!  ( 932 1398 1380 752 -36 -504 )
tmv::UpperTriMatrix<double> um(xm);
std::cout<<"um + m1 = \n"<<um+m1<<std::endl;
//!  um + m1 =
//!  6 6
//!  ( 17 14 5 -6 -13 -22 )

```

```

    //! ( 20 25 16 1 -8 -17 )
    //! ( 0 24 27 12 -9 -12 )
    //! ( 0 0 24 23 2 -25 )
    //! ( 0 0 0 20 13 -14 )
    //! ( 0 0 0 0 12 -3 )
    tmv::LowerTriMatrix<double> lm(xm);
    lm *= m8;
    std::cout<<"lm *= m8 = \n"<<lm<<std::endl;
    //! lm *= m8 =
    //! 6 6
    //! ( 9 0 0 0 0 0 )
    //! ( 80 49 0 0 0 0 )
    //! ( 252 192 81 0 0 0 )
    //! ( 534 440 252 81 0 0 )
    //! ( 744 774 500 224 49 0 )
    //! ( 954 1064 782 396 120 9 )
    tmv::DiagMatrix<double> dm(xm);
    m1 *= dm;
    std::cout<<"m1 *= dm = \n"<<m1<<std::endl;
    //! m1 *= dm =
    //! 6 6
    //! ( 42 84 54 0 0 0 )
    //! ( 60 126 108 18 0 0 )
    //! ( 0 168 162 72 -42 0 )
    //! ( 0 0 216 126 0 -54 )
    //! ( 0 0 0 180 42 -36 )
    //! ( 0 0 0 0 84 -18 )
    return 0;
}
catch (tmv::Error& e)
{
    std::cerr<<e<<std::endl;
    return 1;
}

```

13.5 SymMatrix

```
#include "TMV.h"
// Note: extra include file for SymMatrix
#include "TMV_Sym.h"
// Also need to link with -ltmv_syband
#include <iostream>

int main() try
{
    // Several ways to create and initialize matrices:

    // Create with uninitialized values
    tmv::SymMatrix<double> m1(5);
    // In debug mode, the elements are all set to 888 to make it easier
    // to find errors related to not correctly initializing the matrix.
    std::cout<<"m1 = \n"<<m1<<std::endl;
    //! m1 =
    //! 5 5
    //! ( 888 888 888 888 888 )
    //! ( 888 888 888 888 888 )
    //! ( 888 888 888 888 888 )
    //! ( 888 888 888 888 888 )
    //! ( 888 888 888 888 888 )

    // Initialize with element access:
    for(size_t i=0;i<m1.nrows();i++)
        for(size_t j=0;j<m1.ncols();j++)
            if (i>=j) m1(i,j) = 2.*i-j*j+2.;
    std::cout<<"m1 = \n"<<m1<<std::endl;
    //! m1 =
    //! 5 5
    //! ( 2 4 6 8 10 )
    //! ( 4 3 5 7 9 )
    //! ( 6 5 2 4 6 )
    //! ( 8 7 4 -1 1 )
    //! ( 10 9 6 1 -6 )

    tmv::SymMatrix<double> m2(5,2.); // Create with all 2's.
    std::cout<<"m2 = \n"<<m2<<std::endl;
    //! m2 =
    //! 5 5
    //! ( 2 2 2 2 2 )
    //! ( 2 2 2 2 2 )
    //! ( 2 2 2 2 2 )
    //! ( 2 2 2 2 2 )
    //! ( 2 2 2 2 2 )

    // Create from given elements:
    double mm[25] = {1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,
        16,17,18,19,20,21,22,23,24,25};
    tmv::SymMatrix<double> m3(5,mm); // Default storage is Upper, ColMajor
    std::cout<<"m3 (Upper, ColMajor) = \n"<<m3<<std::endl;
    //! m3 (Upper, ColMajor) =
    //! 5 5
    //! ( 1 6 11 16 21 )
```

```

    //! ( 6 7 12 17 22 )
    //! ( 11 12 13 18 23 )
    //! ( 16 17 18 19 24 )
    //! ( 21 22 23 24 25 )
    tmv::SymMatrix<double, tmv::Upper, tmv::RowMajor> m4(5,mm);
    std::cout<<"m4 (Upper,RowMajor) = \n"<<m4<<std::endl;
    //! m4 (Upper,RowMajor) =
    //! 5 5
    //! ( 1 2 3 4 5 )
    //! ( 2 7 8 9 10 )
    //! ( 3 8 13 14 15 )
    //! ( 4 9 14 19 20 )
    //! ( 5 10 15 20 25 )
    tmv::SymMatrix<double, tmv::Lower, tmv::ColMajor> m5(5,mm);
    std::cout<<"m5 (Lower,ColMajor) = \n"<<m5<<std::endl;
    //! m5 (Lower,ColMajor) =
    //! 5 5
    //! ( 1 2 3 4 5 )
    //! ( 2 7 8 9 10 )
    //! ( 3 8 13 14 15 )
    //! ( 4 9 14 19 20 )
    //! ( 5 10 15 20 25 )
    tmv::SymMatrix<double, tmv::Lower, tmv::RowMajor> m6(5,mm);
    std::cout<<"m6 (Lower,RowMajor) = \n"<<m6<<std::endl;
    //! m6 (Lower,RowMajor) =
    //! 5 5
    //! ( 1 6 11 16 21 )
    //! ( 6 7 12 17 22 )
    //! ( 11 12 13 18 23 )
    //! ( 16 17 18 19 24 )
    //! ( 21 22 23 24 25 )

    // Make from the corresponding portion of a regular Matrix
    // Note – it copies from the correct triangle regardless of
    // the storage order of the two matrices.
    tmv::Matrix<double, tmv::RowMajor> xm(5,5);
    xm <<
        2, 5, 1, 9, 8,
        1, 5, 7, 2, 0,
        3, 9, 6, 8, 4,
        4, 2, 1, 9, 0,
        9, 8, 3, 7, 5;
    tmv::SymMatrix<double, tmv::Lower, tmv::ColMajor> m7(xm);
    std::cout<<"m7 (Lower,ColMajor) = "<<m7<<std::endl;
    //! m7 (Lower,ColMajor) = 5 5
    //! ( 2 1 3 4 9 )
    //! ( 1 5 9 2 8 )
    //! ( 3 9 6 1 3 )
    //! ( 4 2 1 9 7 )
    //! ( 9 8 3 7 5 )

    // Norms, etc.

    std::cout<<"Norm1(m1) = "<<Norm1(m1)<<std::endl;
    //! Norm1(m1) = 32

```

```

std::cout<<"Norm2(m1) = "<<Norm2(m1)<<std::endl;
///! Norm2(m1) = Warning:
///! Calling SymMatrix::Norm2 without previously calling DivideUsing(SV)
///! 24.4968
std::cout<<"m1.doNorm2() = "<<m1.doNorm2()<<std::endl;
///! m1.doNorm2() = 24.4968
std::cout<<"NormInf(m1) = "<<NormInf(m1)<<std::endl;
///! NormInf(m1) = 32
std::cout<<"NormF(m1) = "<<NormF(m1);
std::cout<<" = "<<Norm(m1)<<std::endl;
///! NormF(m1) = 30.0333 = 30.0333
std::cout<<"MaxAbsElement(m1) = "<<MaxAbsElement(m1)<<std::endl;
///! MaxAbsElement(m1) = 10
std::cout<<"Trace(m1) = "<<Trace(m1)<<std::endl;
///! Trace(m1) = 0
std::cout<<"Det(m1) = "<<Det(m1)<<std::endl;
///! Det(m1) = 4866

```

// Views:

```

std::cout<<"m1 = \n"<<m1<<std::endl;
///! m1 =
///! 5 5
///! ( 2 4 6 8 10 )
///! ( 4 3 5 7 9 )
///! ( 6 5 2 4 6 )
///! ( 8 7 4 -1 1 )
///! ( 10 9 6 1 -6 )
std::cout<<"m1.diag() = "<<m1.diag()<<std::endl;
///! m1.diag() = 5 ( 2 3 2 -1 -6 )
std::cout<<"m1.diag(1) = "<<m1.diag(1)<<std::endl;
///! m1.diag(1) = 4 ( 4 5 4 1 )
std::cout<<"m1.diag(-1) = "<<m1.diag(-1)<<std::endl;
///! m1.diag(-1) = 4 ( 4 5 4 1 )
std::cout<<"m1.subSymMatrix(0,3) = \n"<<m1.subSymMatrix(0,3)<<std::endl;
///! m1.subSymMatrix(0,3) =
///! 3 3
///! ( 2 4 6 )
///! ( 4 3 5 )
///! ( 6 5 2 )
std::cout<<"m1.upperTri() = \n"<<m1.upperTri()<<std::endl;
///! m1.upperTri() =
///! 5 5
///! ( 2 4 6 8 10 )
///! ( 0 3 5 7 9 )
///! ( 0 0 2 4 6 )
///! ( 0 0 0 -1 1 )
///! ( 0 0 0 0 -6 )
std::cout<<"m1.lowerTri(tmv::UnitDiag) = \n"<<
    m1.lowerTri(tmv::UnitDiag)<<std::endl;
///! m1.lowerTri(tmv::UnitDiag) =
///! 5 5
///! ( 1 0 0 0 0 )
///! ( 4 1 0 0 0 )
///! ( 6 5 1 0 0 )

```

```

//! ( 8 7 4 1 0 )
//! ( 10 9 6 1 1 )

// Fortran-style indexing:

tmv::SymMatrix<double, tmv::Upper, tmv::ColMajor, tmv::FortranStyle> fm1 = m1;
std::cout<<"fm1 = m1 = \n"<<fm1<<std::endl;
//! fm1 = m1 =
//! 5 5
//! ( 2 4 6 8 10 )
//! ( 4 3 5 7 9 )
//! ( 6 5 2 4 6 )
//! ( 8 7 4 -1 1 )
//! ( 10 9 6 1 -6 )
std::cout<<"fm1(1,1) = "<<fm1(1,1)<<std::endl;
//! fm1(1,1) = 2
std::cout<<"fm1(4,3) = "<<fm1(4,3)<<std::endl;
//! fm1(4,3) = 4
std::cout<<"fm1.subSymMatrix(1,3) = \n"<<fm1.subSymMatrix(1,3)<<std::endl;
//! fm1.subSymMatrix(1,3) =
//! 3 3
//! ( 2 4 6 )
//! ( 4 3 5 )
//! ( 6 5 2 )

// Matrix arithmetic:

tmv::SymMatrix<double> m1pm3 = m1 + m3;
std::cout<<"m1 + m3 = \n"<<m1pm3<<std::endl;
//! m1 + m3 =
//! 5 5
//! ( 3 10 17 24 31 )
//! ( 10 10 17 24 31 )
//! ( 17 17 15 22 29 )
//! ( 24 24 22 18 25 )
//! ( 31 31 29 25 19 )
// Works correctly even if matrices are stored in different order:
tmv::SymMatrix<double> m3pm4 = m3 + m4;
std::cout<<"m3 + m4 = \n"<<m3pm4<<std::endl;
//! m3 + m4 =
//! 5 5
//! ( 2 8 14 20 26 )
//! ( 8 14 20 26 32 )
//! ( 14 20 26 32 38 )
//! ( 20 26 32 38 44 )
//! ( 26 32 38 44 50 )
tmv::SymMatrix<double> m4pm5 = m4 + m5;
std::cout<<"m4 + m5 = \n"<<m4pm5<<std::endl;
//! m4 + m5 =
//! 5 5
//! ( 2 4 6 8 10 )
//! ( 4 14 16 18 20 )
//! ( 6 16 26 28 30 )
//! ( 8 18 28 38 40 )

```

```

//! ( 10 20 30 40 50 )

m1 *= 2.;
std::cout<<"m1 *= 2 = \n"<<m1<<std::endl;
//! m1 *= 2 =
//! 5 5
//! ( 4 8 12 16 20 )
//! ( 8 6 10 14 18 )
//! ( 12 10 4 8 12 )
//! ( 16 14 8 -2 2 )
//! ( 20 18 12 2 -12 )

m1 += m4;
std::cout<<"m1 += m4 = \n"<<m1<<std::endl;
//! m1 += m4 =
//! 5 5
//! ( 5 10 15 20 25 )
//! ( 10 13 18 23 28 )
//! ( 15 18 17 22 27 )
//! ( 20 23 22 17 22 )
//! ( 25 28 27 22 13 )

// Vector outer product
tmv::Vector<double> v = xm.col(0);
tmv::SymMatrix<double> vv = v^v;
std::cout<<"v = "<<v<<std::endl;
//! v = 5 ( 2 1 3 4 9 )
std::cout<<"v^v = \n"<<vv<<std::endl;
//! v^v =
//! 5 5
//! ( 4 2 6 8 18 )
//! ( 2 1 3 4 9 )
//! ( 6 3 9 12 27 )
//! ( 8 4 12 16 36 )
//! ( 18 9 27 36 81 )

// SymMatrix * Vector product
std::cout<<"m1 * v = "<<m1*v<<std::endl;
//! m1 * v = 5 ( 370 431 430 395 364 )
std::cout<<"v * m1 = "<<v*m1<<std::endl;
//! v * m1 = 5 ( 370 431 430 395 364 )

// SymMatrix * Matrix product
tmv::Matrix<double> mlxm = m1 * xm;
std::cout<<"m1 * xm = \n"<<mlxm<<std::endl;
//! m1 * xm =
//! 5 5
//! ( 370 450 260 540 225 )
//! ( 431 547 316 663 292 )
//! ( 430 578 346 694 323 )
//! ( 395 623 396 709 358 )
//! ( 364 656 444 786 373 )
// Note: the product of two symmetric matrices is not symmetric!:
tmv::Matrix<double> m1m5 = m1 * m5;
std::cout<<"m1 * m5 = \n"<<m1m5<<std::endl;
//! m1 * m5 =

```



```

//! 5 5
//! ( 275 630 945 1200 1375 )
//! ( 322 742 1110 1406 1610 )
//! ( 325 760 1123 1418 1625 )
//! ( 310 750 1098 1358 1550 )
//! ( 315 790 1153 1408 1575 )

// Complex matrices:

tmv::SymMatrix<std::complex<double>> cml = m1 * std::complex<double>(1,2);
std::cout<<"cml = m1 * (1+2i) = \n"<<cml<<std::endl;
//! cml = m1 * (1+2i) =
//! 5 5
//! ( (5,10) (10,20) (15,30) (20,40) (25,50) )
//! ( (10,20) (13,26) (18,36) (23,46) (28,56) )
//! ( (15,30) (18,36) (17,34) (22,44) (27,54) )
//! ( (20,40) (23,46) (22,44) (17,34) (22,44) )
//! ( (25,50) (28,56) (27,54) (22,44) (13,26) )
std::cout<<"cml.conjugate() = \n"<<cml.conjugate()<<std::endl;
//! cml.conjugate() =
//! 5 5
//! ( (5,-10) (10,-20) (15,-30) (20,-40) (25,-50) )
//! ( (10,-20) (13,-26) (18,-36) (23,-46) (28,-56) )
//! ( (15,-30) (18,-36) (17,-34) (22,-44) (27,-54) )
//! ( (20,-40) (23,-46) (22,-44) (17,-34) (22,-44) )
//! ( (25,-50) (28,-56) (27,-54) (22,-44) (13,-26) )
std::cout<<"cml.transpose() = \n"<<cml.transpose()<<std::endl;
//! cml.transpose() =
//! 5 5
//! ( (5,10) (10,20) (15,30) (20,40) (25,50) )
//! ( (10,20) (13,26) (18,36) (23,46) (28,56) )
//! ( (15,30) (18,36) (17,34) (22,44) (27,54) )
//! ( (20,40) (23,46) (22,44) (17,34) (22,44) )
//! ( (25,50) (28,56) (27,54) (22,44) (13,26) )
std::cout<<"cml.adjoint() = \n"<<cml.adjoint()<<std::endl;
//! cml.adjoint() =
//! 5 5
//! ( (5,-10) (10,-20) (15,-30) (20,-40) (25,-50) )
//! ( (10,-20) (13,-26) (18,-36) (23,-46) (28,-56) )
//! ( (15,-30) (18,-36) (17,-34) (22,-44) (27,-54) )
//! ( (20,-40) (23,-46) (22,-44) (17,-34) (22,-44) )
//! ( (25,-50) (28,-56) (27,-54) (22,-44) (13,-26) )
std::cout<<"cml.realPart() = \n"<<cml.realPart()<<std::endl;
//! cml.realPart() =
//! 5 5
//! ( 5 10 15 20 25 )
//! ( 10 13 18 23 28 )
//! ( 15 18 17 22 27 )
//! ( 20 23 22 17 22 )
//! ( 25 28 27 22 13 )
std::cout<<"cml.imagPart() = \n"<<cml.imagPart()<<std::endl;
//! cml.imagPart() =
//! 5 5
//! ( 10 20 30 40 50 )
//! ( 20 26 36 46 56 )

```

```

//! ( 30 36 34 44 54 )
//! ( 40 46 44 34 44 )
//! ( 50 56 54 44 26 )
std::cout<<"Norm(cm1) = "<<Norm(cm1)<<std::endl;
//! Norm(cm1) = 227.035

tmv::Vector<std::complex<double>> cv = v * std::complex<double>(-2,1);
cm1 += cv^cv;
std::cout<<"cm1 += cv^cv = \n"<<cm1<<std::endl;
//! cm1 += cv^cv =
//! 5 5
//! ( (17,-6) (16,12) (33,6) (44,8) (79,-22) )
//! ( (16,12) (16,22) (27,24) (35,30) (55,20) )
//! ( (33,6) (27,24) (44,-2) (58,-4) (108,-54) )
//! ( (44,8) (35,30) (58,-4) (65,-30) (130,-100) )
//! ( (79,-22) (55,20) (108,-54) (130,-100) (256,-298) )
cm1 += v^v;
std::cout<<"cm1 += v^v = \n"<<cm1<<std::endl;
//! cm1 += v^v =
//! 5 5
//! ( (21,-6) (18,12) (39,6) (52,8) (97,-22) )
//! ( (18,12) (17,22) (30,24) (39,30) (64,20) )
//! ( (39,6) (30,24) (53,-2) (70,-4) (135,-54) )
//! ( (52,8) (39,30) (70,-4) (81,-30) (166,-100) )
//! ( (97,-22) (64,20) (135,-54) (166,-100) (337,-298) )

tmv::Matrix<std::complex<double>> cx(5,2);
cx.col(0) = v;
cx.col(1) = cv;
cm1 -= cx*cx.transpose();
std::cout<<"cm1 -= cx*cx.transpose() = \n"<<cm1<<std::endl;
//! cm1 -= cx*cx.transpose() =
//! 5 5
//! ( (5,10) (10,20) (15,30) (20,40) (25,50) )
//! ( (10,20) (13,26) (18,36) (23,46) (28,56) )
//! ( (15,30) (18,36) (17,34) (22,44) (27,54) )
//! ( (20,40) (23,46) (22,44) (17,34) (22,44) )
//! ( (25,50) (28,56) (27,54) (22,44) (13,26) )

tmv::HermMatrix<std::complex<double>> cm2 = m1;
cm2.upperTri().offDiag() *= std::complex<double>(1,2);
std::cout<<"cm2 = \n"<<cm2<<std::endl;
//! cm2 =
//! 5 5
//! ( (5,0) (10,20) (15,30) (20,40) (25,50) )
//! ( (10,-20) (13,0) (18,36) (23,46) (28,56) )
//! ( (15,-30) (18,-36) (17,0) (22,44) (27,54) )
//! ( (20,-40) (23,-46) (22,-44) (17,0) (22,44) )
//! ( (25,-50) (28,-56) (27,-54) (22,-44) (13,0) )
std::cout<<"cm2.conjugate() = \n"<<cm2.conjugate()<<std::endl;
//! cm2.conjugate() =
//! 5 5
//! ( (5,-0) (10,-20) (15,-30) (20,-40) (25,-50) )
//! ( (10,20) (13,-0) (18,-36) (23,-46) (28,-56) )
//! ( (15,30) (18,36) (17,-0) (22,-44) (27,-54) )
//! ( (20,40) (23,46) (22,44) (17,-0) (22,-44) )

```

```

//! ( (25,50) (28,56) (27,54) (22,44) (13,-0) )
std::cout<<"cm2.transpose() = \n"<<cm2.transpose()<<std::endl;
//! cm2.transpose() =
//! 5 5
//! ( (5,-0) (10,-20) (15,-30) (20,-40) (25,-50) )
//! ( (10,20) (13,-0) (18,-36) (23,-46) (28,-56) )
//! ( (15,30) (18,36) (17,-0) (22,-44) (27,-54) )
//! ( (20,40) (23,46) (22,44) (17,-0) (22,-44) )
//! ( (25,50) (28,56) (27,54) (22,44) (13,-0) )
std::cout<<"cm2.adjoint() = \n"<<cm2.adjoint()<<std::endl;
//! cm2.adjoint() =
//! 5 5
//! ( (5,0) (10,20) (15,30) (20,40) (25,50) )
//! ( (10,-20) (13,0) (18,36) (23,46) (28,56) )
//! ( (15,-30) (18,-36) (17,0) (22,44) (27,54) )
//! ( (20,-40) (23,-46) (22,-44) (17,0) (22,44) )
//! ( (25,-50) (28,-56) (27,-54) (22,-44) (13,0) )
std::cout<<"cm2.realPart() = \n"<<cm2.realPart()<<std::endl;
//! cm2.realPart() =
//! 5 5
//! ( 5 10 15 20 25 )
//! ( 10 13 18 23 28 )
//! ( 15 18 17 22 27 )
//! ( 20 23 22 17 22 )
//! ( 25 28 27 22 13 )
// imagPart is invalid for hermitian matrix, since the result
// would be anti-symmetric, which we don't have as a matrix type.
std::cout<<"Norm(cm2) = "<<Norm(cm2)<<std::endl;
//! Norm(cm2) = 218.589

cm2 += cv^cv.conjugate();
std::cout<<"cm2 += cv^cv.conjugate() = \n"<<cm2<<std::endl;
//! cm2 += cv^cv.conjugate() =
//! 5 5
//! ( (25,0) (20,20) (45,30) (60,40) (115,50) )
//! ( (20,-20) (18,0) (33,36) (43,46) (73,56) )
//! ( (45,-30) (33,-36) (62,0) (82,44) (162,54) )
//! ( (60,-40) (43,-46) (82,-44) (97,0) (202,44) )
//! ( (115,-50) (73,-56) (162,-54) (202,-44) (418,0) )
cm2 += v^v;
std::cout<<"cm2 += v^v = \n"<<cm2<<std::endl;
//! cm2 += v^v =
//! 5 5
//! ( (29,0) (22,20) (51,30) (68,40) (133,50) )
//! ( (22,-20) (19,0) (36,36) (47,46) (82,56) )
//! ( (51,-30) (36,-36) (71,0) (94,44) (189,54) )
//! ( (68,-40) (47,-46) (94,-44) (113,0) (238,44) )
//! ( (133,-50) (82,-56) (189,-54) (238,-44) (499,0) )
cm2 -= cx*cx.adjoint();
std::cout<<"cm2 -= cx*cx.adjoint() = \n"<<cm2<<std::endl;
//! cm2 -= cx*cx.adjoint() =
//! 5 5
//! ( (5,0) (10,20) (15,30) (20,40) (25,50) )
//! ( (10,-20) (13,0) (18,36) (23,46) (28,56) )
//! ( (15,-30) (18,-36) (17,0) (22,44) (27,54) )
//! ( (20,-40) (23,-46) (22,-44) (17,0) (22,44) )

```

```

    //! ( (25,-50) (28,-56) (27,-54) (22,-44) (13,0) )

    // Can mix SymMatrix with other kinds of matrices:
    tmv::UpperTriMatrix<double> um(xm);
    std::cout<<"um + m1 = \n"<<um+m1<<std::endl;
    //! um + m1 =
    //! 5 5
    //! ( 7 15 16 29 33 )
    //! ( 10 18 25 25 28 )
    //! ( 15 18 23 30 31 )
    //! ( 20 23 22 26 22 )
    //! ( 25 28 27 22 18 )
    tmv::DiagMatrix<double> dm(xm);
    std::cout<<"dm * m1 = \n"<<dm*m1<<std::endl;
    //! dm * m1 =
    //! 5 5
    //! ( 10 20 30 40 50 )
    //! ( 50 65 90 115 140 )
    //! ( 90 108 102 132 162 )
    //! ( 180 207 198 153 198 )
    //! ( 125 140 135 110 65 )

    return 0;
}
catch (tmv::Error& e)
{
    std::cerr<<e<<std::endl;
    return 1;
}

```

14 Known Bugs and Deficiencies (aka To Do List)

Here is a list of various problems with the current version of the TMV library. These are mostly features that are not yet included, rather than bugs per se. If you are looking for issues with getting TMV compiled correctly, you should probably see §12.7 instead where I discuss various known problems that are specific to particular compilers or BLAS or LAPACK distributions.

If you find something to add to this list, or if you want me to bump something to the top of the list, let me know. Not that the list is currently in any kind of priority order, but you know what I mean. Please post a feature request at <http://code.google.com/p/tmv-cpp/issues>, or email the discussion group about what you need at <mailto:tmv-discuss@googlegroups.com>.

1. Symmetric arithmetic

When writing complicated equations involving complex symmetric or hermitian matrices, you may find that an equation that seems perfectly ok does not compile. The reason for this problem is explained in §7.4 in some detail, so you should read about it there. But basically, the workaround is usually to break your equation up into smaller steps that do not require the code to explicitly instantiate any matrices. For example: (this is the example from §7.4)

```
s3 += x*s1 + s2;
```

will not compile if s_1 , s_2 , and s_3 are all complex symmetric, even though it is valid, mathematically. Rewriting this as:

```
s3 += x*s1;  
s3 += s2;
```

will compile and work correctly.

2. Eigenvalues and eigenvectors

The code only finds eigenvalues and eigenvectors for hermitian matrices. I need to add the non-hermitian routines.

3. More special matrix varieties

Block-diagonal, generic sparse (probably both row-based and column-based), block sparse, symmetric and hermitian block diagonal, small symmetric and hermitian... Maybe skew-symmetric and skew-hermitian. Are these worth adding? Let me know.

4. Packed storage

Triangle and symmetric matrices. can be stored in (approximately) half the memory as a full $N \times N$ matrix using what is known as packed storage. There are BLAS routines for dealing with these packed storage matrices, but I don't yet have the ability to create/use such matrices.

5. List initialization for special matrices

I'm not sure what makes the most sense for the list initialization of special matrices. The most straightforward way to implement it would be to just assign the elements in order that they are stored in memory. But what do I do for the memory locations that do not correspond to a location in memory. For example, for an `UpperTriMatrix`, should the initializer be:

```
U << 1, 2, 3,  
      0, 4, 5,  
      0, 0, 6;
```

or

```
U << 1, 2, 3,
      4, 5,
      6;
```

Likewise, a `BandMatrix` has confusing locations for elements that are allocated in memory, but are not actually part of the matrix. Should I require these to be in the list of values? If you have any thoughts about this, feel free to email the discussion group (<mailto:tmv-discuss@googlegroups.com>) about them.

6. Hermitian eigenvector algorithm

There is a faster algorithm for calculating eigenvectors of a hermitian matrix given the eigenvalues, which uses a technique known as a “Relatively Robust Representation”. The native TMV code does not use this, so it is slower than a compilation which calls the LAPACK routine. I think this is the only routine for which the LAPACK version is still significantly faster than the native TMV code.

7. Row-major Bunch-Kaufman

The Bunch-Kaufman decomposition for row-major symmetric/hermitian matrices is currently LDL^\dagger , rather than $L^\dagger DL$. The latter should be somewhat (30%?) faster. The current LDL^\dagger algorithm is the faster algorithm for column-major matrices.²⁶

8. Conditions

Currently, SVD is the only decomposition that calculates the condition of a matrix (specifically, the 2-condition). LAPACK has routines to calculate the 1- and infinity-condition from an LU decomposition (and others). I should add a similar capability.

9. Division error estimates

LAPACK provides these. It would be nice to add something along the same lines.

10. Equilibrate matrices

LAPACK can equilibrate matrices before division. Again, I should include this feature too. Probably as an option (since most matrices don’t need it) as something like `m.Equilibrate()` before calling a division routine.

11. OpenMP

I have rewritten the basic Matrix product algorithm to exploit multiple threads using OpenMP pragma’s. This gives a good boost in speed for non-BLAS compilations, especially since most of the calculation time for the higher-level algorithms is in the MultMM functions. Also, the SVD divide and conquer algorithm uses OpenMP.

But the SVD only uses half of the potential for OpenMP at the moment. I need to reorganize the algorithm a bit to make it more amenable to being parallelized, but it is certainly doable.

Also I should go through the whole code to see which other algorithms might benefit from parallelization. I suspect that all of the so-called “Level 3 BLAS” functions will be amenable to parallelization. I’m not sure which higher level functions (i.e. those normally in a LAPACK library) can be parallelized, but probably some of them.

12. Check for memory throws

Many algorithms are able to increase their speed by allocating extra workspace. Usually this workspace is significantly smaller than the matrix being worked on, so we assume there is enough space for these

²⁶ These comments hold when the storage of the symmetric matrix is in the lower triangle - it is the opposite for storage in the upper triangle.

allocations. However, I should add try-catch blocks to catch any out-of-memory throws and use a less memory-intensive algorithm when necessary.

13. **Permutations**

I should make it easier to use the permutations that are output from `sort()`, as well as from LU and QRP decompositions. I'd like to write `m *= P`, rather than `m.permuteCols(p)`. This just needs to be a wrapper class for the functions that already exist, so it shouldn't be too hard to write.

15 History

Here is a list of the changes from version to version. Whenever a change is not backward compatible, meaning that code using the previous version might be broken, I mark the item with a \times bullet rather than the usual \bullet to indicate this. Also, the bulleted lists are not comprehensive. In most cases, new versions fix minor bugs that I find in the old version. I only list the more significant changes.

Version 0.1 The first matrix/vector library I wrote. It wasn't very good, really. It had a lot of the functionality I needed, like mixing complex/real, SV decomposition, LU decomposition, etc. But it wasn't at all fast for large matrices. It didn't call BLAS or LAPACK, nor were the native routines very well optimized. Also, while it had vector views for rows and columns, it didn't have matrix views for things like transpose. Nor did it have any delayed arithmetic evaluation. And there were no special matrices.

I didn't actually name this one 0.1 until I had what I called version 0.3.

Version 0.2 This was also not named version 0.2 until after the fact. It had most of the current interface for regular Matrix and Vector operations. I added Upper/Lower TriMatrix and DiagMatrix. It also had matrix views and matrix composites to delay arithmetic evaluation. The main problem was that it was still slow. I hadn't included any BLAS calls yet. And while the internal routines at least used algorithms that used unit strides whenever possible, they didn't do any blocking or recursion, which are key for large matrices.

Version 0.3 Finally, I actually named this one 0.3 at the time. The big addition here was BLAS and LAPACK calls, which helped me to realize how slow my internal code really was (although I hadn't updated them to block or recursive algorithms yet). I also added BandMatrix.

Version 0.4 The new version number here was because I needed some added functionality for a project I was working on. In retrospect, it really only deserves a 0.01 increment, since the changes weren't that big. But, oh well.

- Added QR_Downdate. (This was the main new functionality I needed.)
- Improved the numerical accuracy of the QRP decomposition.
- Added the possibility of not storing U,V for the SVD.
- Greatly improved the test suite, and consequently found and corrected a few bugs.
- Added some arithmetic functionality that had been missing (like $m += L*U$).

Version 0.5 The new symmetric matrix classes precipitated a major version number update. I also sped up a lot of the algorithms:

- Added SymMatrix, HermMatrix, and all associated functionality.
- Added blocked versions of most of the algorithms, so the non-LAPACK code runs a lot faster.
- Allowed for loose QRP decomposition.
- Added divideInPlace().

Version 0.51 Some minor improvements:

- Sped up some functions like matrix addition and assignment by adding the LinearView method.
- Added QR_Update, and improved the QR_Downdate algorithm.
- Blocked some more algorithms like TriMatrix multiplication/division, so non-BLAS code runs significantly faster (but still slower than BLAS).

Version 0.52 The first “public” release! And correspondingly, the first with documentation and a web site. A few other people had used previous versions, but since the only documentation was my comments in the .h files, it wasn’t all that user-friendly.

- Added `saveDiv()` and related methods for division control. Also changed the default behavior from saving the decomposition to not saving it.
- Added in-place versions of the algorithms for $S = L^\dagger L$ and $S = LL^\dagger$.

Version 0.53 By popular demand (well, a request by Fritz Stabenau, at least):

- Added the Fortran-style indexing.

Version 0.54 Inspired by my to-do list, which I wrote for Version 0.52, I tackled a few of the items on the list and addressed some issues that people had been having with compiling:

- × Changed from a rudimentary exception scheme (with just one class - `tmv_exception`) to the current more full-featured exception hierarchy. Also added `auto_ptr` usage instead of bald pointers to make the exception throws memory-leak safe.
- Sped up `SymLUDiv` and `SymSVDiv` inverses.
- Added the possibility of compiling a small version of the library and test suite.
- × Consolidated `SymLUDiv` and `HermLUDiv` classes into just `SymLUDiv`, which now checks whether the matrix is hermitian automatically.
- Reduced the number of operations that make temporary matrices when multiple objects use the same storage.
- Specialized Tridiagonal \times Matrix calculation.
- Added `ElementProd` and `AddElementProd` functions for matrices.
- Added CLAPACK and ACML as known versions of LAPACK (the LAPACK calls had previously been specific to the Intel MKL version). Also made the file `TMV_Blas.h` much better organized, so it is easier to tailor the code to a different LAPACK distribution with different calling conventions.

Version 0.60 This revision merits a first-decimal-place increment, since I added a few big features. I also registered it with SourceForge, which is a pretty big milestone as well.

- Added `SmallVector` and `SmallMatrix` with all accompanying algorithms.
- Added `SymBandMatrix` and `HermBandMatrix` with all accompanying algorithms.
- Made arithmetic between any two special matrices compatible, so long as the operation is allowed given their respective shapes. e.g. `U += B` and `U *= B` are allowed if `B` is upper-banded. There are many other expressions that are now legal statements.
- × Changed `QR_Downdate()` to throw an exception rather than return false when it fails.
- Added the GPL License to the end of this documentation, and a copyright and GPL snippet into each file.
- × Changed the `-D` compiler options for changing which types get instantiated. The default is still to instantiate double and float. Now to turn these off, use `-DNO_INST_DOUBLE` and `-DNO_INST_FLOAT` respectively. To add int or long double instantiations, use `-DINST_INT` and `-DINST_LONGDOUBLE` respectively.
- Split up the library into `libtmv.a` and `libtmv_syband.a`. The latter includes everything for the `SymMatrix`, `HermMatrix`, `BandMatrix`, `SymBandMatrix`, and `HermBandMatrix` classes. The former includes everything else including `DiagMatrix` and `Upper/LowerTriMatrix`. Also, I got rid of `libsmalltmv.a`, which was basically the same as the new trimmed down `libtmv.a`.

Version 0.61 A number of updates mostly precipitated by feature requests by me in my own use of the library, as well as some from a few other users. I also did a complete systematic edit of the documentation which precipitated some more changes to make the UI a bit more intuitive.

- × Changed the default storage for the `Matrix` class to `ColMajor` rather than `RowMajor`, since it seems to be more common that column-major storage is the more efficient choice. Therefore, it makes sense to have this be the default.
- × Changed a lot of `size_t` parameters to `int` - mostly variables which are indices of a matrix or vector, like `i` and `j` in `m(i, j)`. These variables used to be of type `size_t`, and now they are `int`, even though they are usually required to be positive.

The reason for this change was that the implementation of these functions often involves multiplying the index by a step size, which is allowed to be negative (and hence is an `int`), so there were lots of casts to `int` for these variables. I decided that it would be better to simply have them be `int` in the first place.

The particular change that is most likely to require modification to existing code involves permutations, which used to be `size_t` arrays, and are now `int` arrays. So if you used them, you might need to change their declarations in your code to `(int [])`.

- × Removed `U.MakeUnitDiag`, since the name was counter-intuitive to what it actually did, and the functionality is more obvious through the `UpperTriMatrixViewOf` function.
- Sped up matrix multiplication for non-blas implementations, including openmp pragmas to allow for multiple threads on machines that support them. The code is now within a factor of 2 or 3 of a good optimized BLAS distribution. So it is still worth it to use BLAS if you have one available, but if you don't have one on your machine, the non-blas code is no longer a order of magnitude slower.
- Changed a few things which prevented Microsoft Visual C++ from compiling successfully. As far as I can tell, these are aspects in which VC++ is not fully compliant with the C++ standard, but there were work-arounds for all of them. Thanks to Andy Molloy for spearheading this effort and doing the lion's share of the work to make the code compatible with the VC++ compiler.

- × Removed the optional index parameter to the non-method versions of `MaxElement`, etc. i.e.

```
vmax = v.MaxElement(&imax)
```

will work, but not

```
vmax = MaxElement(v, &imax)
```

However the functional form without the `imax` parameter still works as before. Basically, this was just a semantic choice. It seems to me that the meaning of the method form with the index parameter is much clearer than the functional form with two arguments.

I also added an optional `scale` parameter to `m.NormSq(scale)`. Again, this optional parameter is only allowed in the method version, not the function `NormSq(m)`

- × Added the explicit decomposition routines. I also got rid of the `SVU`, `SVV` and `SVS` options for `m.divideUsing(...)`, since the point of these was to do the decomposition without calculation `U` and/or `V`. This is now done more intuitively with the explicit decomposition routines. I also added the (hermitian) eigenvalue/eigenvector routines which used to require using the division accessors in non-intuitive ways to calculate.
- Fixed a couple of places where underflow and overflow could cause problems. However, I have not put such very large or very small matrices into the test suite, so there are probably other places where these could be a problem. (Added this to the To Do list.)
- Updated the native TMV code for the singular value decomposition and hermitian eigenvalue calculation to use the divide-and-conquer algorithm.
- Added `m.logDet()` method. With very large matrices, it is common for the determinant to overflow, but often the logarithm of the determinant is sufficient. For example, one may be trying to minimize

a likelihood function that includes the determinant of a matrix. Minimizing the $\log(\text{likelihood})$ is equivalent, which thus involves the log of the determinant.

- × Changed `m.svd().getS()` to return a `DiagMatrix` rather than a `Vector`. (Sorry - this change eluded documentation until version 0.62.)

Version 0.62 This release contains a number of bug fixes and corrections of mistakes in the documentation. I also significantly revamped the `SmallMatrix` class to make it faster at the expense of a few fewer features.

- Corrected an error with `m.divIsSet()`.
- Corrected some errors in the documentation – Thanks to Jake VanderPlas for pointing these out.
- Improved the behavior of the division accessors, also at the suggestion of Jake VanderPlas. (Thanks again!) They used to require that the division object had already been set or they would give an error. So to set the SVD threshold for example, you had to write:

```
m.divideUsing(tmv::SV);  
m.setDiv();  
m.svd().thresh(1.e-5);
```

Jake rightly pointed out that this was not very intuitive. So now the accessors (like `m.svd()`) set the division object appropriately if it is not already set. The above task is now simply performed by just the single line:

```
m.svd().thresh(1.e-5);
```
- Added the `ListInit` method for initializing the values of a `Vector` or `Matrix`. Thanks to Paul Sarli for this suggestion. The implementation is based somewhat on a similar initialization syntax in the Little Template Library (LTL) by Niv Drory et al.
- × Significantly changed the `SmallMatrix` class. See §3.9, on the `SmallMatrix` class in the documentation for the full details of how it works now. But in brief, here are the main changes:
 - `SmallMatrix` and `SmallVector` no longer inherit from `GenMatrix` and `GenVector`.
 - This allowed us to remove all virtual functions and the corresponding `vtable`.
 - Improved the arithmetic so more of the routines correctly do the calculation inline to allow the compiler more opportunity to optimize the calculation.
 - Added inlined division and determinants.
 - Got rid of the “Small” views. Because I got rid of the inheritance in `SmallMatrix`, there is no longer a `GenSmallMatrix` class from which the views and regular objects can both inherit. This makes it a bit unwieldy to write all the different arithmetic combinations of how `SmallMatrix` and `SmallMatrixView` could combine. So I just made all views regular `VectorView` or `MatrixView` objects. These can be assigned to a `SmallVector` or `SmallMatrix`, or used in arithmetic calculations with them, so all the arithmetic statements that use `transpose()` or `row()` or `col()` (for example) will still work, but they might not be done inline anymore.
- Consolidated some of the old header files. Now all `TriMatrix` and `DiagMatrix` calculations are included as part of `TMV.h`, so there is no need to include `TMV_Tri.h` or `TMV_Diag.h` anymore. Although, these header files are still included for backwards compatibility.
- × Removed some of the `ViewOf` commands that were not very clear and which have other, clearer ways of doing the same thing:
 - `d = DiagMatrixViewOf(m)` should now be written
`d = DiagMatrixViewOf(m.diag())`.
 - `U = UpperTriMatrixViewOf(m)` should now be written
`U = m.upperTri()`.
 - `U = UpperTriMatrixViewOf(m, UnitDiag)` should now be written
`U = m.upperTri(UnitDiag)`.

- `U = UpperTriMatrixViewOf(U, UnitDiag)` should now be written
`U = U.viewAsUnitDiag()`.

- Tracked down the problems I had been having with the LAPACK `dstegr` and `sstegr` functions. The segmentation faults were because the algorithm writes to an element one past the end of the E vector. Apparently, this is a documented feature, but it doesn't make any sense to me why they would need to do this. I had previously discovered similar behavior with the `zhetrd` and `dgebrd` functions. Anyway, I fixed this by allocating a vector with one extra element and copying the values.

I also found that there are two separate problems with the `stegr` algorithm where it produces incorrect results sometimes. First, sometimes the `stegr` algorithm seems to return early - before the last update to the D vector and U matrix. I actually figured out how to correct the D vector from some values that remain in the workspace, but I couldn't figure out how to correct the U matrix. This is also a documented feature - `stegr` returns the value of 2 in the `info` variable.

Second, the `stegr` algorithm is apparently not very careful about avoiding overflow or underflow. I found NaN's in the output U matrix for particular input matrices from the float version of the algorithm, `sstegr`. The `info` variable is 0 on output in these cases, so unfortunately `stegr` does not detect the problem. So this has to be checked for by hand.

Since both of these problems are detectable based on the output values (including examining the workspace in the first case), TMV now checks for the problem, and if there was a problem, it calls the `stedc` routine instead.

- Tested the code for memory bugs with Paul Nettle's `mmgr.h` code. There were only a couple of minor memory leaks, which were fixed.
 - Fixed a problem with the OpenMP version of the code giving segmentation faults with `pgCC`.
 - Added the SCons installation method to automatically discover what BLAS and LAPACK libraries are installed on your computer. This should make it easier for most users to install an efficient version of the TMV library without having to fiddle with the Makefile.
 - Added compatibility for Fortran versions of BLAS and LAPACK, including header files that should work with them, `fblas.h` and `flapack.h` are in the `include` directory.
 - Added the CMake installation method. Thanks to Andy Molloy for providing this. Unfortunately, it is not as full featured as the SCons method, but it is pretty straightforward to get a basic TMV library compiled.
 - Added hyperlinks to the PDF documentation. This should make it much easier to locate a particular topic by simply clicking on the link in the Table of Contents. In PDF viewers that support it, you can also navigate using the "drawer" which will show the various sections and let you jump directly to them.
- Also, I added an index, complete with hyperlinks as well, that should make it easier to find a particular topic that you might need help with.
- Fixed other miscellaneous bugs.

Version 0.63 The biggest thing in this version is the new lowercase syntax for the methods, which was in response to feedback from a number of TMV users, who didn't like the old syntax. Hopefully, I'm not responding to the minority here, but I do agree that the new syntax is more conformant to common C++ standards. There are also a couple of bug fixes, and I've started using Google Code for my main website.

- × Changed the names of the methods to start with a lowercase. Many of the methods had started with an uppercase, but this style is not very popular, so I decided to go ahead and change them.
- The free functions all still start with a capital letter. Both use camelCase, with underscores after 2 or 3 letter abbreviations such as LU or SV to help them stand out (since the camel case doesn't do it). This

seems to be a more common standard in the C++ community and is more self-consistent than the style I had been using.

The old names are still here for this version, but are marked as deprecated. Compiling old code with g++ with the flag -Wno-deprecated-declarations will make it compile without warning. But the deprecated versions will disappear eventually, so you will want to convert your existing code over to the new names at some point.

- × Changed the syntax for list initialization. The old style was

```
v = tmv::ListInit, a0 , a1 , a2 , a3 ...
```

The new style is

```
v << a0 , a1 , a2 , a3 ...
```

Also accepted is

```
v << a0 << a1 << a2 << a3 ...
```

This notation seems clearer, and avoids the global variable `tmv::ListInit`. (Again, the old notation is still present, but deprecated.)

- Improved the speed of the native (i.e. non-BLAS) matrix multiplication algorithm. It is now generally within a factor of two or so of optimized BLAS implementations. So it's not such a big performance hit if you compile without BLAS anymore. But if you are using large matrices, then we still recommend compiling with an optimized BLAS library if possible.

The improvements for large matrices involved:

1. Using a blocked algorithm with copying the sub-matrices to new storage so the block calculations are contiguous in memory,
 2. Using optimized kernels for the single-block calculation, including using SSE and SSE2 intrinsics when available,
 3. Implementing complex calculations as sequential real operations, rather than try to do the complex arithmetic directly, which was significantly slower, and
 4. Using OpenMP to split the calculation into as many parts as there are available threads.
- Fixed a couple of bugs involving `SmallMatrix.m.setZero()` didn't return `*this` as it should have, and there were two `DoCopy` helper functions that caused ambiguity problems for some compilers.
 - Disabled the cases where `BlasRowMajor` was used in `CBlas` implementations, since I had problems with it on a particular BLAS version, and I'm not sure if the error is in my code or in that BLAS implementation. So, I removed these pathways for now until I can figure out what might be going on with it. Most calls were with `BlasColMajor` anyway, so this doesn't affect very many calls.
 - Fixed the return type of some methods that erroneously returned `CStyle` views when the source object was `FortranStyle`.
 - Added a way to access the version of TMV being used. You can do this either in C++ code with the function `TMV_Version()`. This returns a string that gives the version being used within your code. There is also an executable, `tmv-version`, that can be used to access the TMV version that is installed at the moment.
 - Added support for `ups`, which is probably completely irrelevant to the majority of users. But if you use `ups` for version control, you can use the command `scons install WITH_UPS=true` to configure `tmv` in `ups`.

Version 0.64 This update mostly addresses problems involving underflow and overflow. This had been item 13 on the to-do list (§14), and it was bumped to top priority from a bug report where underflow problems were leading to an infinite loop in the singular value decomposition code. So I added more matrices to the test suite – specifically a matrix with an extremely large dynamic range of singular values (which reproduced the

behavior of the bug report), a matrix that is nearly zero, and a matrix that is very close to overflow. These tests uncovered quite a few bugs in the code. Some from overflow and underflow issues of course, but also some imprecision issues, and some just plain bugs that I hadn't uncovered before. So I think this version should be much more stable for a wider range of input matrices.

There are also a couple of feature additions as well – most notably the Permutation class.

- Added a some very difficult matrices to the test suite, and found and fixed quite a few bugs as a result involving underflow, overflow, and loss of precision. It is worth noting that when compiled to use LAPACK, rather than the native TMV code, some of these tests (which are enabled with `XTEST=64`) fail. So if you find your program having problems with problematic input matrices, it might be worth compiling TMV with `WITH_LAPACK=false` to see if that helps.
 - Fixed a bug that `U(i, i)` was invalid if `U` is non-const and is `UnitDiag` even if the resulting value is not used in a mutable context. e.g. `T diag_i = U(i, i)`. This used to give an error. Now, `U(i, j)` returns a special reference that can always be used as a value, even if `i==j` and `U` is `UnitDiag`. However, if the user tries to modify the value through this reference type, then it gives an error. This was probably not causing anyone any problems, and the workarounds weren't too onerous, but it seemed worth getting the syntax right.
 - Added a new Permutation class to encapsulate permutation operations. This is basically a wrapper around the old functions like `m.permuteRows(p)`. Now you can write `m = P * m` instead. All of the routines which returned a permutation in an `int* p` object now returns it as a Permutation object. This includes `Vector::sort`, `LU-Decompose`, `QRP-Decompose`, and `LDL-Decompose`. The old functions still work, but are now undocumented.
 - × The `getP()` methods from the LU and QRP divider objects now return a Permutation rather than a `const int*`. As described above, this should make the code clearer when the permutation is used. However, it will break code that uses the old syntax.
 - Added `m.sumElements()` and `m.sumAbsElements()`.
 - Added `v.minAbs2Element()`, `v.maxAbs2Element()` and `m.maxAbs2Element()`.
 - Added `m.unitUpperTri()` and `m.unitLowerTri()`.
 - Made `Swap(m1, m2)` and `Swap(v1, v2)` efficient when the arguments are both complete matrices or vectors, rather than views, so they now take $O(1)$ time rather than $O(MN)$ and $O(N)$ respectively.
 - Added `cView()` and `fView()` methods to switch indexing styles of a matrix or vector.
 - Made arithmetic with `SmallVector` and `SmallMatrix` somewhat more flexible in that the composite objects now (again actually) derive from the normal `GenVector` and `GenMatrix` classes. This removes the "No implicit instantiation" item that had appeared in the list of limitations of the `Small` classes.
- I had removed that inheritance for version 0.62 in my attempt to get rid of the expensive virtual tables from all small operations. However, the `vtable` here is almost never expensive (unlike the one that had been in `SmallVector` and `SmallMatrix`), and it allows for a greater freedom of expression with the `Small` objects.
- Added a version of `MatrixViewOf` with arbitrary steps.
 - Added the three C preprocessor definitions: `TMV_MAJOR_VERSION`, `TMV_MINOR_VERSION`, and `TMV_VERSION_AT_LEAST(major, minor)` to help users customize their code depending on the version of TMV that is installed on a particular computer. Also documented the `TMV_Version` function and the `tmv-version` script, which have been around for a couple versions now, but weren't yet documented.
 - Fixed a bug in BLAS versions of `U*M` and `L*M` when `U` or `L` is real and `M` is complex.

- Fixed a bug in `m.det()` when `m` is a `SmallMatrix<T, 1, 1>`.
- Fixed a bug in an (apparently) uncommon pathway of `Rank2Update`. I say uncommon since it never turned up until I expanded the scope of the test suite to include more combinations of matrix types.
- Researched my `QRDowndate` algorithm to see if it is in the literature already. The only similar algorithm I could find was in Bojanczyk and Steinhardt (1991), a paper which seems to have been overlooked by the matrix algorithms community. It's not quite the same as mine in detail, but it is based on the same basic idea. I added this reference to the discussion in the documentation about my algorithm.

16 License

This software is licensed under the GNU General Public License. This license is sometimes referred to as “copyleft”. This gist of the license (made more specific in the following text) is that you are allowed to copy, modify, and use TMV code with basically no restrictions as far as your own use is concerned, but with a few restrictions regarding further distribution.

The biggest restriction is that if you distribute code that uses the TMV code, then you must distribute it with a copyleft license as well. In other words, you may not use the TMV library with commercial, proprietary software.

The other main restriction on such a distribution is that you must include this document, including the below license, along with your version of the code. If you modified any portions of the library, you must indicate your modifications in the accompanying documentation files.

More information may be obtained from <http://www.gnu.org/copyleft/>.

Version 2, June 1991

Copyright © 1989, 1991 Free Software Foundation, Inc.

51 Franklin St, Fifth Floor, Boston, MA 02110-1301, USA

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

Preamble

The licenses for most software are designed to take away your freedom to share and change it. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change free software—to make sure the software is free for all its users. This General Public License applies to most of the Free Software Foundation’s software and to any other program whose authors commit to using it. (Some other Free Software Foundation software is covered by the GNU Library General Public License instead.) You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for this service if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs; and that you know you can do these things.

To protect your rights, we need to make restrictions that forbid anyone to deny you these rights or to ask you to surrender the rights. These restrictions translate to certain responsibilities for you if you distribute copies of the software, or if you modify it.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must give the recipients all the rights that you have. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

We protect your rights with two steps: (1) copyright the software, and (2) offer you this license which gives you legal permission to copy, distribute and/or modify the software.

Also, for each author’s protection and ours, we want to make certain that everyone understands that there is no warranty for this free software. If the software is modified by someone else and passed on, we want its recipients to know that what they have is not the original, so that any problems introduced by others will not reflect on the original authors’ reputations.

Finally, any free program is threatened constantly by software patents. We wish to avoid the danger that redistributors of a free program will individually obtain patent licenses, in effect making the program proprietary. To prevent this, we have made it clear that any patent must be licensed for everyone’s free use or not licensed at all.

The precise terms and conditions for copying, distribution and modification follow.

GNU GENERAL PUBLIC LICENSE

TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION

0. This License applies to any program or other work which contains a notice placed by the copyright holder saying it may be distributed under the terms of this General Public License. The “Program”, below, refers to any such program or work, and a “work based on the Program” means either the Program or any derivative work under copyright law: that is to say, a work containing the Program or a portion of it, either verbatim or with modifications and/or translated into another language. (Hereinafter, translation is included without limitation in the term “modification”.) Each licensee is addressed as “you”.

Activities other than copying, distribution and modification are not covered by this License; they are outside its scope. The act of running the Program is not restricted, and the output from the Program is covered only if its contents constitute a work based on the Program (independent of having been made by running the Program). Whether that is true depends on what the Program does.

1. You may copy and distribute verbatim copies of the Program’s source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice and disclaimer of warranty; keep intact all the notices that refer to this License and to the absence of any warranty; and give any other recipients of the Program a copy of this License along with the Program.

You may charge a fee for the physical act of transferring a copy, and you may at your option offer warranty protection in exchange for a fee.

2. You may modify your copy or copies of the Program or any portion of it, thus forming a work based on the Program, and copy and distribute such modifications or work under the terms of Section 1 above, provided that you also meet all of these conditions:

- (a) You must cause the modified files to carry prominent notices stating that you changed the files and the date of any change.
- (b) You must cause any work that you distribute or publish, that in whole or in part contains or is derived from the Program or any part thereof, to be licensed as a whole at no charge to all third parties under the terms of this License.
- (c) If the modified program normally reads commands interactively when run, you must cause it, when started running for such interactive use in the most ordinary way, to print or display an announcement including an appropriate copyright notice and a notice that there is no warranty (or else, saying that you provide a warranty) and that users may redistribute the program under these conditions, and telling the user how to view a copy of this License. (Exception: if the Program itself is interactive but does not normally print such an announcement, your work based on the Program is not required to print an announcement.)

These requirements apply to the modified work as a whole. If identifiable sections of that work are not derived from the Program, and can be reasonably considered independent and separate works in themselves, then this License, and its terms, do not apply to those sections when you distribute them as separate works. But when you distribute the same sections as part of a whole which is a work based on the Program, the distribution of the whole must be on the terms of this License, whose permissions for other licensees extend to the entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest your rights to work written entirely by you; rather, the intent is to exercise the right to control the distribution of derivative or collective works based on the Program.

In addition, mere aggregation of another work not based on the Program with the Program (or with a work based on the Program) on a volume of a storage or distribution medium does not bring the other work under the scope of this License.

3. You may copy and distribute the Program (or a work based on it, under Section 2) in object code or executable form under the terms of Sections 1 and 2 above provided that you also do one of the following:

- (a) Accompany it with the complete corresponding machine-readable source code, which must be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
- (b) Accompany it with a written offer, valid for at least three years, to give any third party, for a charge no more than your cost of physically performing source distribution, a complete machine-readable copy of the corresponding source code, to be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
- (c) Accompany it with the information you received as to the offer to distribute corresponding source code. (This alternative is allowed only for noncommercial distribution and only if you received the program in object code or executable form with such an offer, in accord with Subsection b above.)

The source code for a work means the preferred form of the work for making modifications to it. For an executable work, complete source code means all the source code for all modules it contains, plus any associated interface definition files, plus the scripts used to control compilation and installation of the executable. However, as a special exception, the source code distributed need not include anything that is normally distributed (in either source or binary form) with the major components (compiler, kernel, and so on) of the operating system on which the executable runs, unless that component itself accompanies the executable.

If distribution of executable or object code is made by offering access to copy from a designated place, then offering equivalent access to copy the source code from the same place counts as distribution of the source code, even though third parties are not compelled to copy the source along with the object code.

- 4. You may not copy, modify, sublicense, or distribute the Program except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense or distribute the Program is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.
- 5. You are not required to accept this License, since you have not signed it. However, nothing else grants you permission to modify or distribute the Program or its derivative works. These actions are prohibited by law if you do not accept this License. Therefore, by modifying or distributing the Program (or any work based on the Program), you indicate your acceptance of this License to do so, and all its terms and conditions for copying, distributing or modifying the Program or works based on it.
- 6. Each time you redistribute the Program (or any work based on the Program), the recipient automatically receives a license from the original licensor to copy, distribute or modify the Program subject to these terms and conditions. You may not impose any further restrictions on the recipients' exercise of the rights granted herein. You are not responsible for enforcing compliance by third parties to this License.
- 7. If, as a consequence of a court judgment or allegation of patent infringement or for any other reason (not limited to patent issues), conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot distribute so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not distribute the Program at all. For example, if a patent license would not permit royalty-free redistribution of the Program by all those who receive copies directly or indirectly through you, then the only way you could satisfy both it and this License would be to refrain entirely from distribution of the Program.

If any portion of this section is held invalid or unenforceable under any particular circumstance, the balance of the section is intended to apply and the section as a whole is intended to apply in other circumstances.

It is not the purpose of this section to induce you to infringe any patents or other property right claims or to contest validity of any such claims; this section has the sole purpose of protecting the integrity of the

free software distribution system, which is implemented by public license practices. Many people have made generous contributions to the wide range of software distributed through that system in reliance on consistent application of that system; it is up to the author/donor to decide if he or she is willing to distribute software through any other system and a licensee cannot impose that choice.

This section is intended to make thoroughly clear what is believed to be a consequence of the rest of this License.

8. If the distribution and/or use of the Program is restricted in certain countries either by patents or by copyrighted interfaces, the original copyright holder who places the Program under this License may add an explicit geographical distribution limitation excluding those countries, so that distribution is permitted only in or among countries not thus excluded. In such case, this License incorporates the limitation as if written in the body of this License.
9. The Free Software Foundation may publish revised and/or new versions of the General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies a version number of this License which applies to it and “any later version”, you have the option of following the terms and conditions either of that version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of this License, you may choose any version ever published by the Free Software Foundation.

10. If you wish to incorporate parts of the Program into other free programs whose distribution conditions are different, write to the author to ask for permission. For software which is copyrighted by the Free Software Foundation, write to the Free Software Foundation; we sometimes make exceptions for this. Our decision will be guided by the two goals of preserving the free status of all derivatives of our free software and of promoting the sharing and reuse of software generally.

NO WARRANTY

11. BECAUSE THE PROGRAM IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM “AS IS” WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.
12. IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR REDISTRIBUTE THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

END OF TERMS AND CONDITIONS

This library is largely based upon work supported by the National Aeronautics and Space Administration (NASA) under Grant No NAG5-10924, a Keck Foundation grant, and the University of Pennsylvania. Any opinions, findings and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of NASA, the Keck Foundation, or the University of Pennsylvania. This software is made available to you on an “as is” basis with no representations or warranties, express or implied, including but not limited to any warranty of performance, merchantability, fitness for a particular purpose, commercial utility, non-infringement or title. Neither the author nor the organizations providing the support under which the work was developed will be liable to you or any third party with respect to any claim arising from your further development of the software or any products related to or derived from the software, or for lost profits, business interruption, or indirect special or consequential damages of any kind.